

Problem -1

Training data $\rightarrow (x_1, y_1), \dots, (x_n, y_n) \quad x_i \in \mathbb{R}^d$

$$f_{\beta}(x) = \langle \beta, x \rangle$$

$$L_{wss}(\beta) = \sum_{i=1}^n w_i (\langle \beta, x_i \rangle - y_i)^2$$

a) Gradient of $L_{wss}(\beta)$

$$\frac{\partial}{\partial \beta} L_{wss}(\beta) = \frac{\partial}{\partial \beta} \sum_{i=1}^n (w_i (\langle \beta, x_i \rangle - y_i)^2)$$

$$= \sum_{i=1}^n \frac{\partial}{\partial \beta} \left[w_i (\langle \beta, x_i \rangle - y_i)^2 \right]$$

$$= \sum_{i=1}^n 2 w_i (\langle \beta, x_i \rangle - y_i) x_i$$

$$\therefore \boxed{\nabla_{\beta} L_{wss}(\beta) = 2 \sum_{i=1}^n w_i (\langle \beta, x_i \rangle - y_i) x_i}$$

⑥ closed form minimizer β^*

$$\nabla_{\beta} L_{\text{WSS}}(\beta) = 0$$

$$2 \sum_{i=1}^n w_i (\langle \beta^*, x_i \rangle - y_i) x_{ij} = 0 \quad \forall j=1, 2, \dots, d$$

$$\sum_{i=1}^n w_i (\langle \beta^*, x_i \rangle - y_i) x_{ij} = 0$$

$$\Rightarrow \sum_{i=1}^n w_i (\langle \beta^*, x_i \rangle) x_{ij} = \sum_{i=1}^n w_i x_{ij} y_i \quad (j=1, \dots, d)$$

$$\Rightarrow \sum_{i=1}^n w_i x_{ij} x_i^T \beta^* = \sum_{i=1}^n w_i x_{ij} y_i \quad (j=1, \dots, d)$$

$$\Rightarrow (X^T W X) \beta^* = X^T W y.$$

$$\Rightarrow \boxed{\beta^* = (X^T W X)^{-1} X^T W y}$$

Problem -2

(a) Minimize $L(m) = \sum_{i=1}^n (y_i - m)^2$

$$L(m) = \sum_{i=1}^n ((y_i - \bar{y}) + (\bar{y} - m))^2$$

$$= \sum_{i=1}^n (y_i - \bar{y})^2 + 2(\bar{y} - m) \sum_{i=1}^n (y_i - \bar{y}) + n(\bar{y} - m)^2$$

$$\sum_{i=1}^n (y_i - \bar{y}) = 0$$

$$\Rightarrow L(m) = \underbrace{\sum_{i=1}^n (y_i - \bar{y})^2}_{\text{Constant w.r.t } m} + \underbrace{n(\bar{y} - m)^2}_{\text{Need to minimize with } m}$$

$\therefore L(m)$ is minimized at $m = \bar{y}$

$\therefore \boxed{m = \bar{y}}$ is the unique minimizer

(b) Minimize $L(m) = \max_i |y_i - m|$

Let $y_{(1)} = \min_i (y_i)$ & $y_{(n)} = \max_i y_i$

claim: A minimizer is $m^* = \frac{y_{(1)} + y_{(n)}}{2} \rightarrow$ midrange

If $m < y_{(1)}$ \rightarrow the worst error is $|y_{(n)} - m| = y_{(n)} - m$
which strictly decreases as you move m right toward $y_{(1)}$

Similarly for $m > y_{(n)}$ moving left. So any minimizer lies in $(y_{(1)}, y_{(n)})$

For any $m \in [y_{(1)}, y_{(n)}]$ the worst error comes from the extremes.

$$\therefore L(m) = \max \{m - y_{(1)}, y_{(n)} - m\}.$$

The maximum is smallest when the two terms are equal.

$$m - y_{(1)} = y_{(n)} - m \Rightarrow m = \frac{y_{(1)} + y_{(n)}}{2}$$

$$\textcircled{c} \quad L(m) = \sum_{i=1}^n |y_i - m| \rightarrow \text{median}$$

Let the data be sorted $\rightarrow y_{(1)} \leq y_{(2)} \leq \dots \leq y_{(n)}$

Let consider pairs around the center.

for $i = 1, 2, \dots, \lfloor \frac{n}{2} \rfloor$, pair $(y_{(i)}, y_{(n-i+1)})$

$$\text{Let } g_i(m) = |y_{(i)} - m| + |y_{(n-i+1)} - m|$$

$$\text{If } m \leq y_{(i)} : g_i(m) = (y_{(i)} - m) + (y_{(n-i+1)} - m) = y_{(i)} + y_{(n-i+1)} - 2m$$

$g_i(m)$ decreases as m increases

$$\text{If } m \geq y_{(n-i+1)} : g_i(m) = (m - y_{(i)}) + (m - y_{(n-i+1)}) = 2m - (y_{(i)} + y_{(n-i+1)})$$

$g_i(m)$ increases as m increases

$$\text{If } y_{(i)} \leq m \leq y_{(n-i+1)} : g_i(m) = (m - y_{(i)}) + (y_{(n-i+1)} - m) = y_{(n-i+1)} - y_{(i)}$$

$g_i(m) \rightarrow \text{constant}$ ($\&$ it's minimum)

$L(m) = \sum_i g_i(m)$ is minimized when m lies in the intersection of all these minimizing intervals.

$$\bigcap_{i=1}^{\lfloor n/2 \rfloor} [y_{(i)}, y_{(n-i+1)}]$$

If n is odd, the intersection collapses to the single point $y_{(\lfloor n/2 \rfloor)}$ \rightarrow the middle value (median)

\Rightarrow Unique minimizer: $m^* = y_{(\lfloor n/2 \rfloor)}$

If n is even, the intersection is the interval $(y_{(\lfloor n/2 \rfloor)}, y_{(\lceil n/2 \rceil)})$

\Rightarrow Every m^* in this interval minimizes $L(m)$

\therefore Any median minimizer $\sum_i |y_i - m|$

Hence proven

Problem-3

(a) $f(x) = \begin{cases} a_1 + s_1 x_i : x_i < \lambda \\ a_2 + s_2 x_i : x_i \geq \lambda \end{cases}$ $a_1 + s_1 \lambda = a_2 + s_2 \lambda$

$$a_2 + s_2 \lambda = a_1 + s_1 \lambda \Rightarrow a_2 = a_1 + (s_1 - s_2) \lambda$$

for $x_i \geq \lambda$

$$f(x) = a_2 + s_2 x_i = a_1 + (s_1 - s_2) \lambda + s_2 x_i$$

$$= \underbrace{a_1 + s_1 \lambda - s_2 \lambda}_{\text{new intercept}} + s_2 x_i$$

Thus the model is equivalent to

$$f(x) = \begin{cases} a_1 + s_1 x_i & : x_i < \lambda \\ a_1 + s_1 \lambda - s_2 \lambda + s_2 x_i & : x_i \geq \lambda \end{cases}$$

$$\textcircled{b} \quad h_\lambda(x) = (x - \lambda)_+ = \max\{0, x - \lambda\}$$

$$f(x) = \beta_0 + \beta_1 x + \beta_2 h_\lambda(x).$$

$$\text{If } x < \lambda \rightarrow h_\lambda(x) = 0$$

$$f(x) = \beta_0 + \beta_1 x \Rightarrow a_1 = \beta_0 \quad s_1 = \beta_1$$

$$\text{If } x \geq \lambda \rightarrow h_\lambda(x) = x - \lambda$$

$$f(x) = \beta_0 + \beta_1 x + \beta_2 (x - \lambda) = (\beta_0 + \beta_2 \lambda) + (\beta_1 + \beta_2) x$$

$$f(x) = (\beta_0 - \beta_2 \lambda) + (\beta_1 + \beta_2) x$$

$$\Rightarrow s_2 = \beta_1 + \beta_2 \quad ; \quad a_1 = \beta_0$$

$$\text{Let } X \in \mathbb{R}^{n \times 3} \Rightarrow x_i = [1, x_i, h_\lambda(x_i)]$$

{ target vector $y = (y_1, \dots, y_n)^T$

$$\hat{\beta} = \arg \min_{\beta} \|y - X\beta\|_2^2$$

$$\hat{\beta} = (\hat{\beta}_0, \hat{\beta}_1, \hat{\beta}_2) \Rightarrow \boxed{\hat{a}_1 = \hat{\beta}_0} \quad \boxed{\hat{s}_1 = \hat{\beta}_1} \quad \boxed{\hat{s}_2 = \hat{\beta}_1 + \hat{\beta}_2}$$

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from pathlib import Path

csv_path = 'https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/auto-mpg.data'
df = pd.read_csv(
    csv_path, delim_whitespace=True,
    names=['mpg', 'cylinders', 'displacement',
           'horsepower', 'weight', 'acceleration',
           'model year', 'origin', 'car name'])

print(df.shape)
df.head()

```

/tmp/ipython-input-3835944077.py:7: FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and will be removed in a future version. Use ``sep='\\s+'`` instead.

```

df = pd.read_csv(csv_path, delim_whitespace=True, names=['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', 'acceleration', 'model year', 'origin', 'car name'])
(398, 9)

```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	ford torino

```

df['horsepower'] = pd.to_numeric(df['horsepower'], errors='coerce')

df = df[['mpg', 'horsepower']].dropna().copy()

df = df[df['horsepower'].apply(lambda v: np.isfinite(v))]
df = df[df['mpg'].apply(lambda v: np.isfinite(v))]

y = df['mpg'].to_numpy().astype(float)
x = df['horsepower'].to_numpy().astype(float)

lam = 100.0
h = np.maximum(0.0, x - lam)
X = np.column_stack([np.ones_like(x), x, h])

beta_hat, residuals, rank, svals = np.linalg.lstsq(X, y, rcond=None)
b0, b1, b2 = beta_hat

a1 = b0
s1 = b1
s2 = b1 + b2

def f_hat(x_val):
    return b0 + b1 * x_val + b2 * np.maximum(0.0, x_val - lam)

```



```

x_left = np.linspace(x.min(), min(lam, x.max()), 200)
x_right = np.linspace(max(lam, x.min()), x.max(), 200)

y_left = a1 + s1 * x_left
y_right = (a1 + s1*lam - s2*lam) + s2 * x_right

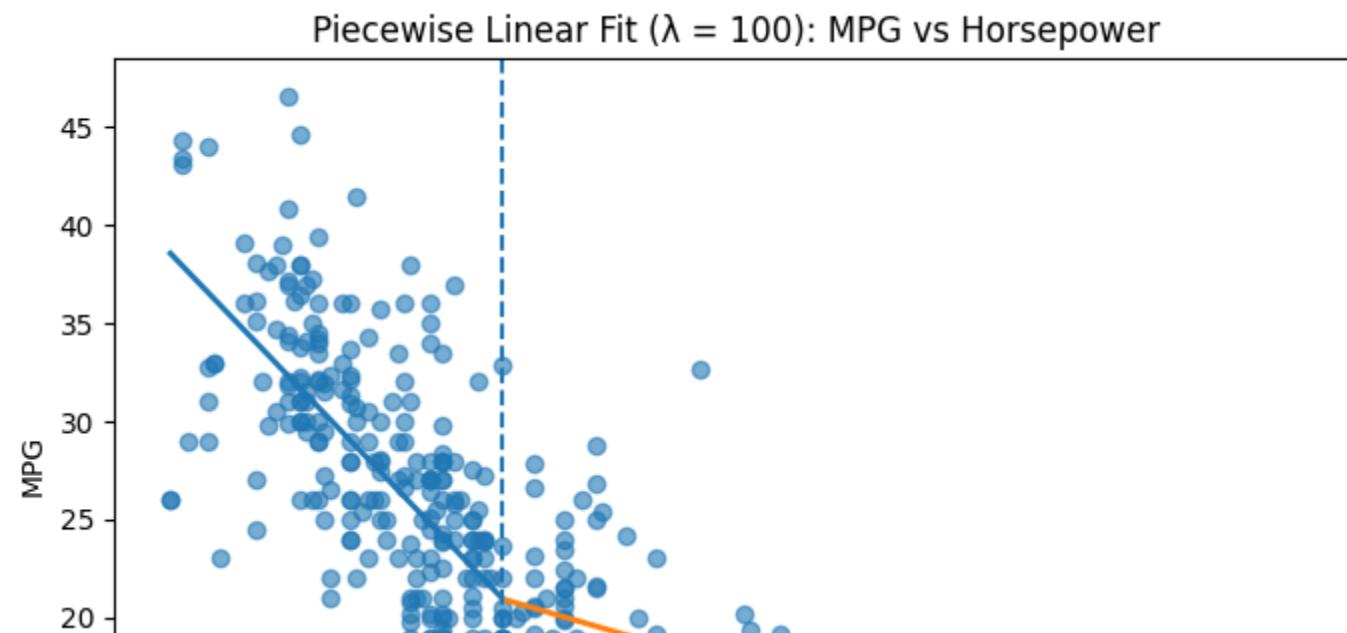
print(7*\n')
plt.figure(figsize=(7, 5))
plt.scatter(x, y, alpha=0.6)
plt.plot(x_left, y_left, linewidth=2)
plt.plot(x_right, y_right, linewidth=2)
plt.axvline(lam, linestyle='--')

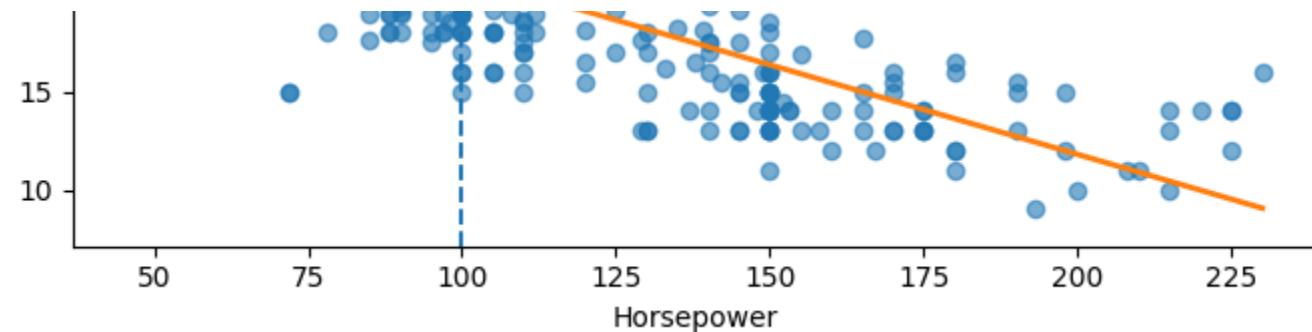
plt.xlabel('Horsepower')
plt.ylabel('MPG')
plt.title('Piecewise Linear Fit ( $\lambda = 100$ ): MPG vs Horsepower')
plt.tight_layout()
plt.show()

print("Estimated parameters:")
print(f"a1 = {a1:.4f}")
print(f"s1 = {s1:.4f}")
print(f"s2 = {s2:.4f}")

print("\nPiecewise model:")
print(f"f(x) = a1 + s1*x,           for x < {lam}")
print(f"f(x) = a1 + s1*{lam} - s2*{lam} + s2*x,  for x >= {lam}")

```





Estimated parameters:

$$a_1 = 53.5772$$

$$s_1 = -0.3264$$

$$s_2 = -0.0914$$

Piecewise model:

$$f(x) = a_1 + s_1*x, \quad \text{for } x < 100.0$$

$$f(x) = a_1 + s_1*100.0 - s_2*100.0 + s_2*x, \quad \text{for } x \geq 100.0$$

Start coding or [generate with AI](#).

problem - 4

(a) $X = \begin{pmatrix} 1 & x_{1,1} & x_{1,2} & \dots & x_{1,d} \\ 1 & x_{2,1} & x_{2,2} & \dots & x_{2,d} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n,1} & x_{n,2} & \dots & x_{n,d} \end{pmatrix}$

$$\Rightarrow X = [1 \ Z] \quad 1 \in \mathbb{R}^n, Z \in \mathbb{R}^{n \times d}$$

The linear predictor $\hat{y} = X\beta = \beta_0 1 + Z\beta$

$$\beta = (\beta_0, \beta) \in \mathbb{R}^{1+d}$$

Let column means be $\mu = \frac{1}{n} Z^T 1 \in \mathbb{R}^d$

Let $Z_c = Z - 1\mu^T$ (center each feature)

$$\beta_0 1 + Z\beta = \beta_0 1 + (Z_c + 1\mu^T)\beta = (\beta_0 + \mu^T \beta) 1 + Z_c \beta$$

$$\text{For centered features } \beta'_0 = \beta_0 + \mu^T \beta \quad \beta' = \beta$$

Conversely, (β'_0, β') on Z_c , choose $\beta_0 = \beta'_0 - \mu^T \beta'$, $\beta = \beta'$ to get same \hat{y} on Z . \therefore The hypothesis classes coincide

\therefore The prediction sets achieve the same minimum training loss before and after centering.

(b) Let $\sigma = (\sigma_1, \dots, \sigma_d) \rightarrow$ column standard deviations

$$S = \text{diag}(\sigma_1, \dots, \sigma_d)$$

$$Z_s = ZS^\dagger$$

$$\text{For any } (\beta_0, \beta) \quad \beta_0 I + Z\beta = \beta_0 I + Z_s(S\beta)$$

So for scaled features

$$\beta'_0 = \beta_0 \quad \beta' = S\beta \quad \text{to get some } \hat{y}.$$

The map is invertible, so the prediction sets coincide.

The minimum training loss is unchanged & likewise the test loss.

② What if ℓ_1 or ℓ_∞ loss instead of ℓ_2 ?

The above arguments did not rely on the squared loss.
Only on the fact that the feature transform are
invertible & that the intercept column is present

Because the set of achievable predictions $\{\hat{y}\}$ is
identical before and after an invertible affine
reparametrization, the minimum possible value of
any loss that depends only on residuals

$$r = y - \hat{y} \text{ is the same.}$$

With ℓ_1 or ℓ_∞ , centering/normalizing still cannot
improve the optimal loss - they only reparametrize
the same prediction functions

\therefore The optimal loss is invariant