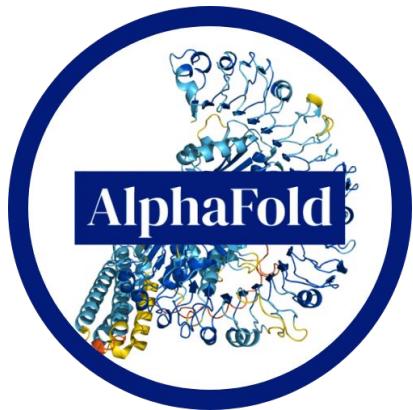


# NYU CS-GY 6923

# Machine Learning

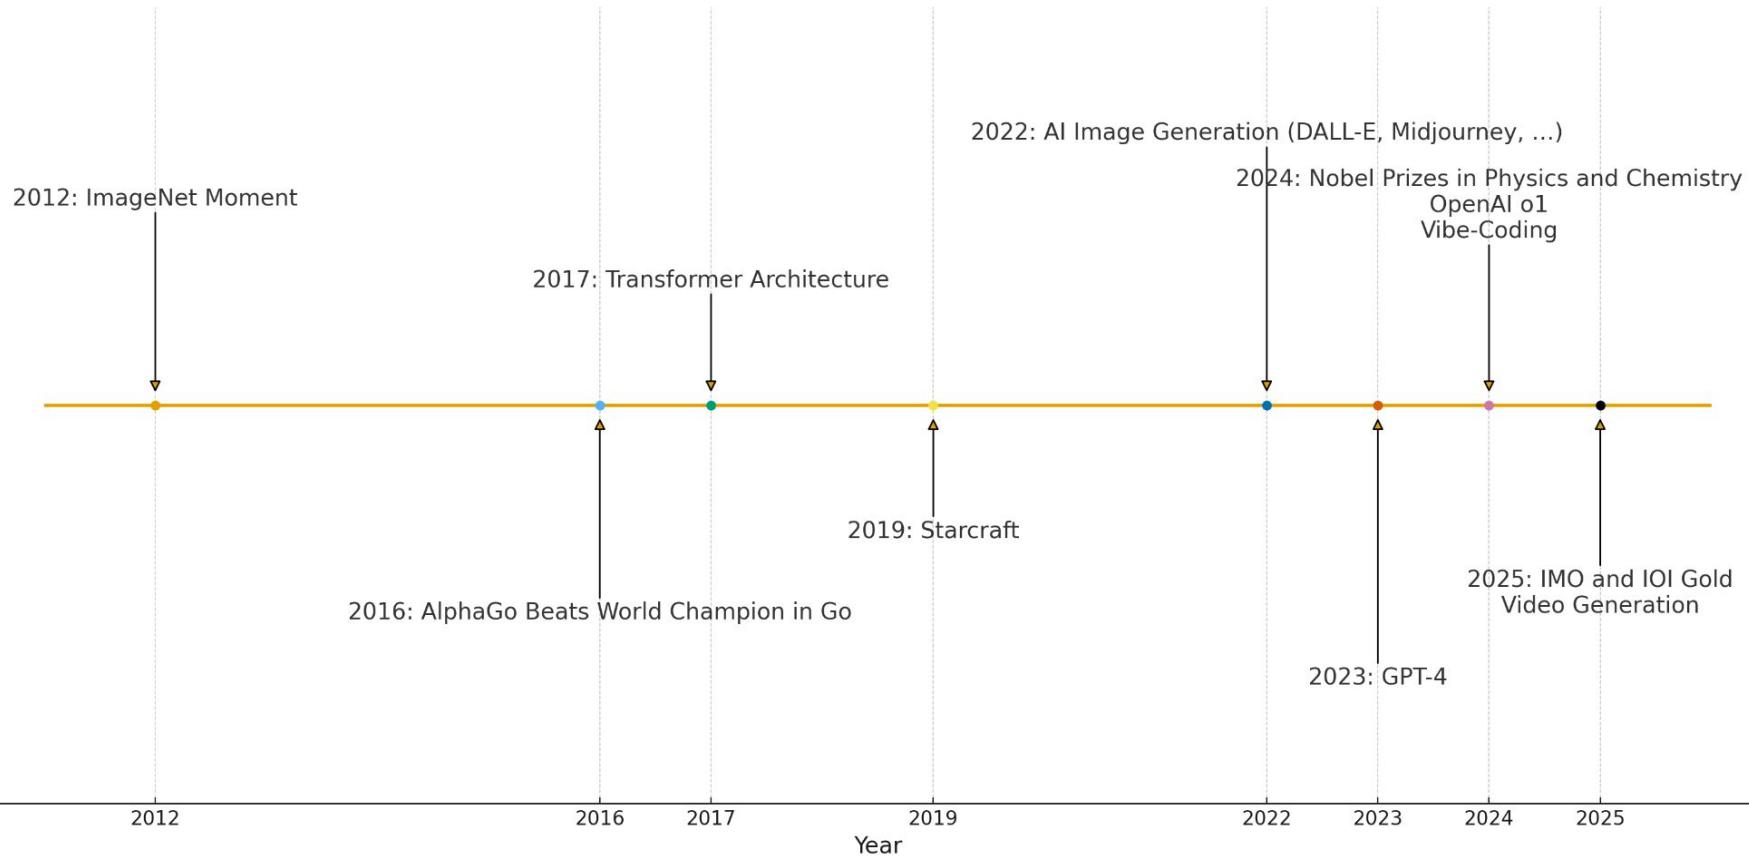
Prof. Pavel Izmailov



# Highlights

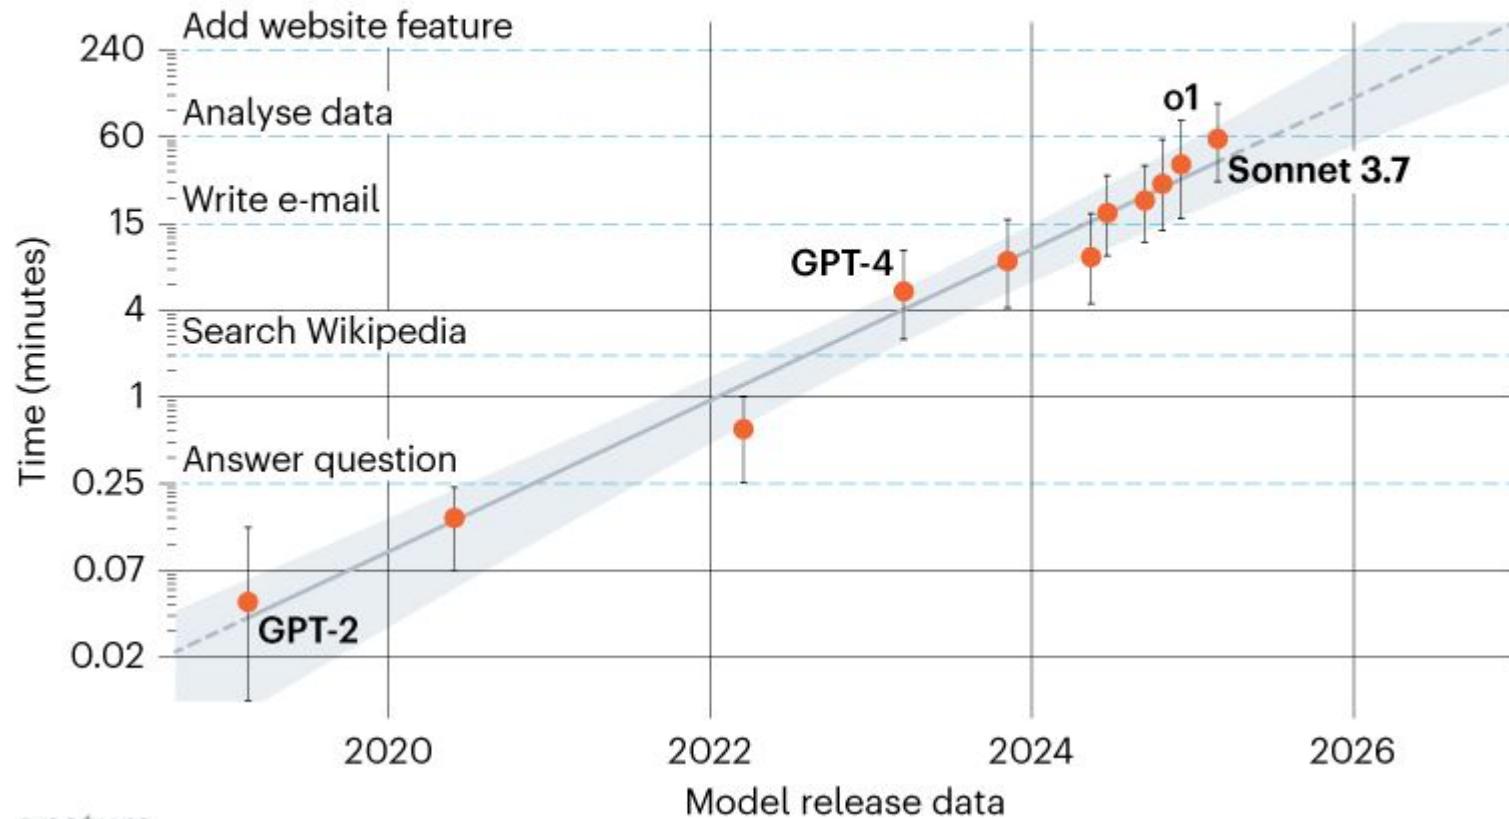
- 2012: ImageNet Moment
- 2016: AlphaGo Beats World Champion in Go
- 2017: Transformer Architecture
- 2019: Starcraft
- 2022: AI Image Generation (DALL-E, Midjourney, ...)
- 2023: GPT-4
- 2024:
  - Nobel Prizes in Physics and Chemistry
  - OpenAI o1
  - Vibe-Coding
- 2025 IMO and IOI Gold

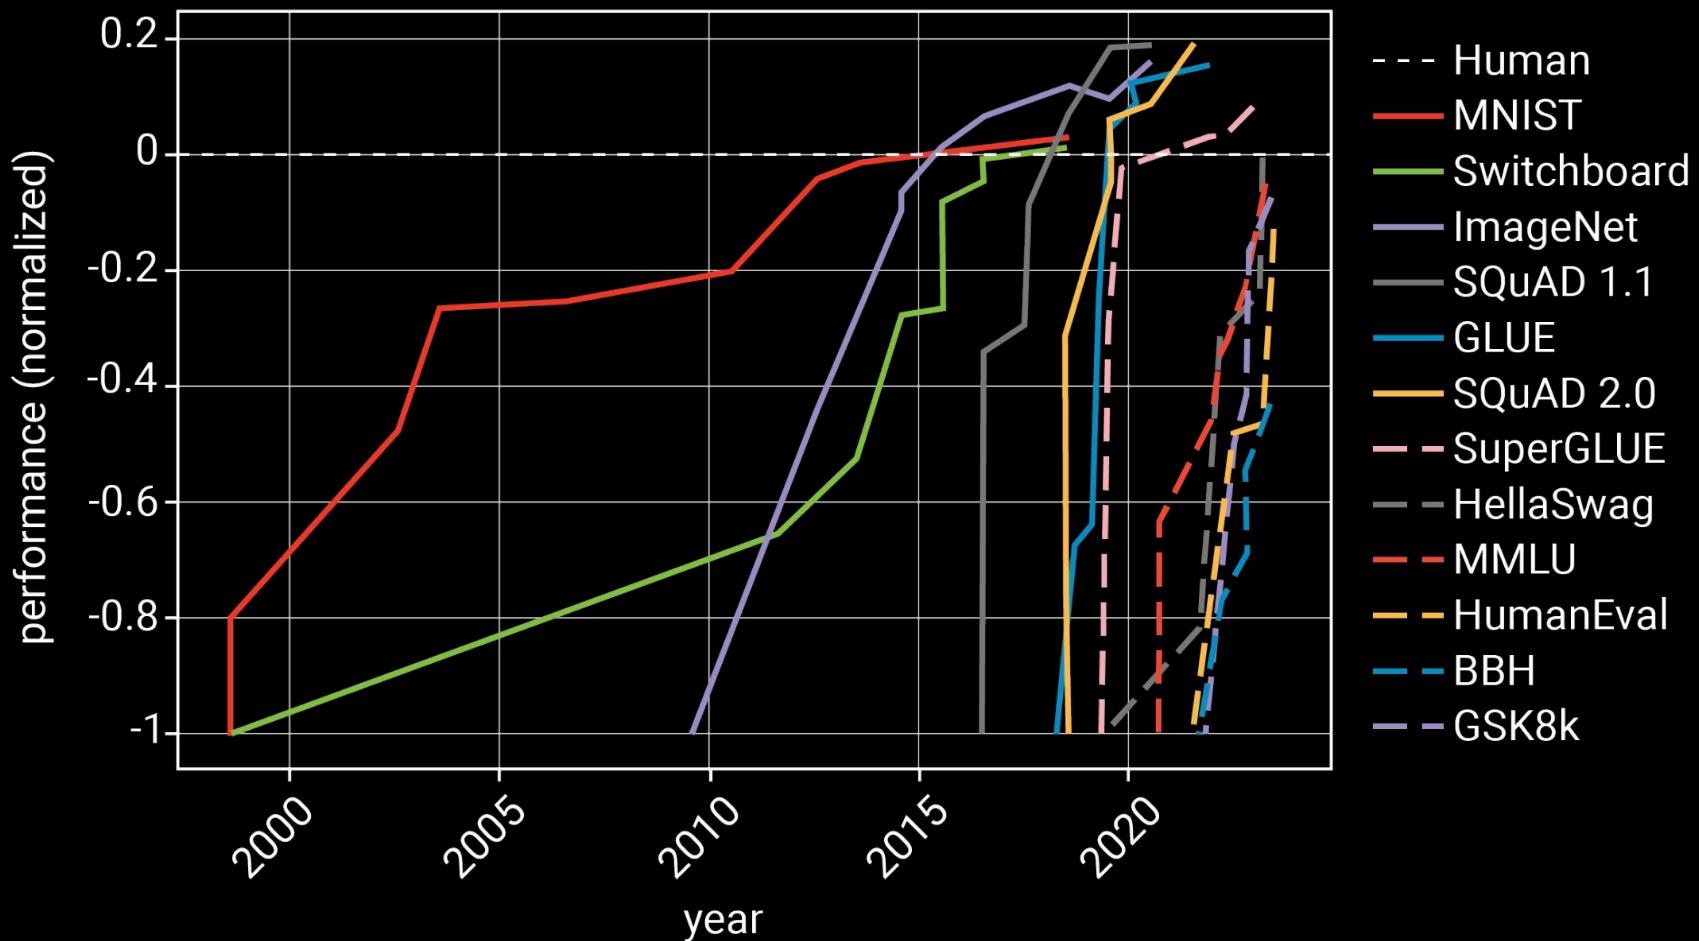
## AI Milestones Timeline (2012–2025)



# IMPROVING PERFORMANCE

Since 2019, the time humans take to complete tasks that AI models can complete at 50% accuracy has doubled every seven months.





[Kiela et al. 2023]

# Goals of this class

- Give you a foundation to understand the main ideas in modern ML
- By the end of the class, you should be able to:
  - Read NeurIPS papers
  - Implement and apply most of the core ML methods
  - Understand basic theoretical analysis and math behind these methods

# Grading

- Coding
  - Take-home labs: 20%
- Theory
  - Problem sets: 20%
  - Midterm exam: 25%
  - Final exam: 25%
- Participation: 10%
  - Ask questions, attend office hours, post on Ed.

# Team



**Majid Daliri**

*Course Assistant*



**Harsha Mupparaju**

*Course Assistant*



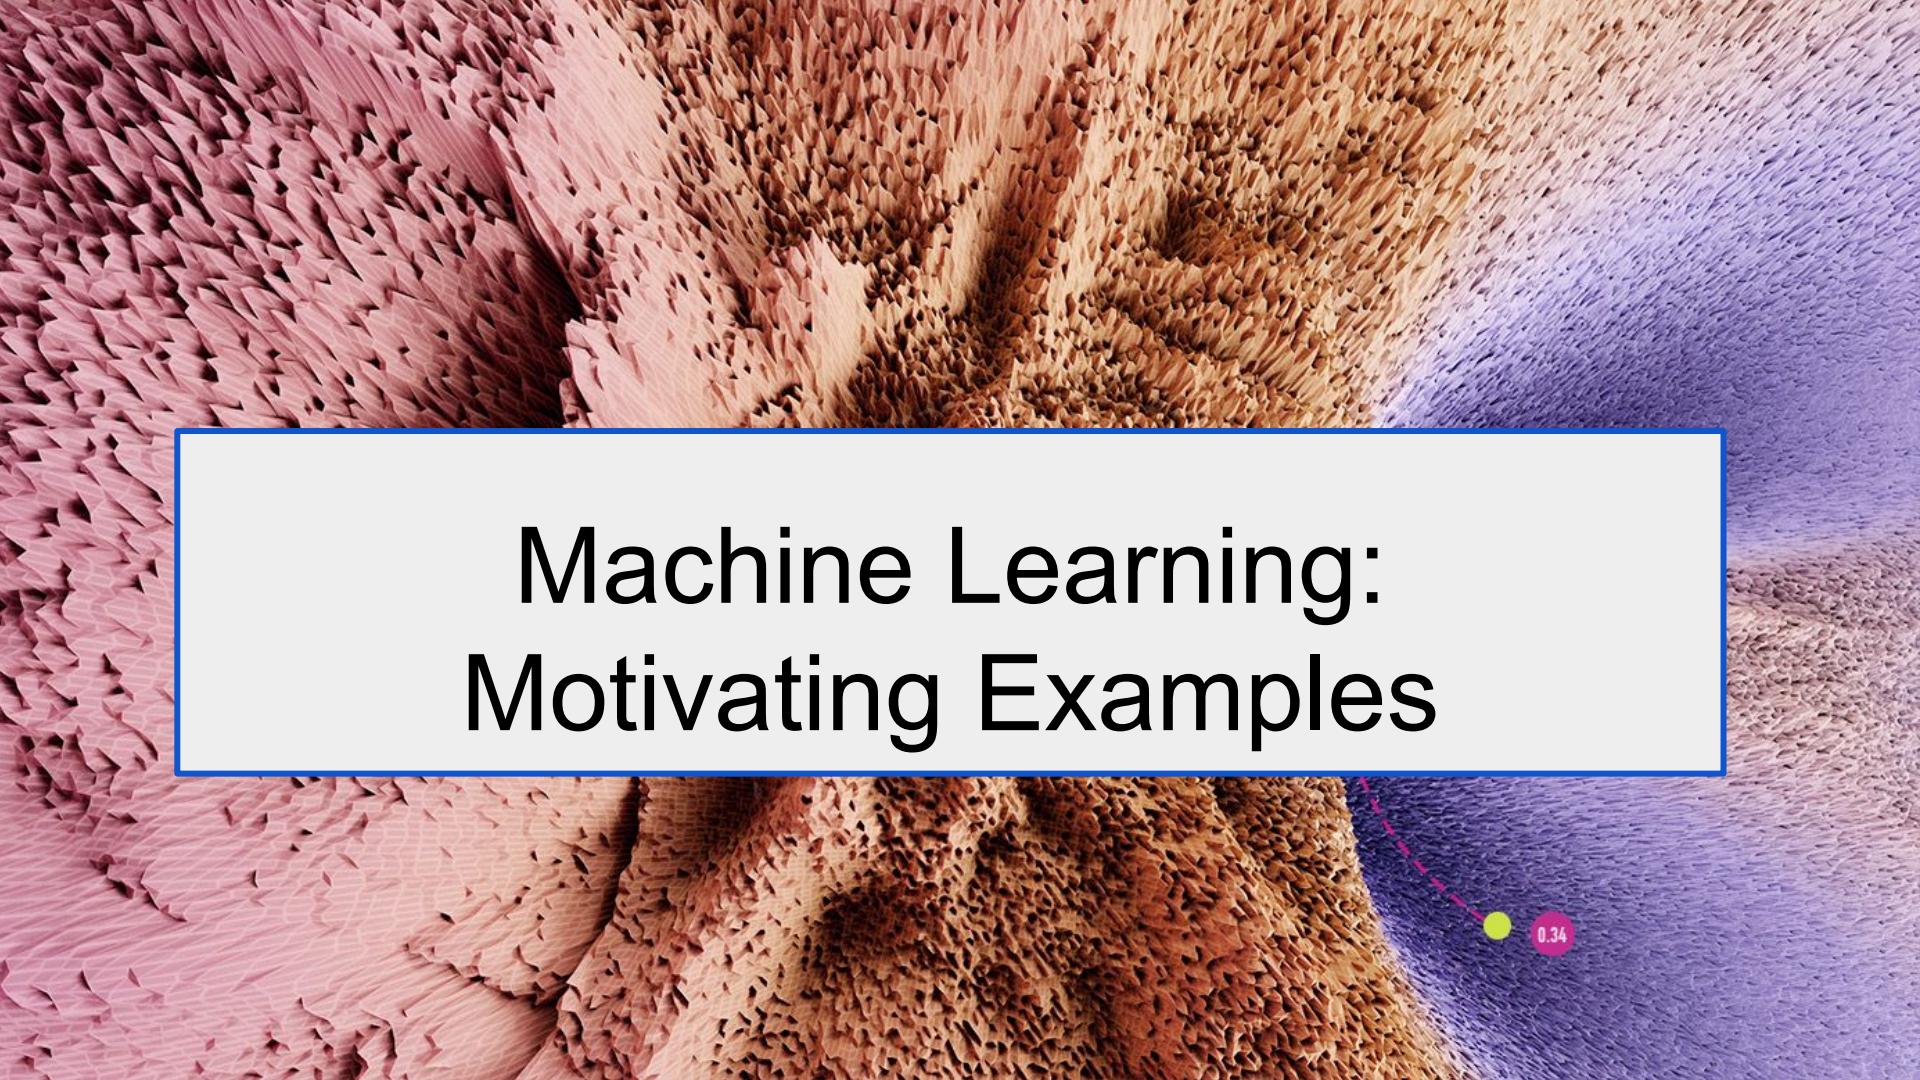
**Shubham Rastogi**

*Course Assistant*

- Up to date info: [izmailovpavel.github.io/ny\\_fall25\\_ml\\_grad](https://izmailovpavel.github.io/ny_fall25_ml_grad)
- Follow Ed discussion, go to office hours

# Policy

- Exams
  - Midterm: 10/17
  - Final: 12/19
- Collaboration
  - Fine to discuss problem sets and assignments, but **must write solutions independently**
  - Zero-tolerance policy for copied solutions
- AI tools allowed, make sure they don't impede your learning



# Machine Learning: Motivating Examples

0.34

# MNIST



Goal: recognize handwritten digits

- Images: 28x28 pixels, black and white
- Labels: 0, ..., 9

How can a computer classify them?

- Idea: hand-craft rules, e.g. if contains two circles then 8, ...

# Rule-based Systems

- Dominant in AI in 1970s and 80s
- Hard to implement heuristics
  - How to define a circle? A line?
- A lot of manual work
- Unreliable: Impossible to cover all corner cases
- Not scalable
  - What is your heuristic for telling a cheetah from a lion?



# ELIZA

He says I am depressed much of the time.

I AM SORRY TO HEAR YOU ARE DEPRESSED

INPUT

It's true. I am unhappy.

DO YOU THINK COMING HERE WILL HELP YOU NOT TO BE UNHAPPY

INPUT

I need some help, that much seems certain.

WHAT WOULD IT MEAN TO YOU IF YOU GOT SOME HELP

INPUT

Perhaps I could learn to get along with my mother.

TELL ME MORE ABOUT YOUR FAMILY

INPUT

My mother takes care of me.

WHO ELSE IN YOUR FAMILY TAKES CARE OF YOU

INPUT

My father.

YOUR FATHER

INPUT

# Rule-based Systems

ELIZA:

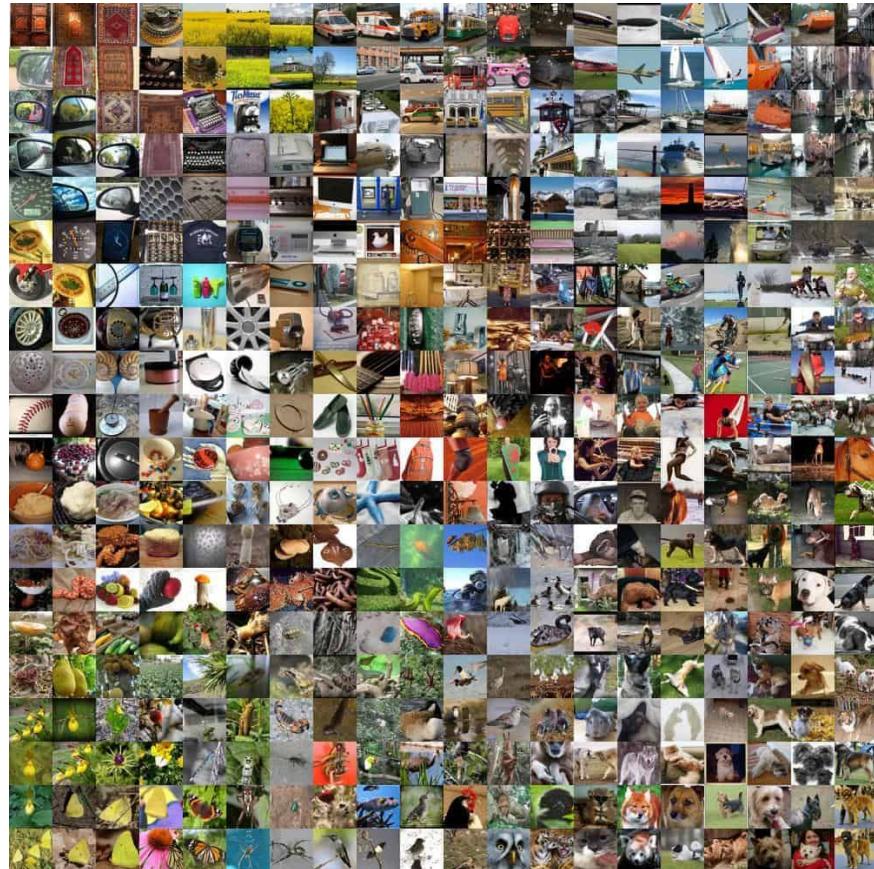
- Try here: <https://www.masswerk.at/elizabot/>
- Built in 1966
- Uses hand-crafted rules
- Impressive for the time, but by modern standards very bad

Rule-Based Systems

- Used in Siri, Alexa, ...

# Rule-based Systems

- Dominant in AI in 1970s and 80s
- Hard to implement heuristics
  -
- A lot of manual work
- Unreliable: Impossible to cover all corner cases
- Not scalable



# Machine Learning: sketch



Use data to train computers to make predictions.

- Construct a **training dataset**
- Define a **model**: mathematical function that maps from inputs (28x28 black and white pixels) to outputs (labels 0,...,9)
  - The model depends on **parameters**
- **Train** the model: set the parameters so that predictions on training data are good
- Use the model to make predictions on new data

*Whiteboard: logistic regression example*

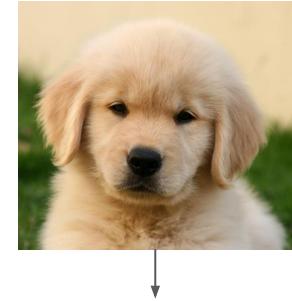
# Machine Learning: challenges



- Data: collect good training data
- Model: decide on the model class
- Training: find good parameters
- Generalization: make sure good train accuracy translates into good test predictions

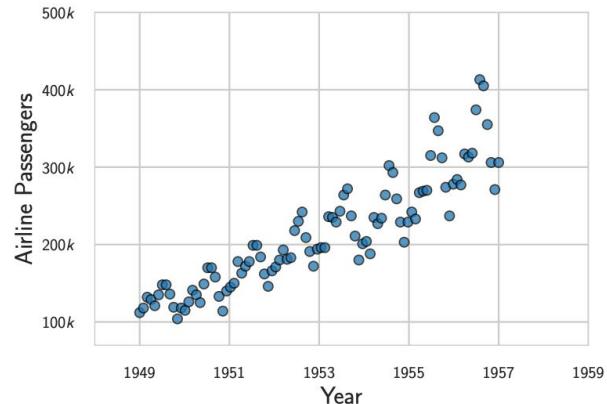
# Supervised learning

Classification: predict a discrete label



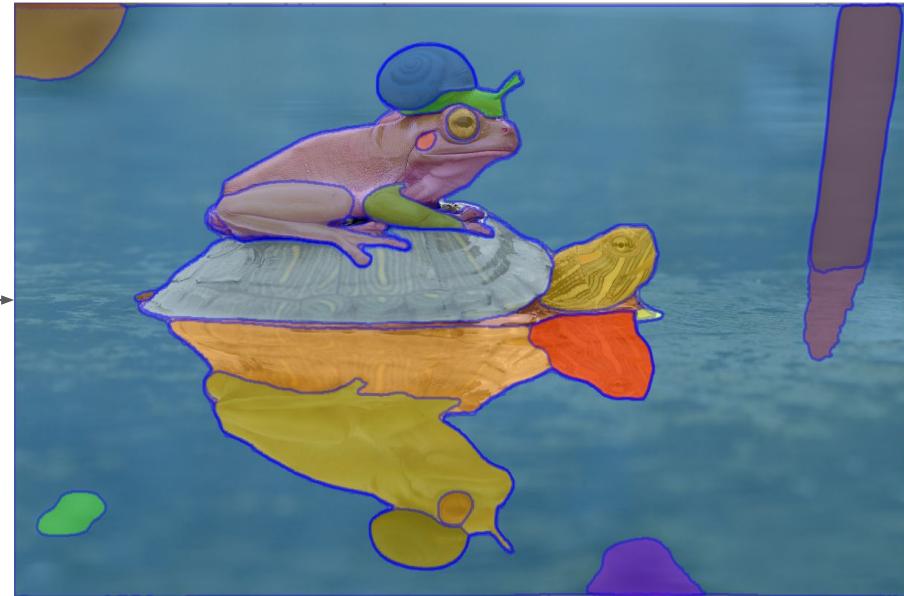
$y = Cat$

Regression: predict a continuous value



# Supervised learning

Segmentation: identify locations of objects



<https://segment-anything.com/demo>

# Other types of learning

- Unsupervised learning: no labels
- Self-supervised learning: also no labels
- Reinforcement learning: learn to act from interaction with environment
- ...

What is ChatGPT?

# Other types of learning

- Unsupervised learning: no labels
- Self-supervised learning: also no labels
- Reinforcement learning: learn to act from interaction with environment
- ...

What is ChatGPT?

- Pre-training is supervised learning / classification
- Post-training is reinforcement learning (RLHF)

# Simple Linear Regression

0.34

# Simple linear regression

- Demo in: `demo_auto_mpg.ipynb`
- Predict highway miles per gallon of a car given engine characteristics

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	ford torino
5	15.0	8	429.0	198.0	4341.0	10.0	70	1	ford galaxie 500

↑  
*y*



*X*

# Simple linear regression

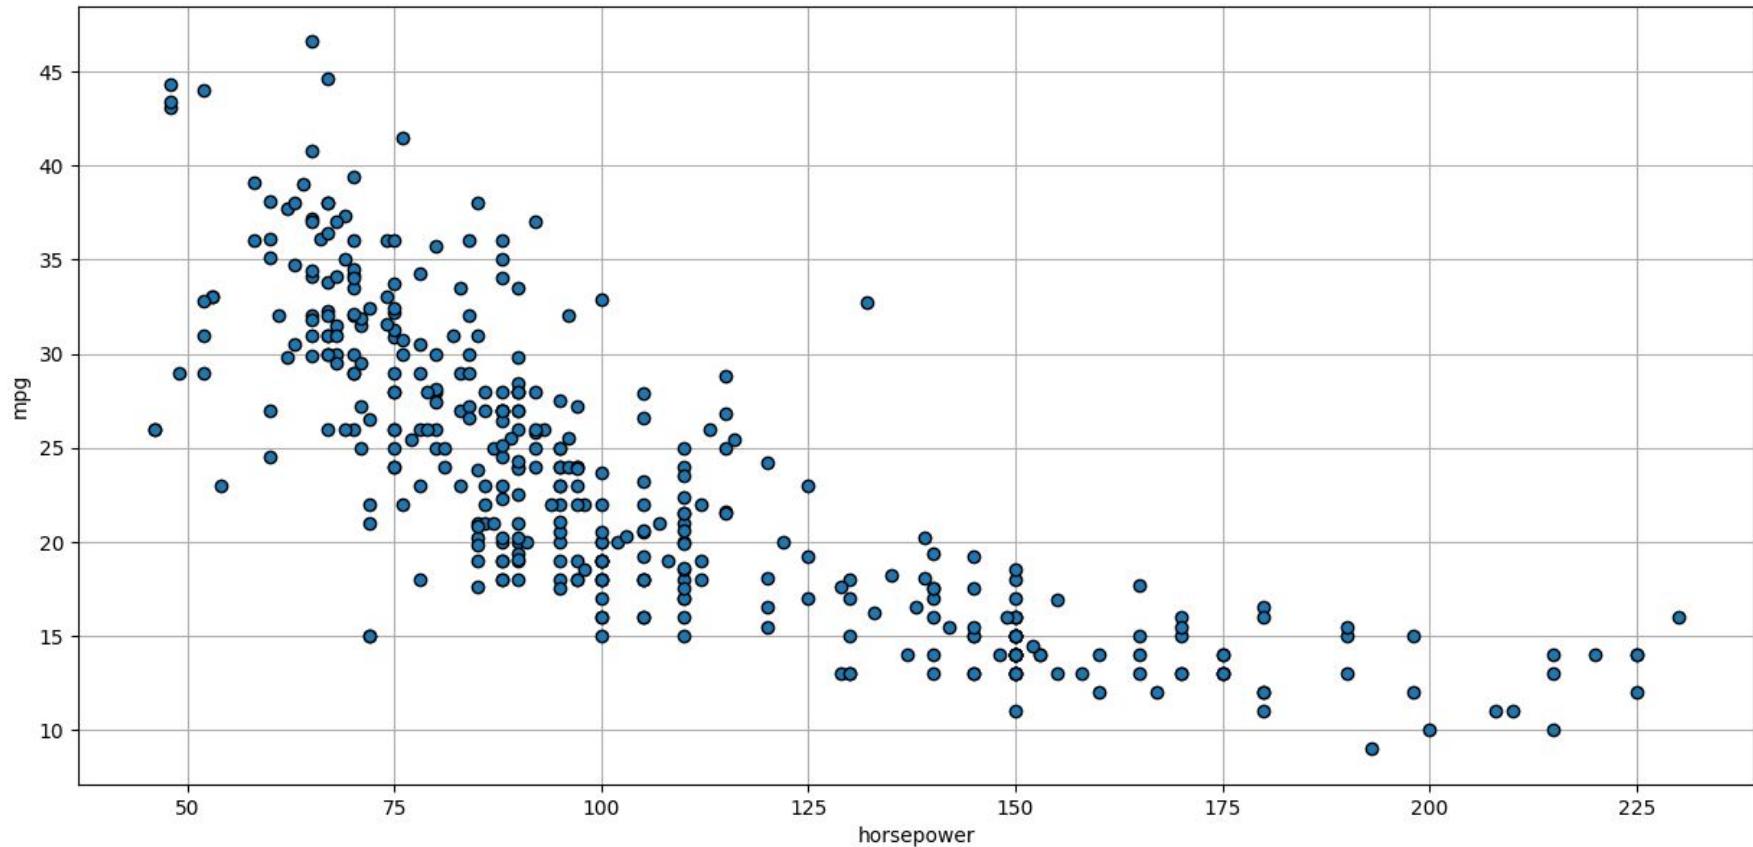
- Demo in: `pavel_demo_auto_mpg.ipynb`
- Predict highway miles per gallon of a car given engine characteristics

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin	car name
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	ford torino
5	15.0	8	429.0	198.0	3411.0	10.0	70	1	ford galaxie 500

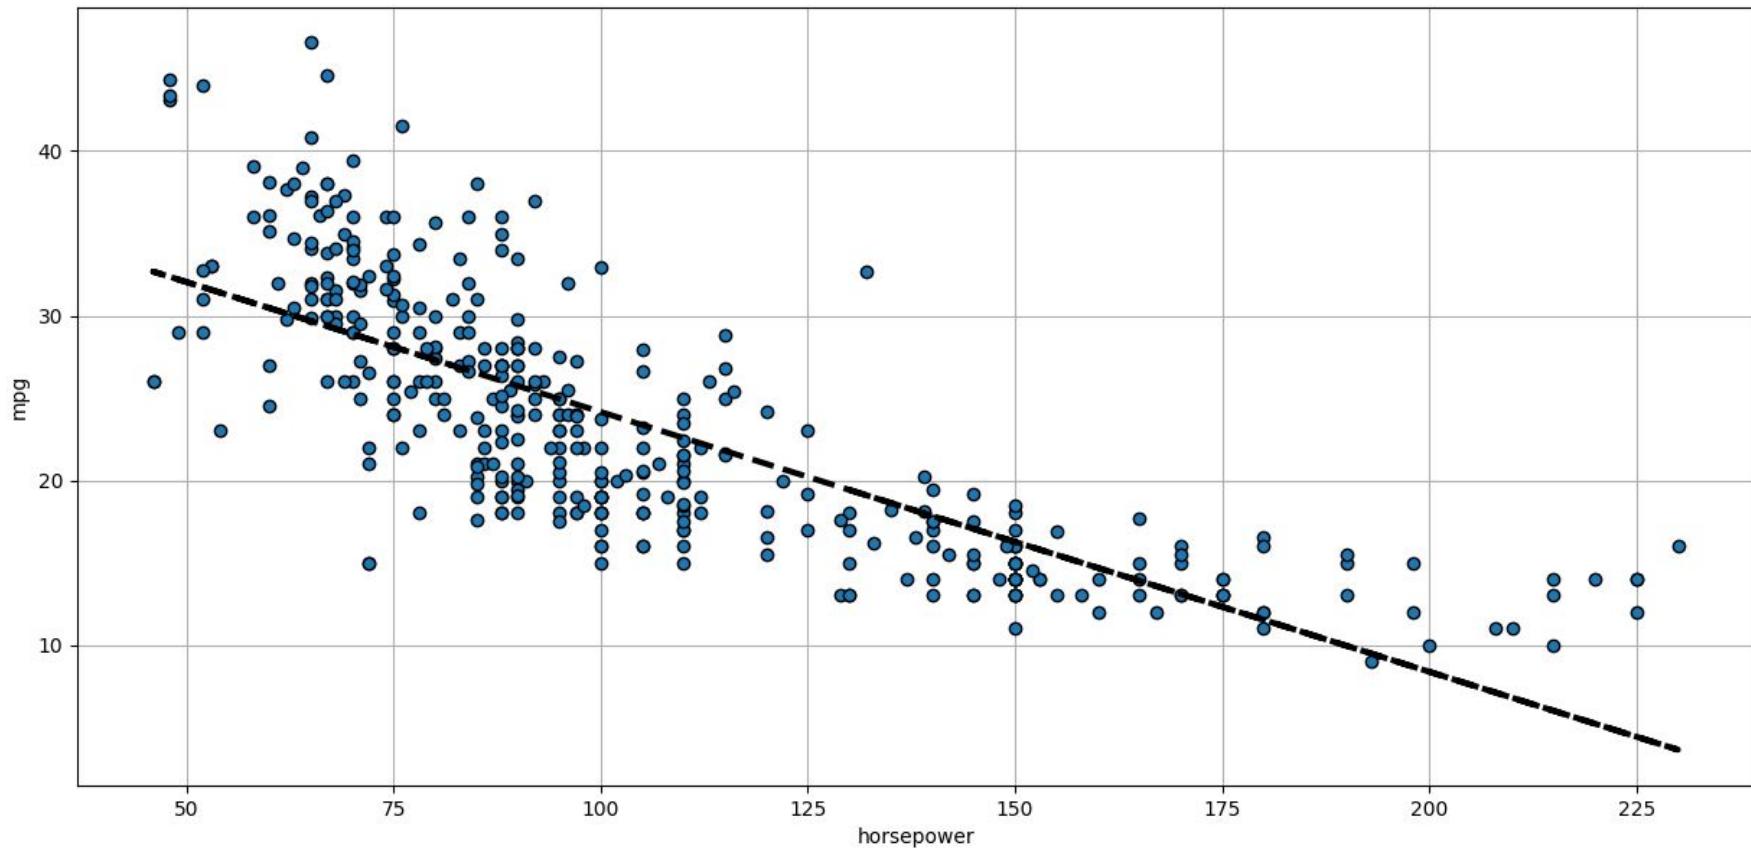
↑  
*y*

*x*

# Simple linear regression



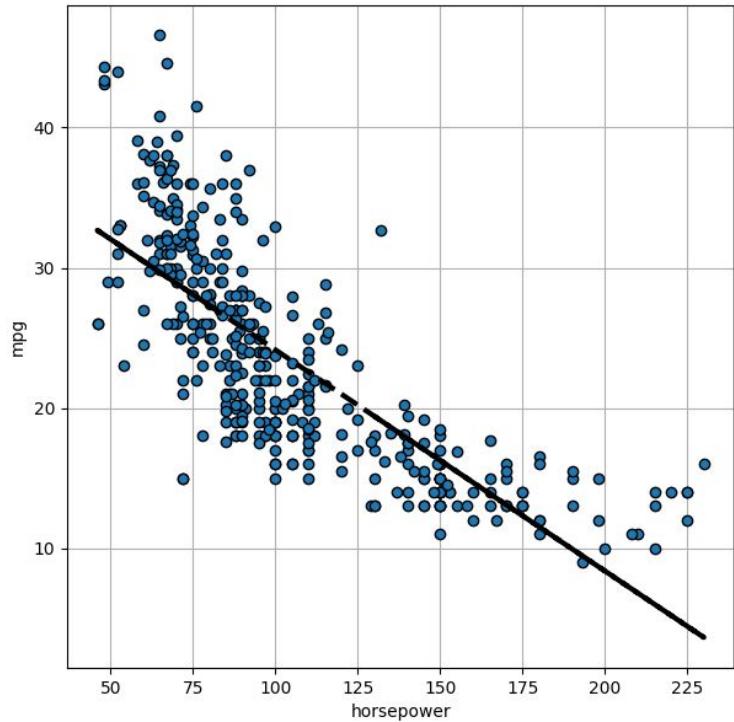
# Simple linear regression



# Simple linear regression

*Whiteboard:*

- *Dataset*
- *Model*
- *Model Parameters*
- *Loss Functions*
- *Training (ERM)*

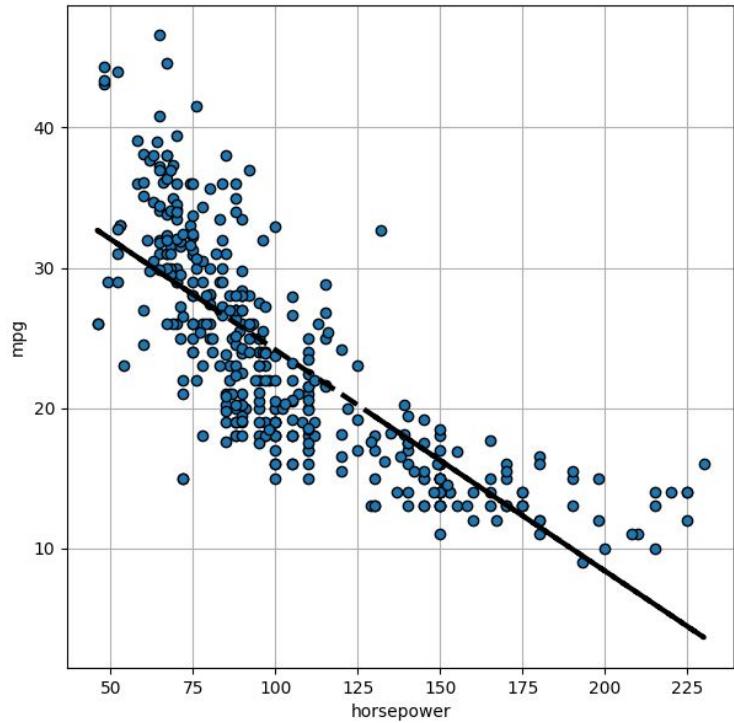


# Simple linear regression

$$L(w, b) = \sum_{i=1}^n (y_i - \hat{y}_i)^2 = \sum_{i=1}^n (y_i - wx_i - b)^2$$

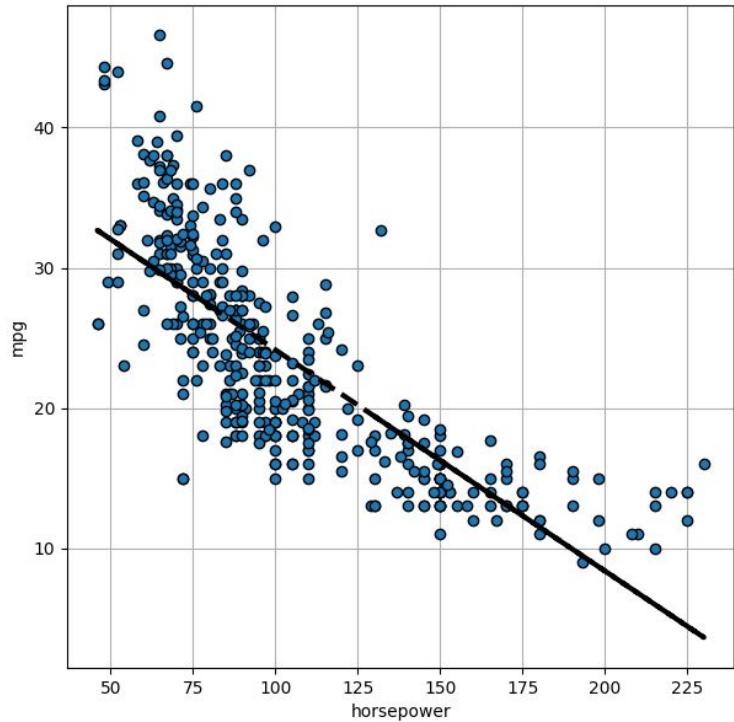
$$\hat{w} = \frac{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})}{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{b} = \bar{y} - \hat{w}\bar{x}$$



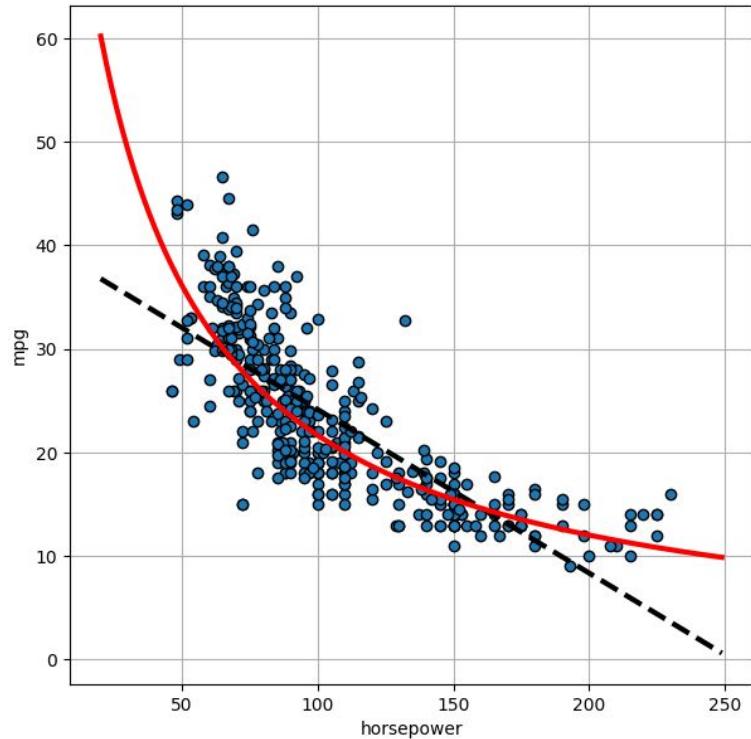
# Simple linear regression: comments

- In general we cannot always find optimal parameters analytically
- Data can be non-linear



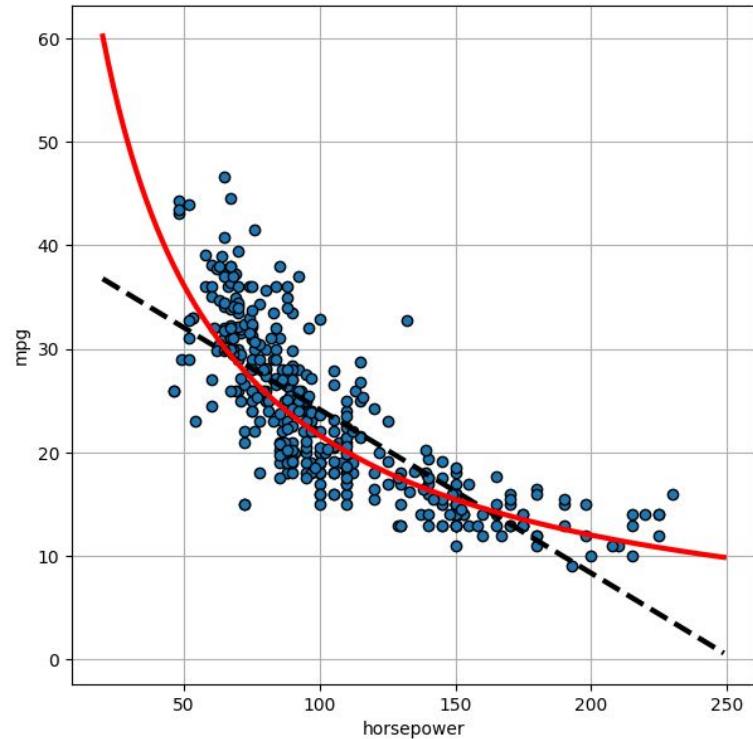
# Simple linear regression: comments

- In general we cannot always find optimal parameters analytically
- Data can be non-linear
  - Can preprocess x and y, e.g.  
 $y' = 1/y$



# Simple linear regression: comments

- In general we cannot always find optimal parameters analytically
- Data can be non-linear
  - Can preprocess x and y, e.g.  
 $y' = 1/y$
- You may need more than one feature
  - Think about MNIST



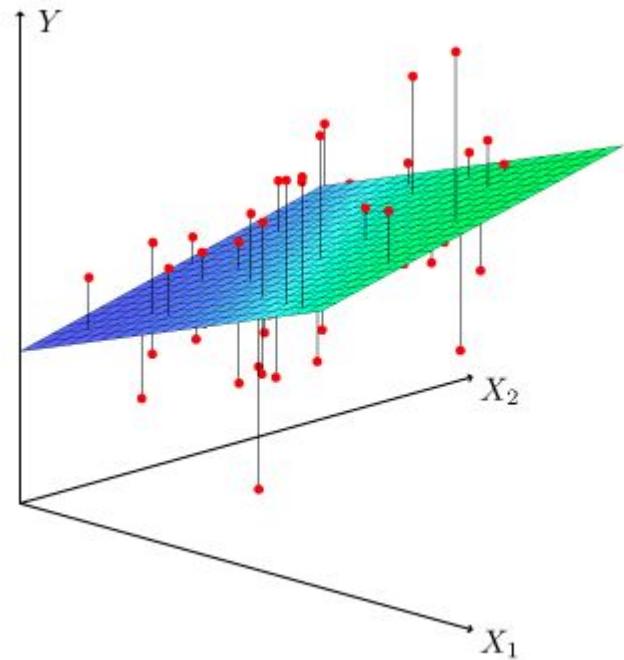
# Linear Regression With Multiple Features

0.34

# General linear regression

*Whiteboard:*

- *Dataset*
- *Model*
- *Model Parameters*
- *Loss Functions*
- *Training (ERM)*



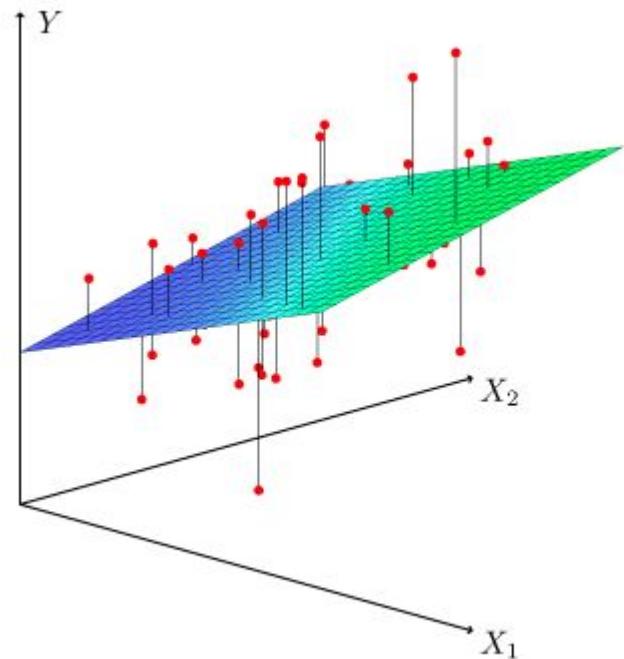
# General linear regression

$$\hat{w} = (X^T X)^{-1} (X^T Y)$$

Compare to the formulas we got for the simple linear regression:

$$\hat{w} = \frac{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})}{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{b} = \bar{y} - \hat{w}\bar{x}$$



# General linear regression

$$\hat{w} = (X^T X)^{-1} (X^T Y)$$

Compare to the formulas we got for the simple linear regression:

$$\hat{w} = \frac{\frac{1}{n} \sum_{i=1}^n (y_i - \bar{y})(x_i - \bar{x})}{\frac{1}{n} \sum_{i=1}^n (x_i - \bar{x})^2}$$

$$\hat{b} = \bar{y} - \hat{w}\bar{x}$$

# General linear regression

```
import pandas as pd
from sklearn.linear_model import LinearRegression

y_col = "mpg"
x_cols = ["cylinders", "displacement", "horsepower",
"weight", "acceleration", "model year"]
df_selected = df[[y_col] + x_cols]
df_selected = df_selected.dropna()

reg = LinearRegression()
reg.fit(X=df_selected[x_cols], y=df_selected[y_col])
print(reg.coef_)
```



# General linear regression

```
# Create DataFrame with Ferrari 308 GTB/GTS data
data = {
    'mpg': [17], # Combined estimate
    'cylinders': [8],
    'displacement': [177],
    'horsepower': [240],
    'weight': [3050], # pounds, mid-range estimate
    'acceleration': [7.0], # 0-60 mph in seconds
    'model year': [78] # Representative year
}

ferrari_df = pd.DataFrame(data)
reg.predict(ferrari_df[x_cols])
```



Predicts 22!

# General linear regression

*Whiteboard:*

- *Categorical features*
- *Transformations of features*