





# NYU CS-GY 6923

# Machine Learning

Prof. Pavel Izmailov

# Today

- Recap overfitting and underfitting
- Regularization
- Classification and logistic regression



# Overfitting, Underfitting and Generalization

0.34

# Overfitting and underfitting


- Data generated by

$$y_{\text{true}} = 1.2 \cdot x + 0.5 \cdot \sin(2\pi x) + 0.3 \cdot x^2$$

$$y = y_{\text{true}} + \varepsilon, \quad \varepsilon \sim N(0, 0.2^2)$$

- Polynomial features


$$[1, x, x^2, \dots, x^d]$$




# Overfitting and underfitting

## Underfitting vs Good Fit vs Overfitting


Underfitting (degree=1)  
Train MSE: 0.104, Test MSE: 0.058



Good Fit (degree=3)  
Train MSE: 0.025, Test MSE: 0.019




Overfitting (degree=15)  
Train MSE: 0.004, Test MSE: 0.605




# Validation

- We don't have access to the true test loss in practice
- Leave some of the train data as a held-out validation set
- Use validation loss as a proxy for generalization



# Model selection




# Connection to systems of linear equations

- We can view fitting a linear model as solving a system of linear equations

$$Xw = y \Leftrightarrow Ax = b$$

- What determines when we can solve a linear system?



# Connection to systems of linear equations


$$Xw = y \Leftrightarrow Ax = b$$

- If  $n > m$  and  $X$  full-rank then generally no solution
- Instead, we solve a least-squares problem

$$\min_w \|y - Xw\|^2$$

- Solution is given by the Moore–Penrose inverse

$$w^* = (X^T X)^{-1} X^T y$$



# Connection to systems of linear equations


$$Xw = y \Leftrightarrow Ax = b$$

- If  $n > m$  and  $X$  full-rank then generally no solution
- Instead, we solve a least-squares problem

$$\min_w \|y - Xw\|^2$$

- Solution is given by the Moore–Penrose inverse

$$w^* = (X^T X)^{-1} X^T y$$



$$[(X^T X)^{-1} X^T] \cdot X = (X^T X)^{-1} (X^T X) = I.$$

# Connection to systems of linear equations

$$Xw = y \Leftrightarrow Ax = b$$

- If  $n = m$  and  $X$  full-rank then unique solution
- We can fit the data perfectly in one unique way
- Pseudo-Inverse simplifies to

$$w^* = X^{-1}y$$


The diagram illustrates the matrix equation  $Xw = y$ . It features three vertical rectangles. The first rectangle on the left is blue and labeled with a large italicized 'X' in its center. Above it is the letter 'm' and below it is the letter 'n'. To its right is a small black multiplication sign. Next is a second rectangle colored light red, labeled with a large italicized 'w' in its center. Above it is the letter 'm' and below it is the letter 'n'. To its right is an equals sign '='. Finally, to the right of the equals sign is a third rectangle colored light green, labeled with a large italicized 'y' in its center. Above it is the letter 'n'.

# Connection to systems of linear equations

$$Xw = y \Leftrightarrow Ax = b$$

- If  $n < m$  and  $X$  full-rank then many solutions fit data perfectly
- Our formula no longer works...

$$w^* = (X^T X)^{-1} X^T y$$



- We can solve an optimization problem


$$\min_w \|w\|_2, \text{ s.t. } Xw = y$$

# Connection to systems of linear equations

$$Xw = y \Leftrightarrow Ax = b$$

- If  $n < m$  and  $X$  full-rank then many solutions fit data perfectly
- Our formula no longer works...

$$w^* = (X^T X)^{-1} X^T y$$



- We can solve an optimization problem




$$\min_w \|w\|_2, \text{ s.t. } Xw = y$$

- Turns out, the solution is still given by the pseudo-inverse (defined via SVD)!

$$w^* = X^+ y$$

# Connection to systems of linear equations

## Linear Systems Perspective on Under/Over-Parameterization



Linear System:  $Ax = b$

Equations (n): 8  
Unknowns (m): 2  
Rank of A: 2  
Condition #: 3.89e+00

Residual  $\|Ax-b\|$ : 0.7859

System Type:  
• OVERRDETERMINED ( $n > m$ )  
• No exact solution  
• Use least squares  
• High bias, low variance

Linear System:  $Ax = b$

Equations (n): 8  
Unknowns (m): 8  
Rank of A: 8  
Condition #: 2.68e+05

Residual  $\|Ax-b\|$ : 0.0000

System Type:  
• WELL-DETERMINED ( $n \approx m$ )  
• Unique solution  
• Balanced bias-variance

Linear System:  $Ax = b$

Equations (n): 8  
Unknowns (m): 11  
Rank of A: 8  
Condition #: 3.98e+04

Residual  $\|Ax-b\|$ : 0.0000

System Type:  
• UNDERDETERMINED ( $n < m$ )  
• Infinite solutions  
• Zero residual possible  
• Low bias, high variance


# Connection to systems of linear equations

- Equations  $\leftrightarrow$  datapoints, variables  $\leftrightarrow$  features
- If we have many equations and few features, we will often underfit
  - When will we not underfit here?
- When we have more ( $\geq$ ) features than data, we will fit the data perfectly
  - In most cases, this would mean overfitting
- Pseudo-inverse gives a general form of optimal solution  $w^* = X^+y$

# Regularization


# Regularization Motivation

- Go back to our regression problem




# Regularization Motivation

- Go back to our regression problem
- Let's fit a degree-15 polynomial



# Regularization Motivation


- Go back to our regression problem
- Let's fit a degree-15 polynomial
- We are not in the over-parameterized regime, but close to it
  - We are overfitting badly



```
9.9E-02 + 0.0E+00 x^1 + 5.6E+02 x^2 - 3.0E+04 x^3 + 6.7E+05 x^4 - 8.2E+06 x^5 + 6.4E+07 x^6 -  
3.4E+08 x^7 + 1.3E+09 x^8 - 3.4E+09 x^9 + 6.6E+09 x^10 - 9.6E+09 x^11 + 1.0E+10 x^12 - 7.4E+09 x^13
```

# Regularization Motivation

- Go back to our regression problem
- Let's fit a degree-15 polynomial
- We are not in the over-parameterized regime, but close to it
  - We are overfitting badly




The model has to use big weights to fit the data

$$9.9\text{E}-02 + 0.0\text{E}+00 \cdot x^1 + 5.6\text{E}+02 \cdot x^2 - 3.0\text{E}+04 \cdot x^3 + 6.7\text{E}+05 \cdot x^4 - 8.2\text{E}+06 \cdot x^5 + 6.4\text{E}+07 \cdot x^6 - 3.4\text{E}+08 \cdot x^7 + 1.3\text{E}+09 \cdot x^8 - 3.4\text{E}+09 \cdot x^9 + 6.6\text{E}+09 \cdot x^{10} - 9.6\text{E}+09 \cdot x^{11} + 1.0\text{E}+10 \cdot x^{12} - 7.4\text{E}+09 \cdot x^{13}$$

# Regularization Motivation

- We change features a bit and get a lot of change in the predictions
- This happens when the weights are large
- What can we do to prevent large weights?



The model has to use big weights to fit the data

$$9.9\text{E-}02 + 0.0\text{E+}00 x^1 + 5.6\text{E+}02 x^2 - 3.0\text{E+}04 x^3 + 6.7\text{E+}05 x^4 - 8.2\text{E+}06 x^5 + 6.4\text{E+}07 x^6 - 3.4\text{E+}08 x^7 + 1.3\text{E+}09 x^8 - 3.4\text{E+}09 x^9 + 6.6\text{E+}09 x^{10} - 9.6\text{E+}09 x^{11} + 1.0\text{E+}10 x^{12} - 7.4\text{E+}09 x^{13}$$

# Regularization Motivation

What can we do to prevent large weights?

$$\min_w \|Xw - y\|^2 + \lambda \|w\|^2$$

We add a *regularization* term that penalizes large weights!

- We trade off fitting the data vs the norm of the weights.
- We can choose different forms of regularization.

# Regularization

$$\min_w \|Xw - y\|^2 + \lambda \|w\|^2$$

# Regularization

$$\min_w \|Xw - y\|^2 + \lambda \|w\|^2$$

$$J(w) = (Xw - y)^T(Xw - y) + \lambda w^T w$$

$$\begin{aligned} (Xw - y)^T(Xw - y) &= (w^T X^T - y^T)(Xw - y) \\ &= w^T X^T X w - w^T X^T y - y^T X w + y^T y \\ &= w^T X^T X w - 2w^T X^T y + y^T y \end{aligned}$$

$$\begin{aligned} J(w) &= w^T (X^T X + \lambda I) w - 2w^T X^T y \\ &\quad + y^T y \end{aligned}$$

- $\frac{\partial}{\partial w}(w^T A w) = 2Aw$  for symmetric matrix  $A$
- $\frac{\partial}{\partial w}(a^T w) = a$  for vector  $a$

$$\frac{\partial J(w)}{\partial w} = 2(X^T X + \lambda I)w - 2X^T y$$

$$w^* = (X^T X + \lambda I)^{-1} X^T y$$

# Regularization

$$\min_w \|Xw - y\|^2 + \lambda \|w\|^2$$

$$w^* = (X^T X + \lambda I)^{-1} X^T y$$

- When  $\lambda$  is 0, recover our standard LR solution
- When  $\lambda > 0$ ,  $w^*$  always exists,  $(X^T X + \lambda I)$  is always invertible
- What happens when  $\lambda \rightarrow \infty$ ?


# Regularization

$$\min_w \|Xw - y\|^2 + \lambda \|w\|^2$$

$$w^* = (X^T X + \lambda I)^{-1} X^T y$$

- When  $\lambda$  is 0, recover our standard LR solution
- When  $\lambda > 0$ ,  $w^*$  always exists,  $(X^T X + \lambda I)$  is always invertible
- What happens when  $\lambda \rightarrow \infty$ ?
  - $w^* \rightarrow 0$


# Regularization



- Bigger  $\lambda \rightarrow$  simpler solution
- When  $\lambda$  is too big, we get *underfitting*
- We now have a way of going from underfitting to overfitting without changing the model family!


# Model selection

- No regularization
- Overfitting



# Model selection


- With regularization
- No overfitting!
- Strength of regularization should be tuned via validation loss



# Regularization geometry

$$\min_w \|Xw - y\|^2 + \lambda \|w\|^2$$

- Let's consider  $\min \|Xw - y\|^2$ , s.t.  $\|w\| < \alpha$ .
- On the level set of the loss, we will find the point with the lowest norm.
- L2-regularization is also known as weight decay\*, it shrinks weights towards the origin
- L2-regularized linear regression is also known as Ridge regression




\*Strictly speaking, weight decay is not identical to L2 reg, but in practice it often is.

# Regularization geometry

$$\min \|Xw - y\|^2 + \lambda \|w\|_1$$


$$\min \|Xw - y\|^2, \text{ s.t. } \|w\|_1 < \alpha.$$

- The level set of L1-norm is a polyhedron.
- The minimizer will end up in a vertex of the polyhedron.
- L1-regularization produces sparse solutions!
- L1 reg is also known as LASSO (least absolute shrinkage and selection operator)
- We don't have a closed-form solution for LASSO!
  - Need to use gradient descent or fancy custom methods




# Regularization geometry


**Effect of Regularization on High-Degree Polynomial (degree=15)**



Regularization  $\alpha=0.01$   
Train MSE: 0.085, Test MSE: 0.047



# Regularization geometry




Ridge regression chooses low weights

$$\begin{aligned} & 5.5E-01 + 0.0E+00 x^1 + 6.0E-01 x^2 - 6.0E-01 x^3 \\ & - 5.3E-01 x^4 - 2.2E-01 x^5 + 5.0E-02 x^6 + \\ & 2.2E-01 x^7 + 3.0E-01 x^8 + 3.2E-01 x^9 + 3.0E-01 \\ & x^{10} + 2.4E-01 x^{11} + 1.8E-01 x^{12} + 1.0E-01 x^{13} \\ & + 2.1E-02 x^{14} - 5.9E-02 x^{15} - 1.4E-01 x^{16} \end{aligned}$$

Lasso fit is choosing a sparse solution


$$6.1E-01 + 4.6E-01 x^{11} + 7.2E-02 x^{12}$$



# Regularization geometry



Elastic Net: mix L1 + L2

$$\min_w \|Xw - y\|^2 + \lambda_1 \|w\|_2^2 + \lambda_2 \|w\|_1$$



# Regularization: Summary

- Regularization is a core idea in ML
  - L2 regularization is used everywhere, from HFT stock price models to LLMs
- Train flexible models without overfitting
- Strength of the regularization is a key parameter that needs to be tuned via (cross-)validation
- L1 regularization can do automatic feature selection




# Logistic Regression

8.26

0.34

# Classification




Classification:

$$x \in \mathbb{R}^{28^2}, \quad y \in \{0, 1, \dots, 9\}$$

- Inputs: arbitrary, e.g. images
- Labels: discrete set

# Classification




*Binary Classification:*

$$x \in \mathbb{R}^{28^2}, \quad y \in \{0, 1\}$$

- Inputs: arbitrary, e.g. images
- Labels: two classes

# Likelihood



Form of  $z$  can be  
arbitrary

Distribution of  $y | z$   
can be arbitrary


$$z = Xw, \quad p(y|z) = N(z, I)$$

$$\log p(y|X, w) = \log p(y|z) = \log N(z, 1) =$$

$$-\frac{1}{2} \|z - y\|^2 + C = -\frac{1}{2} \|Xw - y\|^2 + C.$$

- We can think of a model as defining the distribution of labels  $y$  given the parameters  $w$  of the model and inputs  $X$

# Likelihood



Form of  $z$  can be arbitrary

Distribution of  $y|z$  can be arbitrary


$$z = Xw, \quad p(y|z) = N(z, I)$$

$$\log p(y|X, w) = \log p(y|z) = \log N(z, 1) =$$

$$-\frac{1}{2} \|z - y\|^2 + C = -\frac{1}{2} \|Xw - y\|^2 + C.$$

- We can think of a model as defining the distribution of labels  $y$  given the parameters  $w$  of the model and inputs  $X$

# Likelihood



Form of  $z$  can be arbitrary

Distribution of  $y | z$  can be arbitrary

$$z = Xw, \quad p(y|z) = N(z, I)$$


$$\log p(y|X, w) = \log p(y|z) = \log N(z, 1) =$$

$$-\frac{1}{2} \|z - y\|^2 + C = -\frac{1}{2} \|Xw - y\|^2 + C.$$

- We can think of a model as defining the distribution of labels  $y$  given the parameters  $w$  of the model and inputs  $X$
- This is a powerful way of defining losses
  - Regularization = prior

# Classification Loss

$$p(y_i|z_i) = \sigma(z_i) \quad \sigma(z_i) = \frac{1}{1 + e^{-z_i}}$$



- Likelihood given a score (logit) is the sigmoid (logistic) function


$$z = Xw$$

- If the logit  $z$  is linear, we get logistic regression

$$p(y_i|X_i, w) = \sigma(X_i w),$$
$$L(w) = \log(p(y|X, w)) = \sum_{i=1}^N \log (\sigma(X_i w)).$$

# Classification Loss


## Linear Regression



$$\log p(y|X, w) = \log p(y|z) = \log N(z, 1) =$$

$$-\frac{1}{2} \|z - y\|^2 + C = -\frac{1}{2} \|Xw - y\|^2 + C.$$


## Logistic Regression



$$p(y_i|X_i, w) = \sigma(X_i w),$$
$$L(w) = \log(p(y|X, w)) = \sum_{i=1}^N \log (\sigma(X_i w)).$$

# Classification Loss

## Linear Regression




$$\log p(y|X, w) = \log p(y|z) = \log N(z, 1) =$$

$$-\frac{1}{2} \|z - y\|^2 + C = -\frac{1}{2} \|Xw - v\|^2 + C.$$

$+ \lambda \|w\|^2$

## Logistic Regression




$$p(y_i|X_i, w) = \sigma(X_i w),$$
$$L(w) = \log(p(y|X, w)) = \sum_{i=1}^N \log (\sigma(X_i w)).$$

$+ \lambda \|w\|^2$

# Classification Loss

- Unlike linear regression, we cannot solve logistic regression in a closed form
- We will use gradient descent to minimize the loss


## Logistic Regression



$$p(y_i|X_i, w) = \sigma(X_i w),$$
$$L(w) = \log(p(y|X, w)) = \sum_{i=1}^N \log (\sigma(X_i w)).$$

$+ \lambda \|w\|^2$


# Logistic regression geometry



$$p(y_i|X_i, w) = \sigma(X_i w),$$

- The decision boundary ( $p = 0.5$ ) is a linear subspace
$$xw = 0$$
- The decision boundary goes through  $x=0$
- How can we separate the decision boundary from 0?
- How can we make the decision boundary non-linear?

# Logistic regression geometry




$$p(y_i | X_i, w) = \sigma(X_i w),$$


- The decision boundary ( $p = 0.5$ ) is a linear subspace  
 $xw = 0$
- The decision boundary goes through  $x=0$
- How can we separate the decision boundary from 0? Add a bias term
- How can we make the decision boundary non-linear? Use non-linear features

# Logistic regression geometry


**Concentric Circles**  
Logistic Regression (Acc: 0.510)



**Half Moons**  
Logistic Regression (Acc: 0.845)




**Multiple Blobs**  
Logistic Regression (Acc: 1.000)




What features would work well here?

# Logistic regression geometry




With degree-2 polynomial features, we can achieve perfect accuracy.

# Logistic regression: MNIST




Binary classification accuracy with logistic regression: 87.1%

# Logistic regression: MNIST



- Confusion matrix: predicted vs actual label distribution
- Accuracy: fraction of correct answers

# Logistic regression: MNIST



- Confusion matrix: predicted vs actual label distribution
- Accuracy: fraction of correct answers

---


$$\frac{\text{Green Box}}{\text{Green Box} + \text{Red Box}}$$

# Classification metrics

Other metrics for classification:

- Precision
- Recall


$$\bullet F_1 = \frac{2}{\text{recall}^{-1} + \text{precision}^{-1}} = 2 \frac{\text{precision} \cdot \text{recall}}{\text{precision} + \text{recall}}$$



How many retrieved items are relevant?  
How many relevant items are retrieved?

$$\text{Precision} = \frac{\text{green}}{\text{red} + \text{green}}$$
      
$$\text{Recall} = \frac{\text{green}}{\text{green} + \text{light grey}}$$

# Multiclass classification



Classification:

$$x \in \mathbb{R}^{28^2}, \quad y \in \{0, 1, \dots, 9\}$$

- Inputs: arbitrary, e.g. images
- Labels: discrete set

*How to define a likelihood?*

# Multiclass classification

*Binary*

*Classification*

$$y \in \{0, 1\}$$

$$z \in \mathbb{R}$$

$$p(y | z) = \sigma(z)$$

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

---

*Multiclass*

*Classification*

$$y \in \{0, 1, \dots, C - 1\}$$

$$z \in \mathbb{R}^C$$

$$p(y | z) = \text{softmax}(z)$$


# Multiclass classification

*Multiclass  
Classification*

$$y \in \{0, 1, \dots, C - 1\}$$

$$z \in \mathbb{R}^C$$

$$p(y | z) = \text{softmax}(z)$$



$$\sum_{c=1}^C \text{softmax}(z)_c = 1.$$

# Multiclass classification

*Multiclass  
Classification*

$$y \in \{0, 1, \dots, C - 1\}$$


$$z \in \mathbb{R}^C$$

$$p(y | z) = \text{softmax}(z)$$

$$z_i = X_i w, \quad p(y_i | z_i) = \text{softmax}(z_i)_{y_i}$$



$$X_i \in \mathbb{R}^d, w \in \mathbb{R}^{d \times C}, z_i \in \mathbb{R}^C,$$

$$y_i \in \{0, 1, \dots, C\}.$$



$$\sum_{c=1}^C \text{softmax}(z)_c = 1.$$

# Multiclass classification




Multiclass Classification Accuracy: 92%


# Multiclass classification

Misclassified Examples


True: 8, Pred: 5  
Confidence: 0.839




True: 4, Pred: 9  
Confidence: 0.374




True: 3, Pred: 2  
Confidence: 0.475




True: 9, Pred: 7  
Confidence: 0.514




True: 5, Pred: 8  
Confidence: 0.746



True: 2, Pred: 8  
Confidence: 0.845



True: 9, Pred: 7  
Confidence: 0.891



True: 5, Pred: 6  
Confidence: 0.686

