

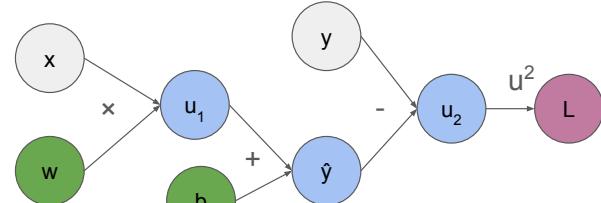
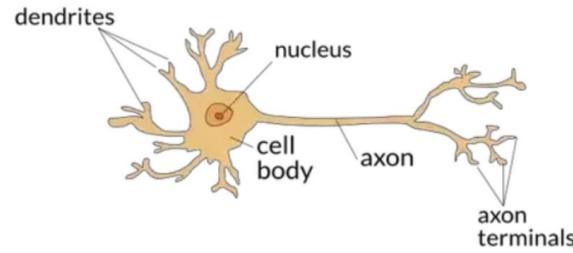
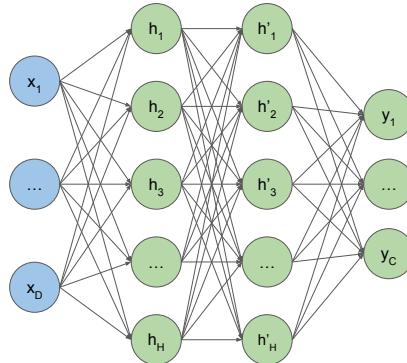
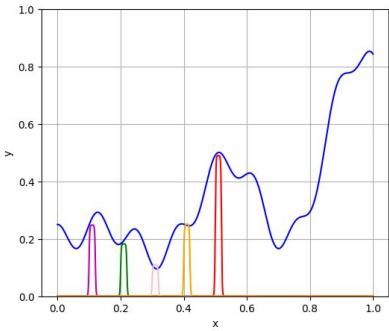
# NYU CS-GY 6923

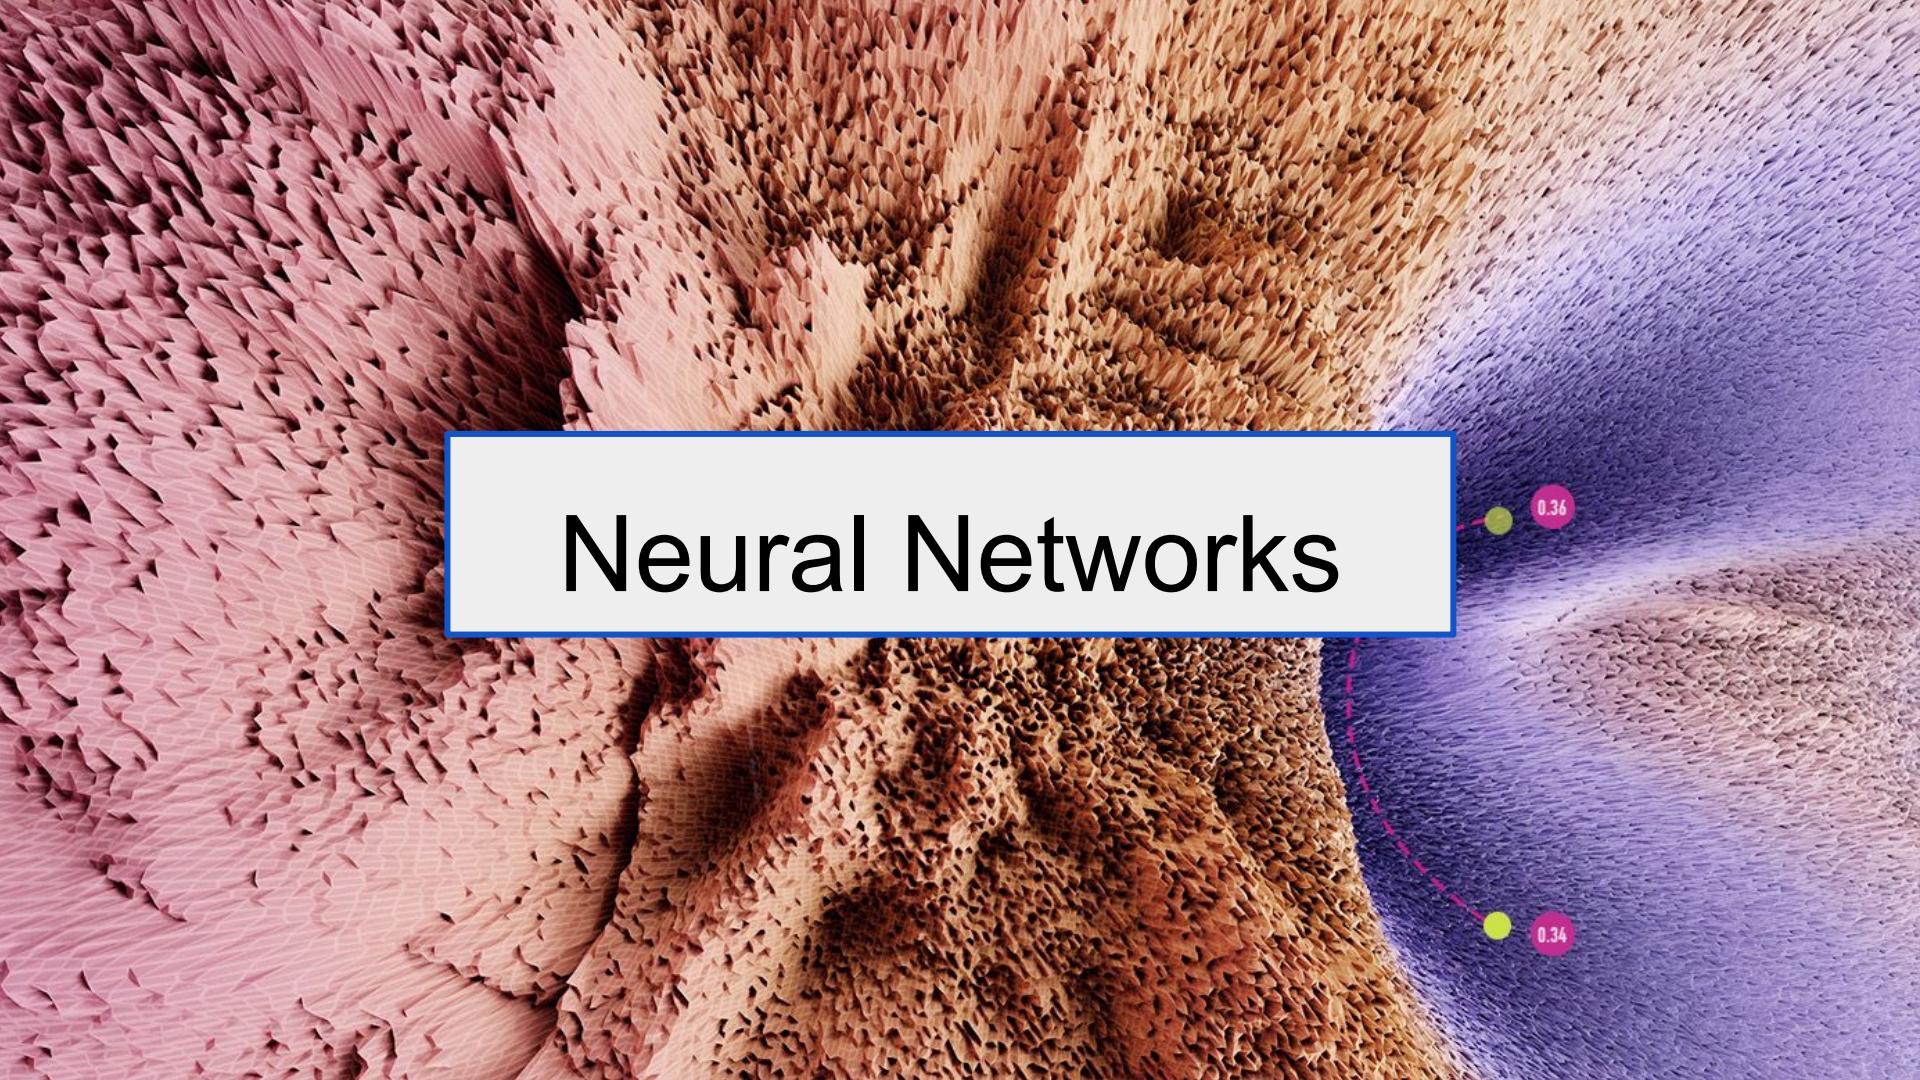
# Machine Learning

Prof. Pavel Izmailov

# Today

- Neural Networks
- Universal Approximation
- Backpropagation
- Autograd
- Demo



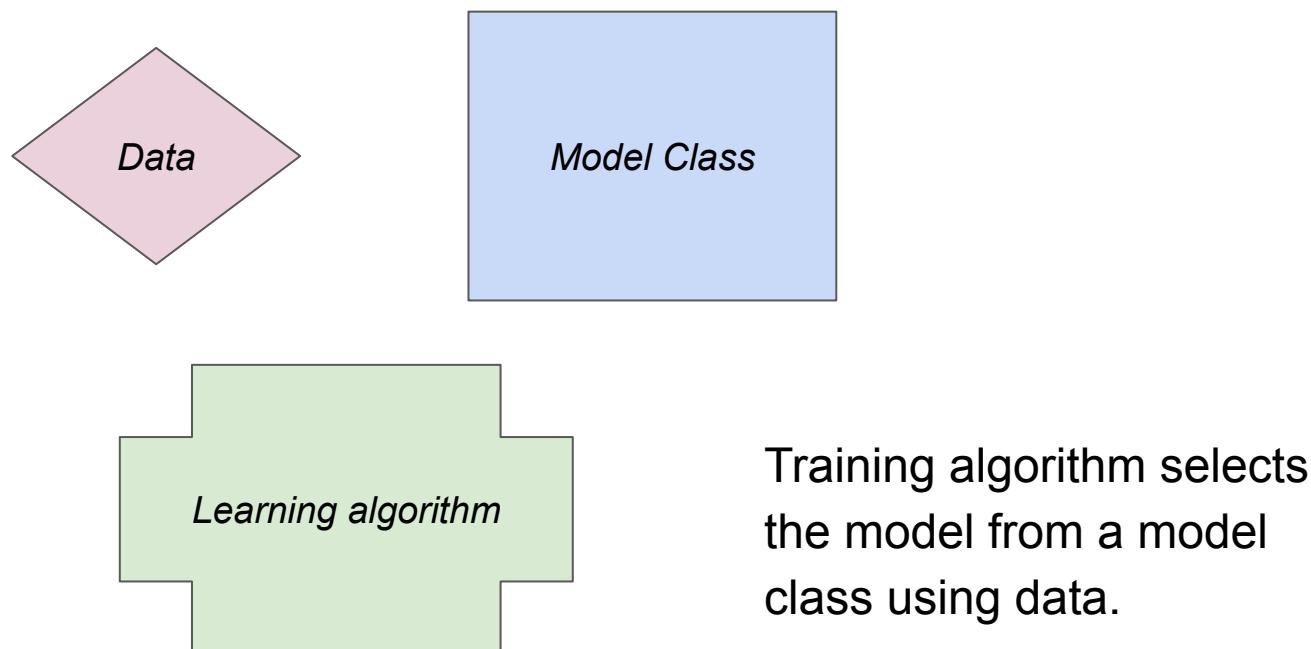


# Neural Networks

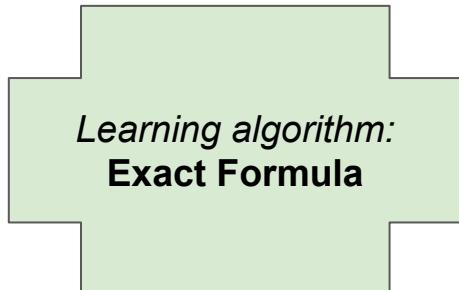
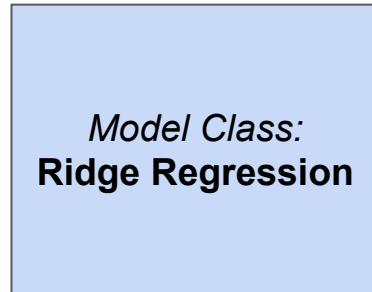
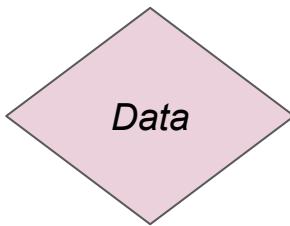
0.36

0.34

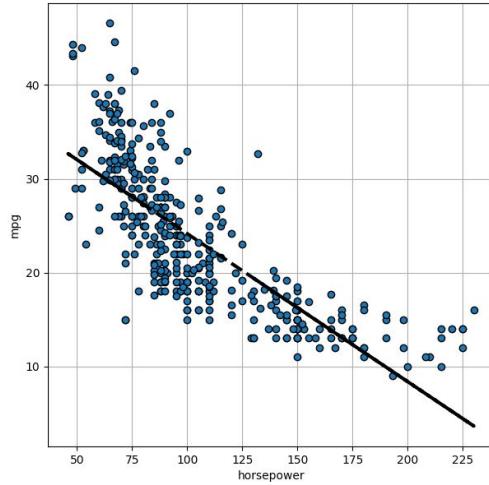
# General Picture of ML



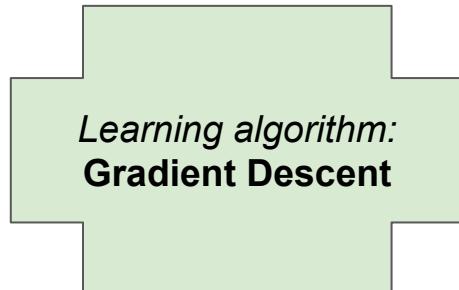
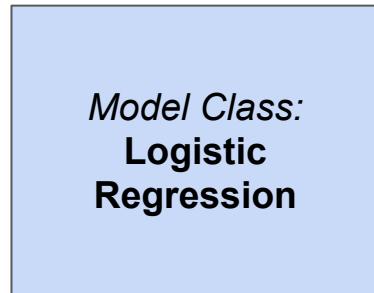
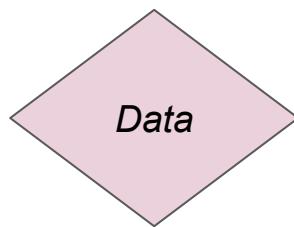
# Example: Linear Regression



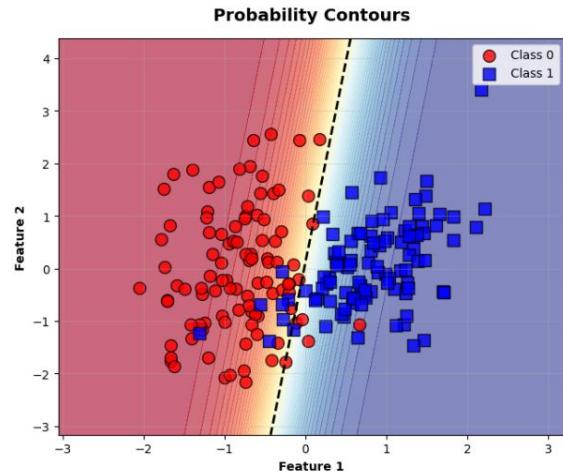
Training algorithm selects  
the model from a model  
class using data.



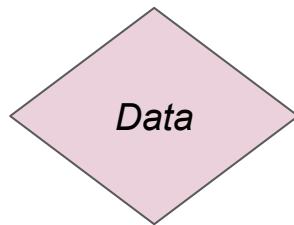
# Example: Logistic Regression



Training algorithm selects  
the model from a model  
class using data.

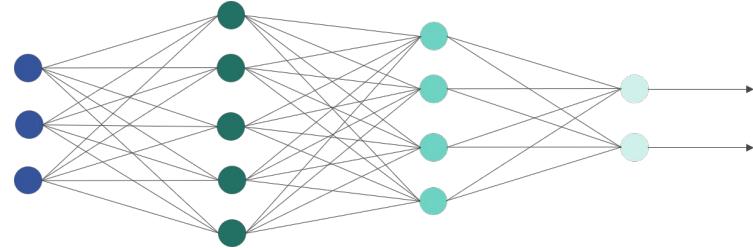


# GD-Based Learning



*Model Class:*  
**Any Model  
Differentiable wrt  
Parameters**

*Learning algorithm:*  
**Gradient Descent**



Training algorithm selects  
the model from a model  
class using data.

*Gradient Descent is an extremely powerful  
learning algorithm*

# GD-Based Learning

Gradient-based learning:

- Define a model:  $M(X_i, w)$  producing predictions based on data
  - $w$  represents parameters of the model
  - $M(X_i, w)$  typically needs to be differentiable with respect to  $w$
- Define a loss function  $L(w)$ 
  - Typically,  $L(w) = \sum_i L(M(X_i, w), y_i)$
- Minimize using gradient descent to find optimal  $w$

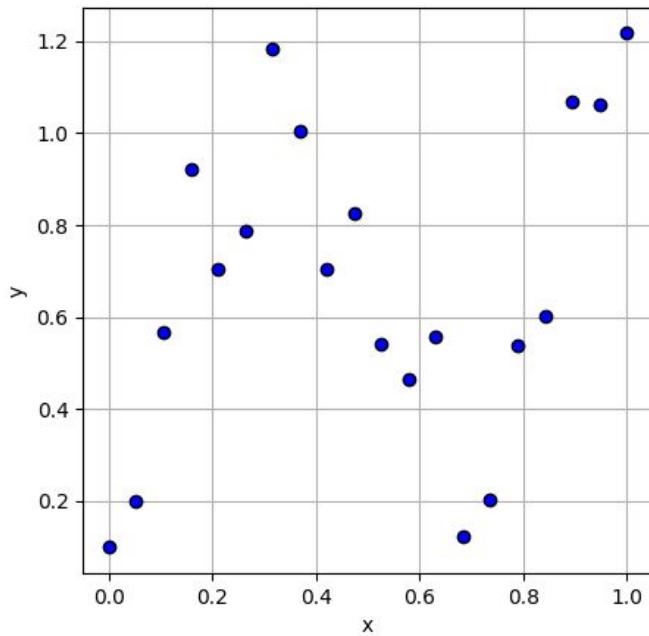
# GD-Based Learning

Gradient-based learning:

- Define a model:  $M(X_i, w)$  producing predictions based on data
  - $w$  represents parameters of the model
  - $M(X_i, w)$  typically needs to be differentiable with respect to  $w$
- Define a loss function  $L(w)$ 
  - Typically,  $L(w) = \sum_i L(M(X_i, w), y_i)$
- Minimize using gradient descent to find optimal  $w$

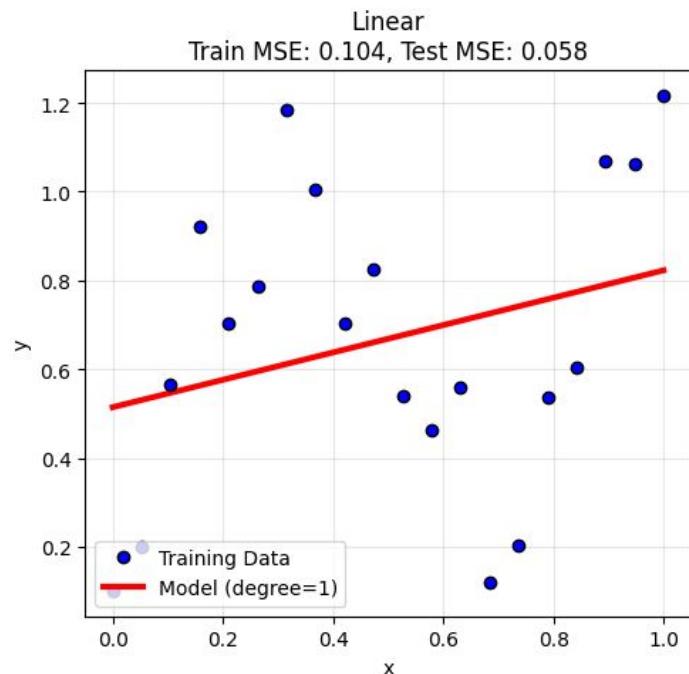
# Neural Networks: Motivation

- Non-linear data



# Neural Networks: Motivation

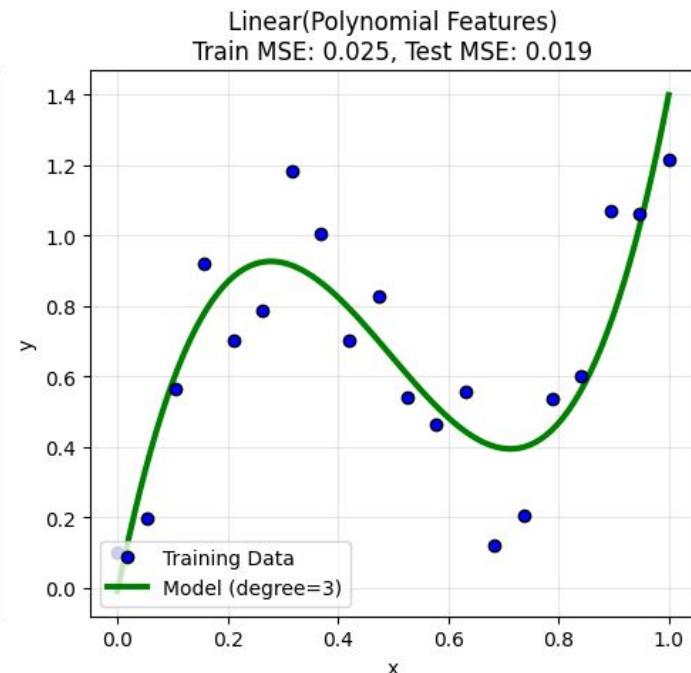
- Non-linear data
- Linear fit doesn't work



# Neural Networks: Motivation

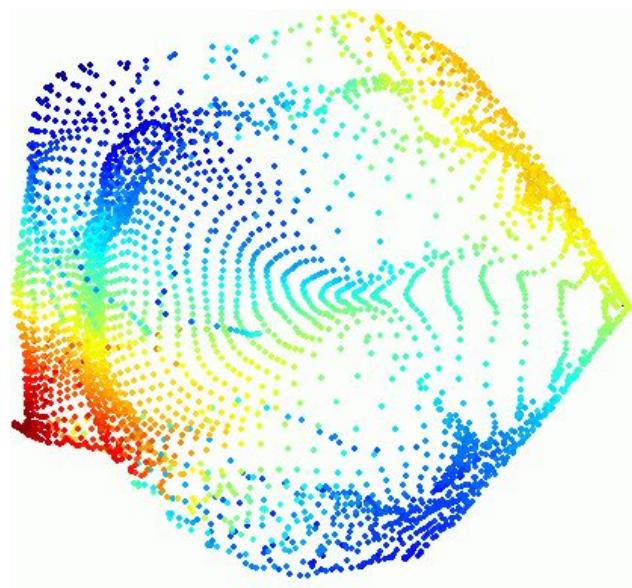
- Non-linear data
- Linear fit doesn't work
- Old solution: use polynomial features

$$(1, x, x^2, \dots, x^d)$$



# Neural Networks: Motivation

- Non-linear data
- Linear fit doesn't work
- Old solution: use polynomial features
$$(1, x, x^2, \dots, x^d)$$
- But what if our data is very high-dimensional and we don't know what features to use?



# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = W_1 x + b_1 \quad \xleftarrow{\text{Hidden features}}$$
$$\hat{y} = W_2 h + b_2 \quad \xleftarrow{\text{Output}}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- Will this model be able to learn features?

# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = W_1 x + b_1 \quad \xleftarrow{\text{Hidden features}}$$
$$\hat{y} = W_2 h + b_2 \quad \xleftarrow{\text{Output}}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- Will this model be able to learn features?
  - No, it is equivalent to a linear model!

$$W_2(W_1 x + b_1) + b_2 = (W_2 W_1)x + (b_2 + W_2 b_1) = W'x + b'$$

# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = W_1 x + b_1$$

Hidden features

$$\hat{y} = W_2 h + b_2$$

Output

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

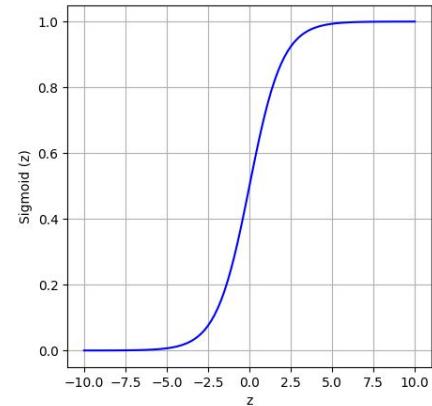
- Will this model be able to learn features?

- No, it is equivalent to a linear model!

$$W_2(W_1 x + b_1) + b_2 = (W_2 W_1)x + (b_2 + W_2 b_1) = W'x + b'$$

- We need a non-linearity!

$$h = \sigma(W_1 x + b_1)$$



# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = \sigma(W_1x + b_1) \quad \text{Hidden features}$$
$$\hat{y} = W_2h + b_2 \quad \text{Output}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- What loss can we use?

# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = \sigma(W_1x + b_1) \quad \text{Hidden features}$$
$$\hat{y} = W_2h + b_2 \quad \text{Output}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- What loss can we use? MSE

$$L(W_1, b_1, W_2, b_2) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = \sigma(W_1x + b_1) \quad \text{Hidden features}$$
$$\hat{y} = W_2h + b_2 \quad \text{Output}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- What loss can we use? MSE

$$L(W_1, b_1, W_2, b_2) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- How many parameters?

- Train: minimize loss with Adam / SGD

# Feature Engineering → Feature Learning

- Let's *learn* the features!

$$h = \sigma(W_1x + b_1) \quad \text{Hidden features}$$
$$\hat{y} = W_2h + b_2 \quad \text{Output}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- What loss can we use? MSE

$$L(W_1, b_1, W_2, b_2) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- Train: minimize loss with Adam / SGD

- How many parameters?  
DH + H + H + 1

# Feature Engineering → Feature Learning

- Let's *learn* the features!

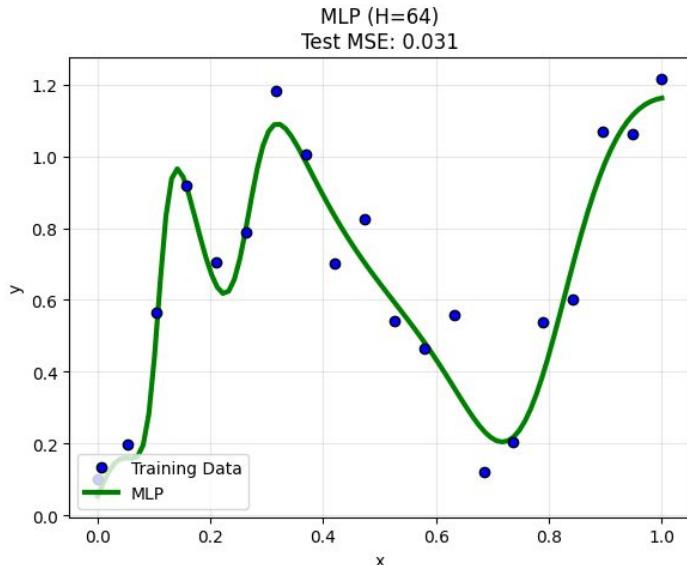
$$h = \sigma(W_1x + b_1) \quad \text{Hidden features}$$
$$\hat{y} = W_2h + b_2 \quad \text{Output}$$

$$h \in \mathbb{R}^H, x \in \mathbb{R}^D, \hat{y} \in \mathbb{R}$$

- What loss can we use? MSE

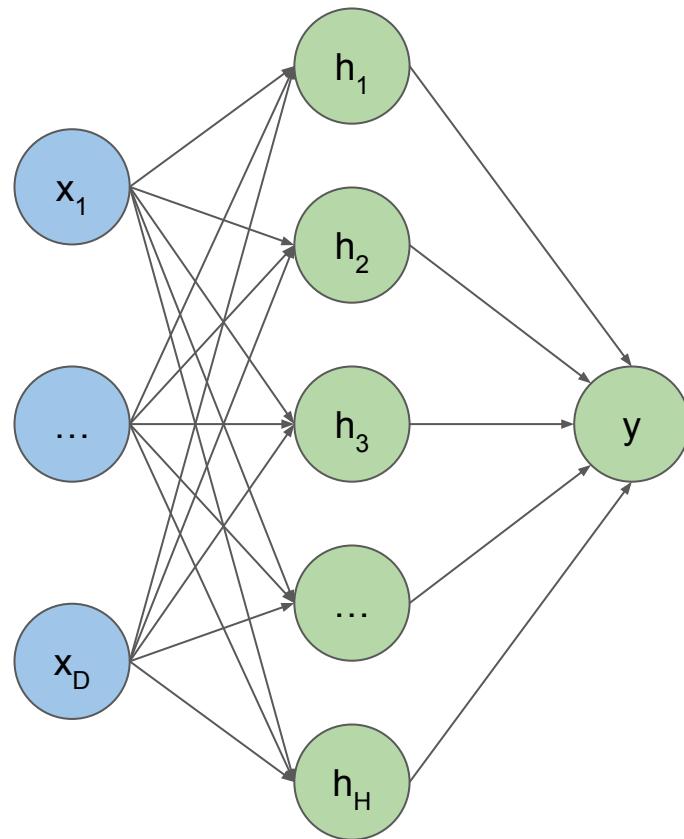
$$L(W_1, b_1, W_2, b_2) = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

- Train: minimize loss with Adam / SGD

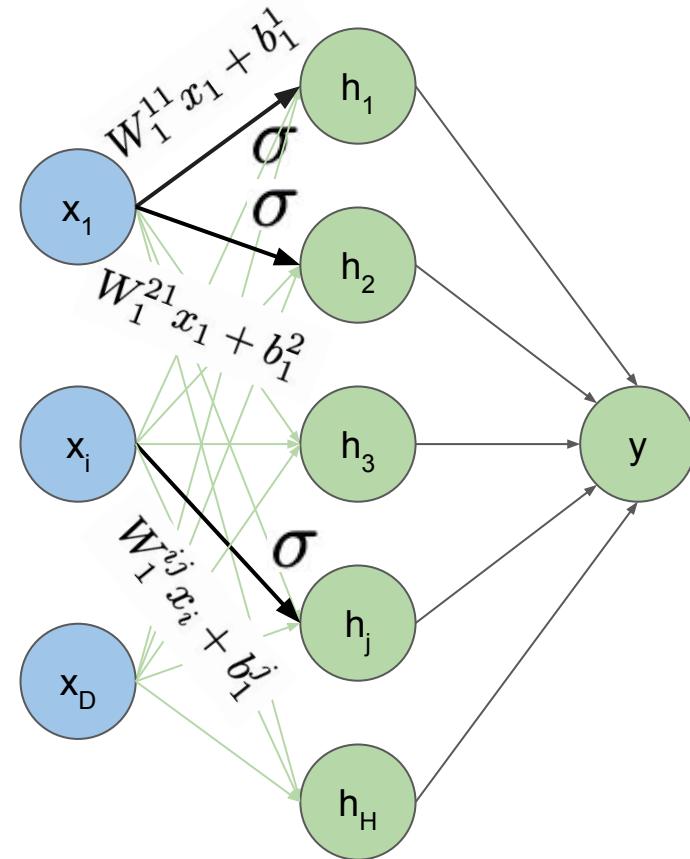


- How many parameters?  
 $DH + H + H + 1$

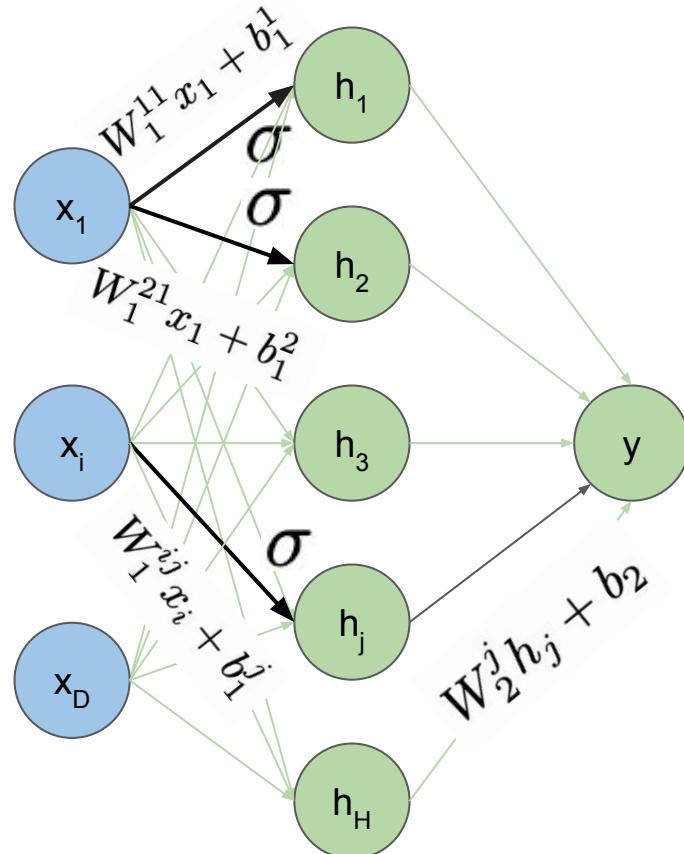
# Neural Networks



# Neural Networks



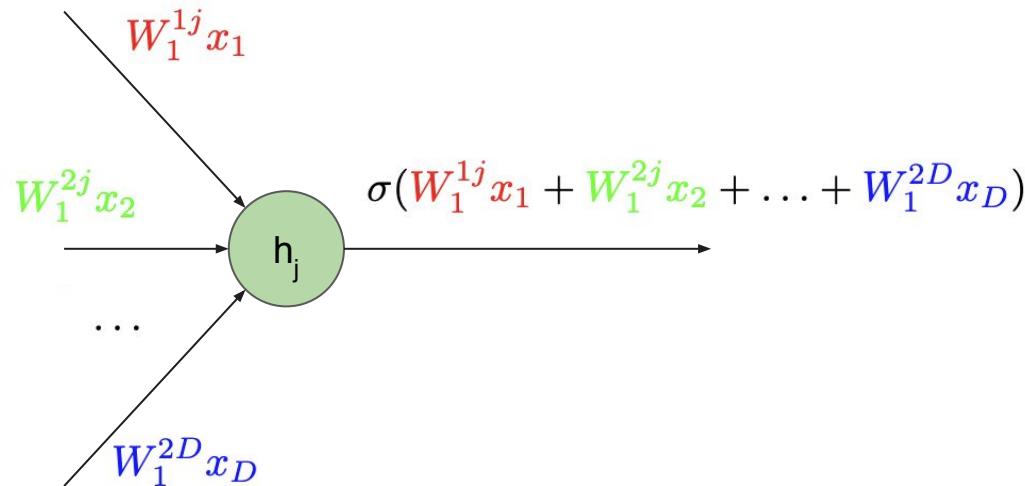
# Neural Networks



$$\hat{y} = W_2\sigma(W_1x + b_1) + b_2$$

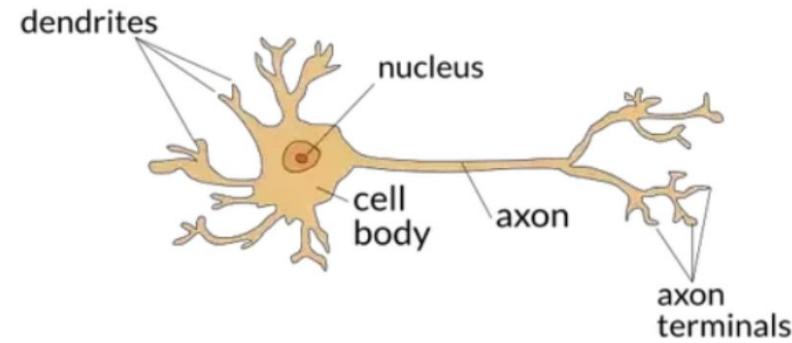
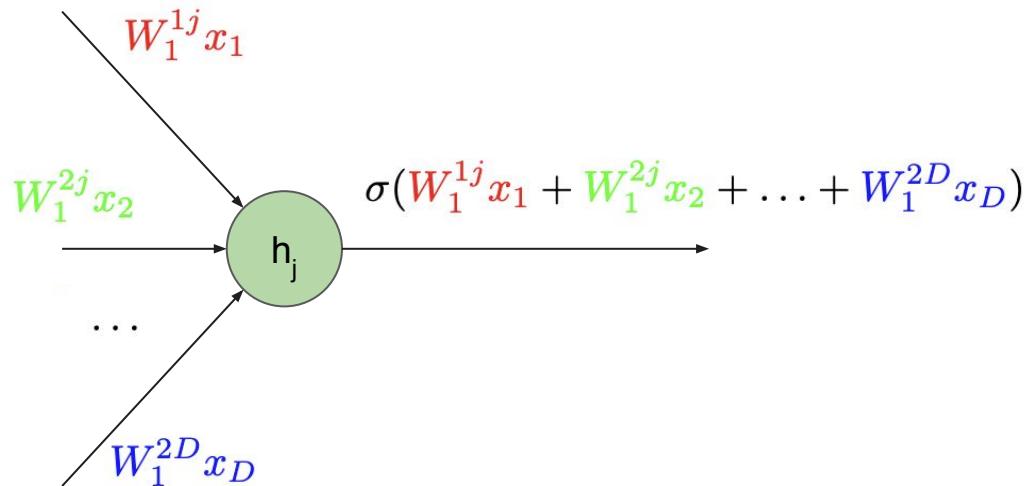
- Linear model on top of learned *hidden features  $h$*
- Hidden features are computed from input as a linear transform followed by element-wise non-linearity

# Neurons



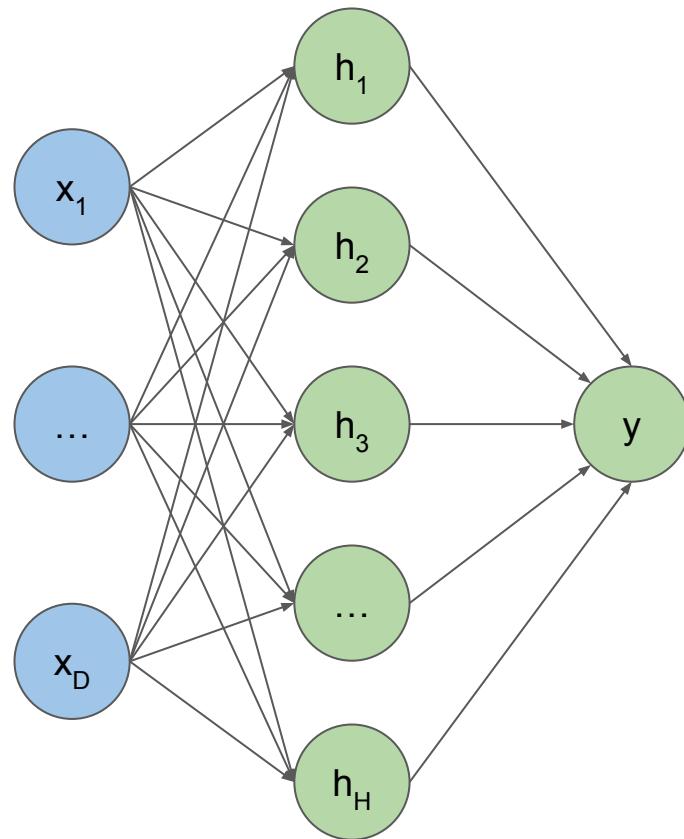
- Each node in the network is called a neuron
- Neurons aggregate their inputs and output an *activation*

# Neurons

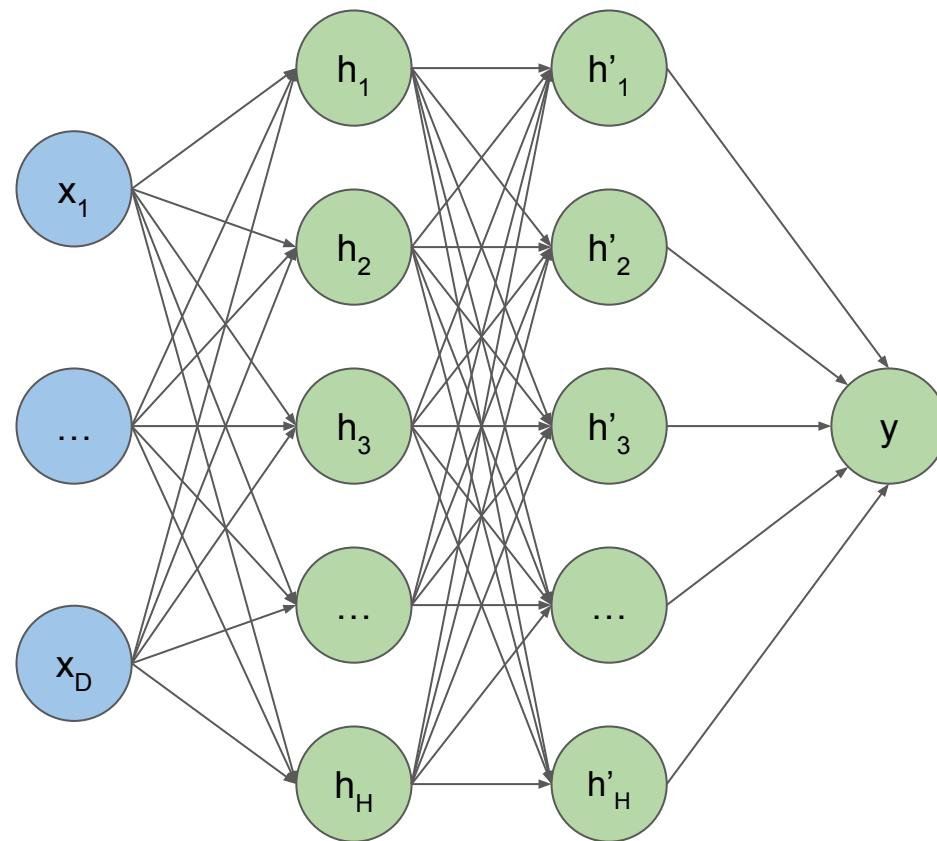


- Each node in the network is called a neuron
- Neurons aggregate their inputs and output an *activation*
- Rough analogy with a biological neuron

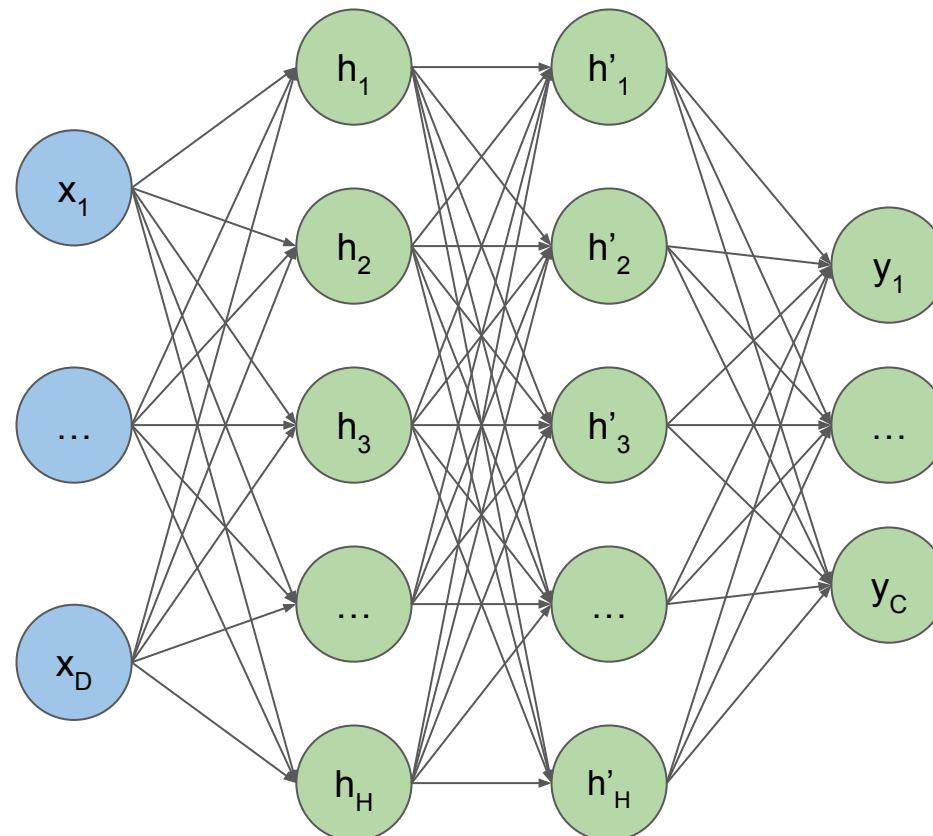
# Neural Networks



# Neural Networks: 2 hidden layers

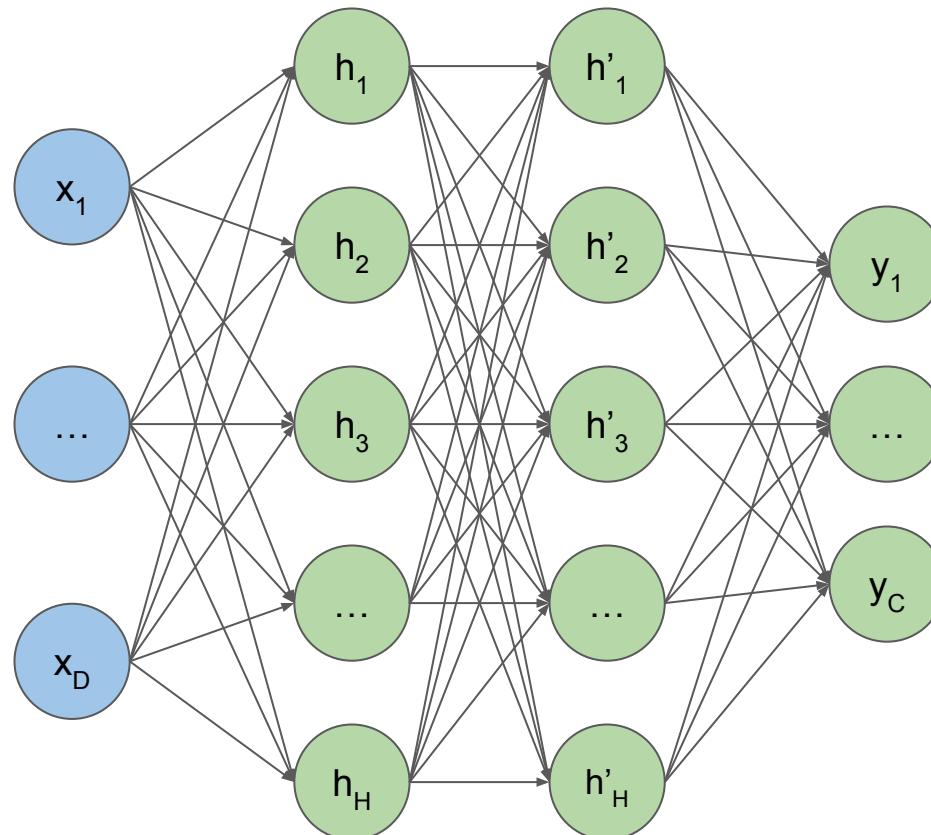


# Neural Networks: multiple outputs



- How many parameters does this model have as a function of  $D, H, C$ ?

# Neural Networks: multiple outputs



- How many parameters does this model have as a function of  $D, H, C$ ?

$$DH + H + HH + H + HC + C$$

# Universal Approximation Theorem

See <http://neuralnetworksanddeeplearning.com/>, chapter 4

# Universal Approximation Theorem

**Universal approximation theorem**—Let  $C(X, \mathbb{R}^m)$  denote the set of continuous functions from a subset  $X$  of a Euclidean  $\mathbb{R}^n$  space to a Euclidean space  $\mathbb{R}^m$ . Let  $\sigma \in C(\mathbb{R}, \mathbb{R})$ . Note that  $(\sigma \circ x)_i = \sigma(x_i)$ , so  $\sigma \circ x$  denotes  $\sigma$  applied to each component of  $x$ .

Then  $\sigma$  is not polynomial if and only if for every  $n \in \mathbb{N}$ ,  $m \in \mathbb{N}$ , compact  $K \subseteq \mathbb{R}^n$ ,  $f \in C(K, \mathbb{R}^m)$ ,  $\varepsilon > 0$  there exist  $k \in \mathbb{N}$ ,  $A \in \mathbb{R}^{k \times n}$ ,  $b \in \mathbb{R}^k$ ,  $C \in \mathbb{R}^{m \times k}$  such that

$$\sup_{x \in K} \|f(x) - g(x)\| < \varepsilon$$

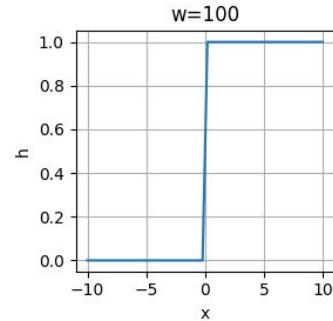
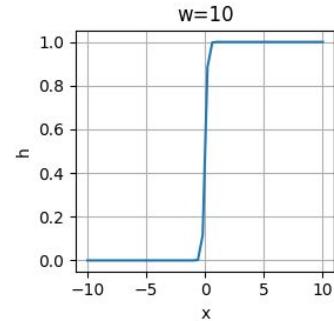
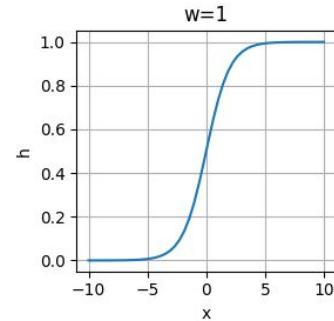
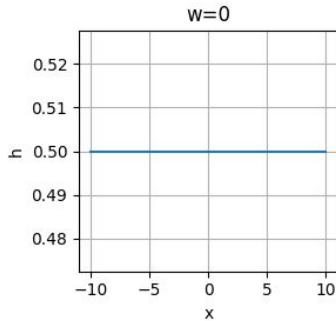
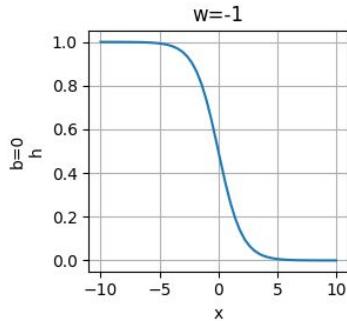
$$\text{where } g(x) = C \cdot (\sigma \circ (A \cdot x + b))$$

**Informal version:** a neural network with sufficient width with a single hidden layer can approximate any continuous function up to arbitrary precision.

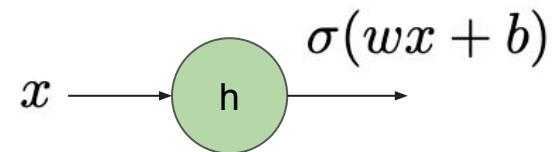
# Universal Approximation Theorem

Intuitive proof sketch following [Michael Nielsen's book](#):

$$h = \sigma(wx + b)$$



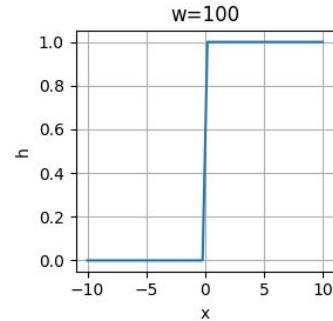
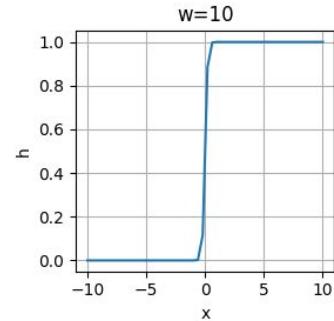
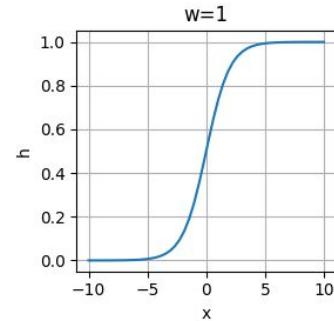
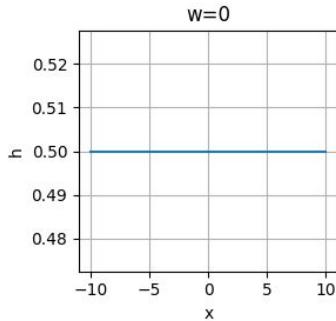
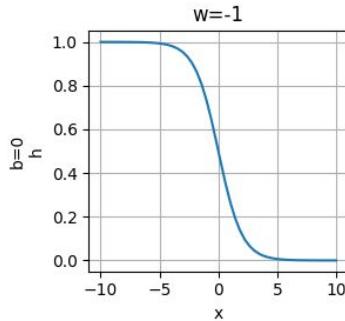
Let's look at the activations of a single neuron and change  $w$



# Universal Approximation Theorem

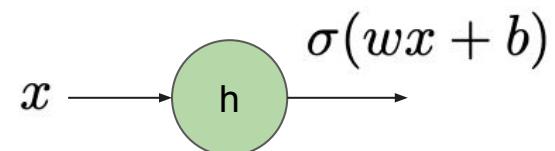
Intuitive proof sketch following [Michael Nielsen's book](#):

$$h = \sigma(wx + b)$$



Let's look at the activations of a single neuron and change  $w$

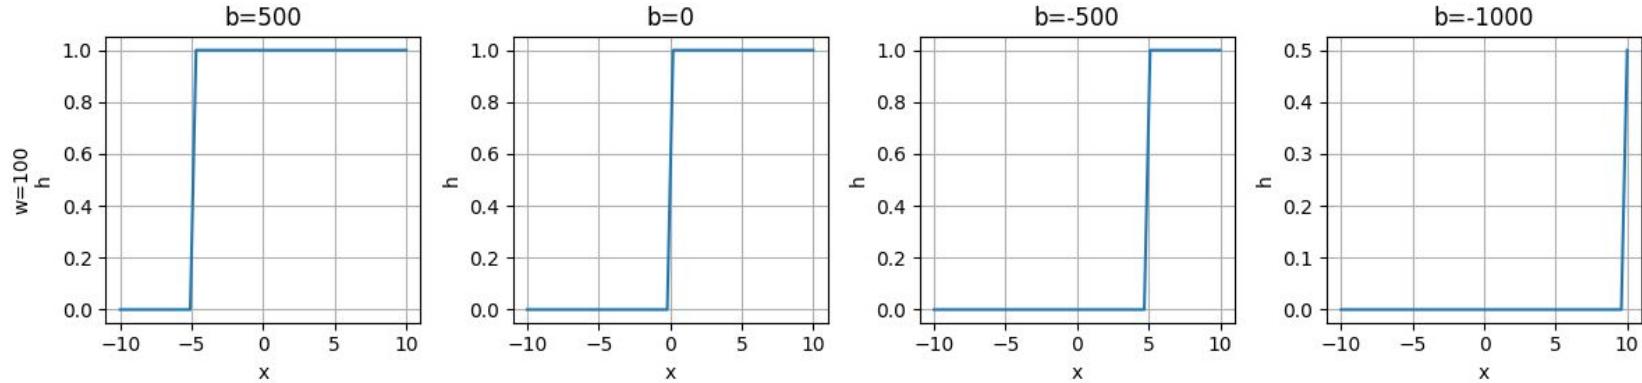
- Larger  $w \Rightarrow$  closer to step function



# Universal Approximation Theorem

Intuitive proof sketch following [Michael Nielsen's book](#):

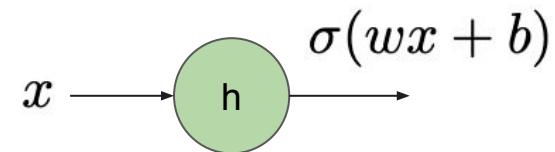
$$h = \sigma(wx + b)$$



Now, let's fix  $w$  and change  $b$ .

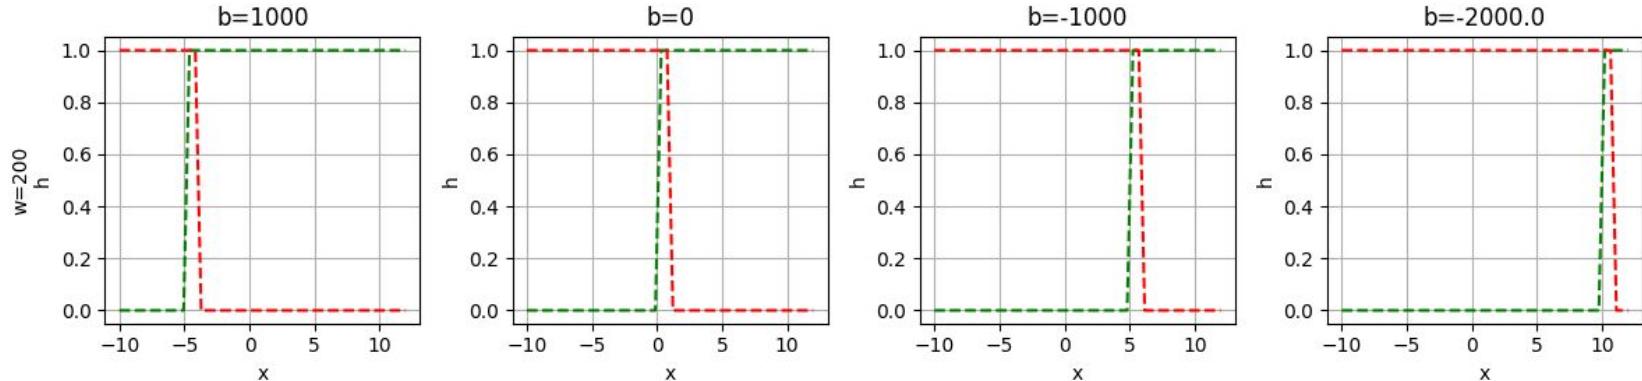
$$\sigma(w\hat{x} + b) = 0.5 \Leftrightarrow b = -w\hat{x}.$$

- $b$  shifts the curve horizontally



# Universal Approximation Theorem

Intuitive proof sketch following [Michael Nielsen's book](#):



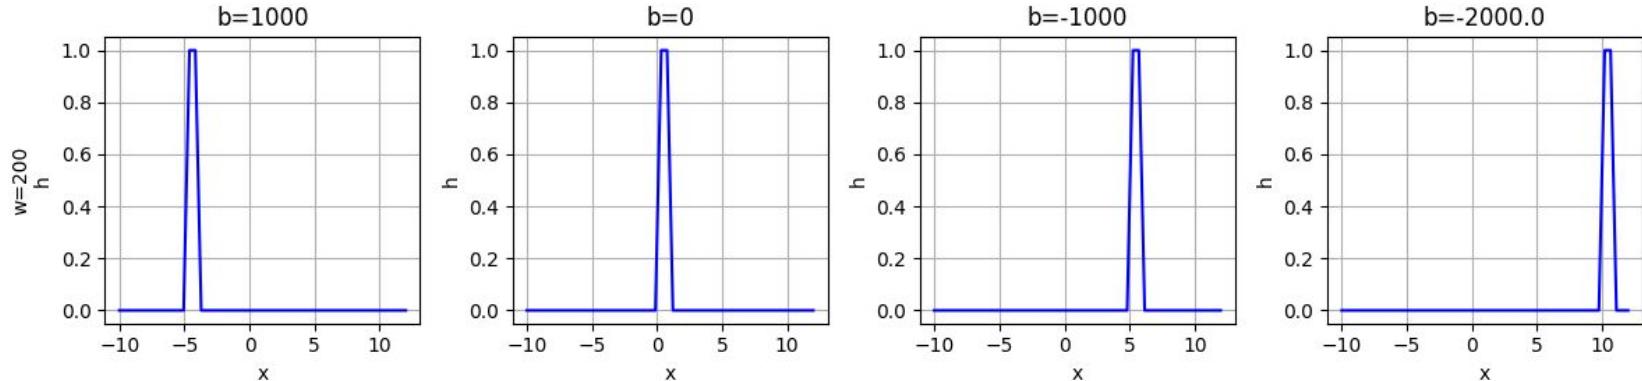
Now, let's put two of these together, but pointing in the opposite directions!

$$\sigma(wx - w\hat{x}) + \sigma(-wx + w(\hat{x} + \epsilon)) - 1$$

- We get a bump at a desired position!

# Universal Approximation Theorem

Intuitive proof sketch following [Michael Nielsen's book](#):



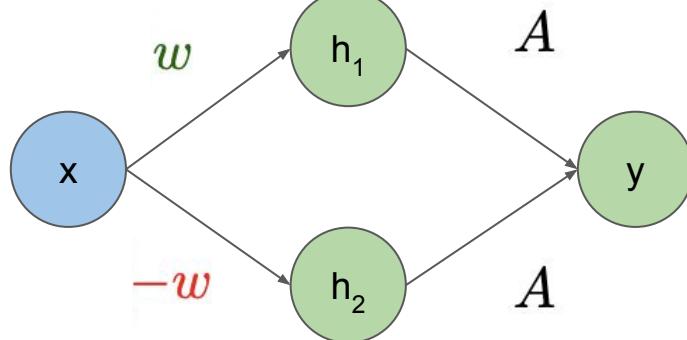
Now, let's put two of these together, but pointing in the opposite directions!

$$\sigma(wx - w\hat{x}) + \sigma(-wx + w(\hat{x} + \epsilon)) - 1$$

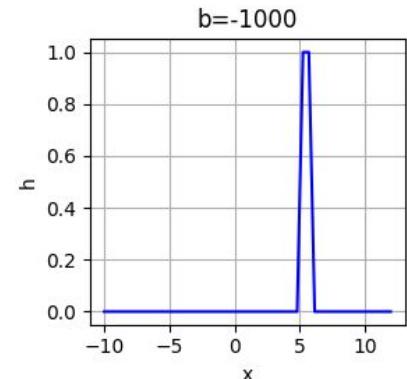
- We get a bump at a desired position!

# Universal Approximation Theorem

Intuitive proof sketch following [Michael Nielsen's book](#):

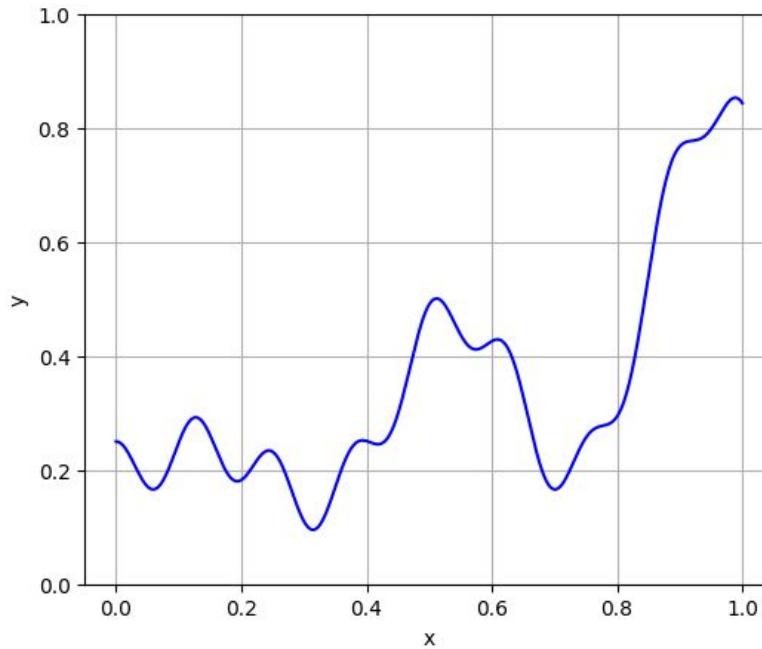


$$A \cdot (\sigma(wx - w\hat{x}) + \sigma(-wx + w(\hat{x} + \epsilon)) - 1)$$



- We can create a bump with a neural net with two hidden units

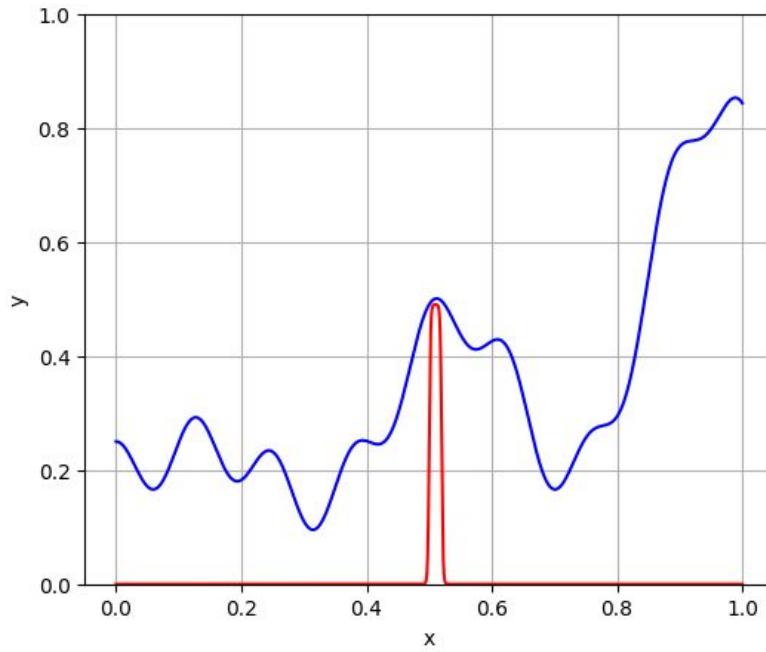
# Universal Approximation Theorem



Let's approximate a function  $f$  with a collection of bumps.

$$\sum_{j=1}^M f(\hat{x}_j) \cdot (\sigma(wx - w\hat{x}_j) + \sigma(-wx + w(\hat{x}_j + \epsilon)) - 1)$$

# Universal Approximation Theorem

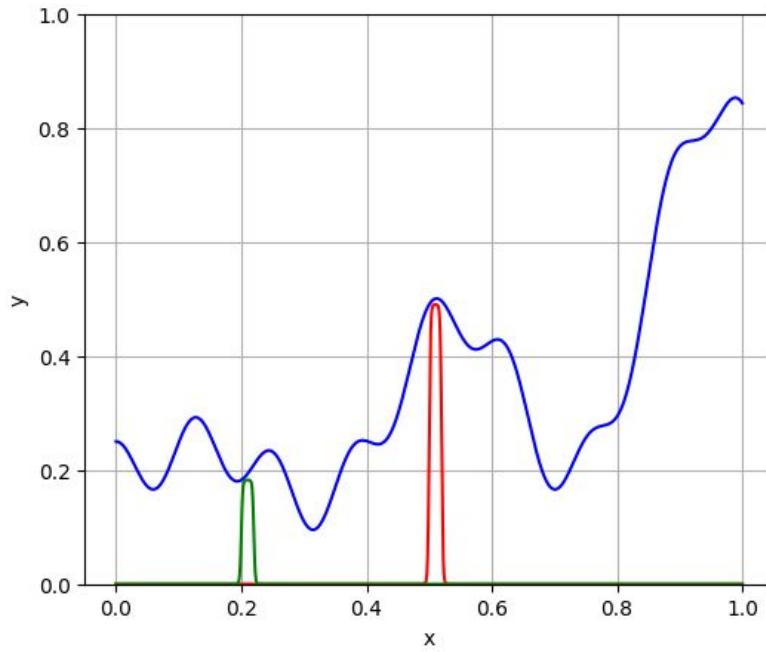


Let's approximate a function  $f$  with a collection of bumps.

- Start with a bump at 0.5

$$\sum_{j=1}^M f(\hat{x}_j) \cdot (\sigma(wx - w\hat{x}_j) + \sigma(-wx + w(\hat{x}_j + \epsilon)) - 1)$$

# Universal Approximation Theorem

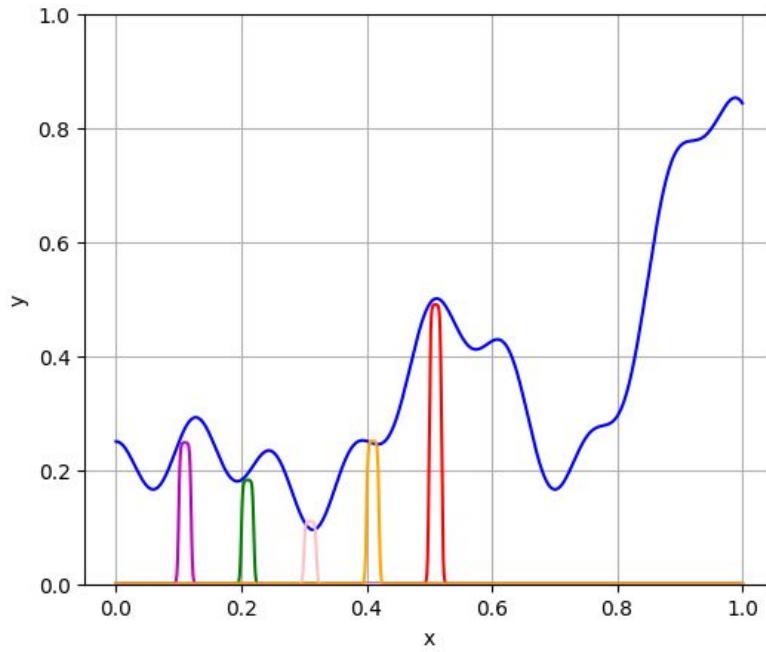


Let's approximate a function  $f$  with a collection of bumps.

- Start with a bump at 0.5
- Add another one at 0.2
- 

$$\sum_{j=1}^M f(\hat{x}_j) \cdot (\sigma(wx - w\hat{x}_j) + \sigma(-wx + w(\hat{x}_j + \epsilon)) - 1)$$

# Universal Approximation Theorem

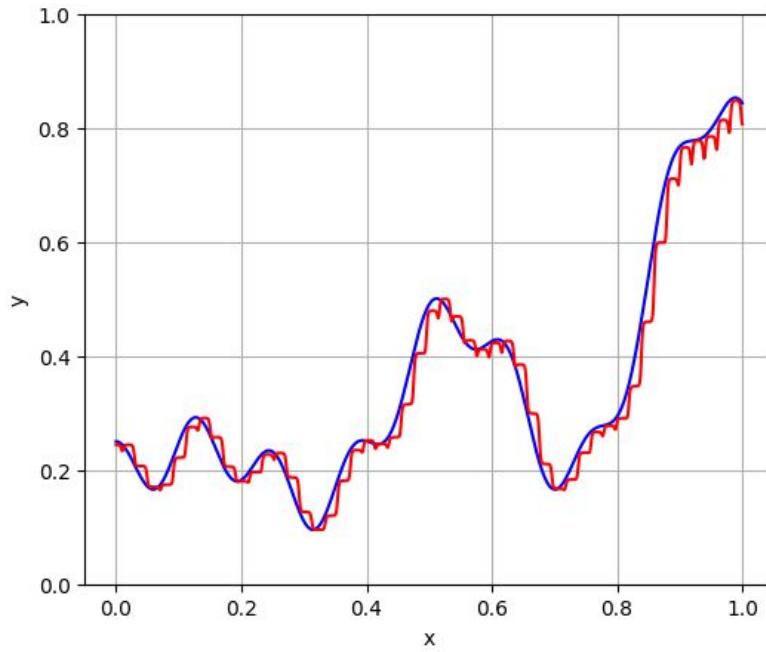


Let's approximate a function  $f$  with a collection of bumps.

- Start with a bump at 0.5
- Add another one at 0.2
- Keep going...

$$\sum_{j=1}^M f(\hat{x}_j) \cdot (\sigma(wx - w\hat{x}_j) + \sigma(-wx + w(\hat{x}_j + \epsilon)) - 1)$$

# Universal Approximation Theorem

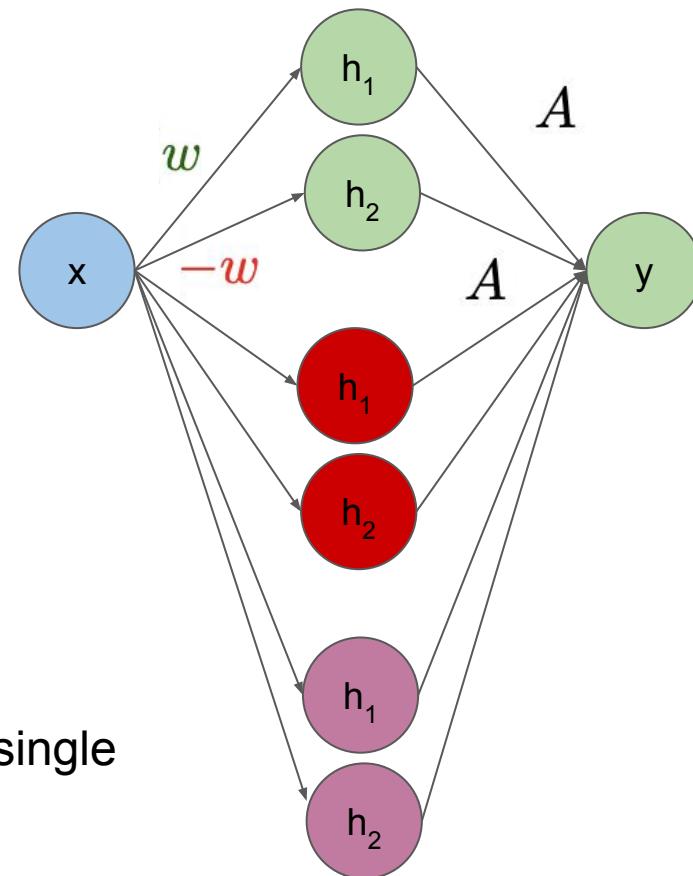
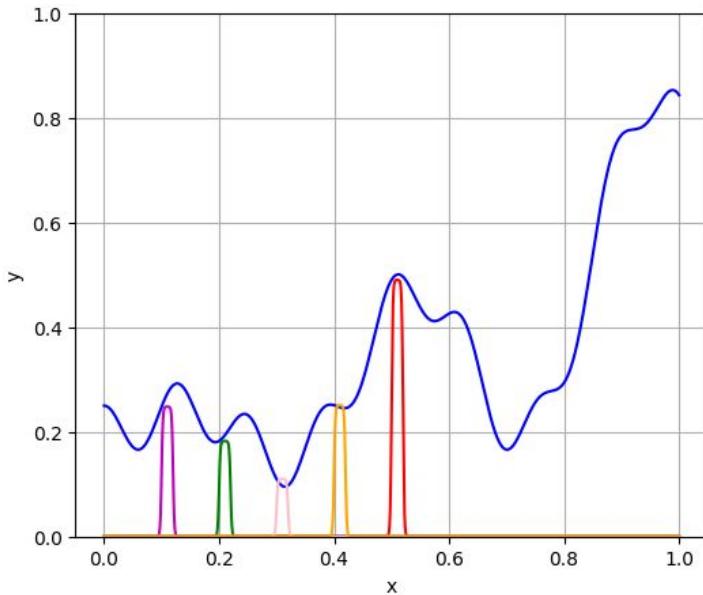


Let's approximate a function  $f$  with a collection of bumps.

- Start with a bump at 0.5
- Add another one at 0.2
- Keep going...
- 100 sigmoid pairs!

$$\sum_{j=1}^M f(\hat{x}_j) \cdot (\sigma(wx - w\hat{x}_j) + \sigma(-wx + w(\hat{x}_j + \epsilon)) - 1)$$

# Universal Approximation Theorem

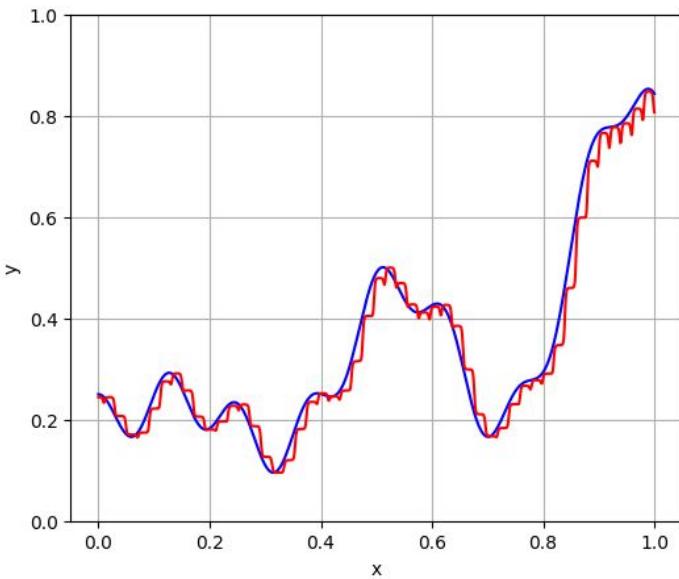


Putting the bumps together, we still get a single layer network.

# Universal Approximation Theorem

The theorem is still true if

- Multiple inputs
- Multiple outputs
- Almost arbitrary non-linearity



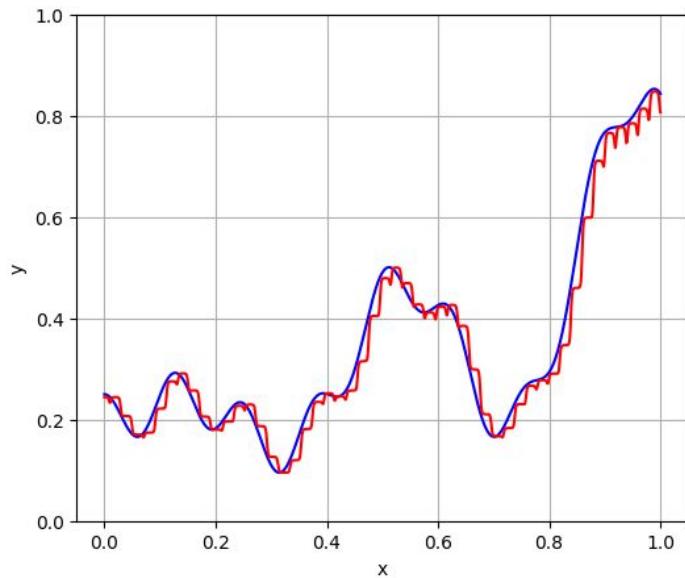
# Universal Approximation Theorem

The theorem is still true if

- Multiple inputs
- Multiple outputs
- Almost arbitrary non-linearity

Is this how neural networks approximate functions?

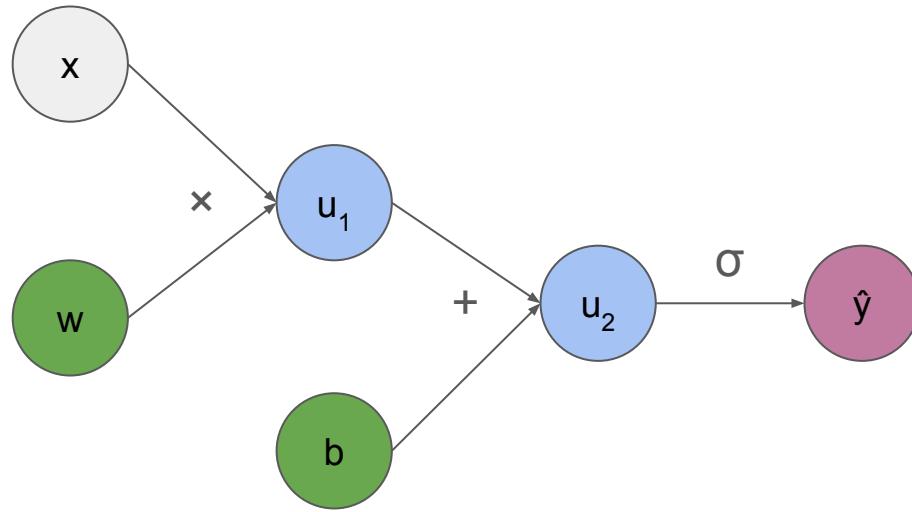
- No, this is a theoretical construction; models are more parameter-efficient
- In practice, depth makes a difference



# Backpropagation

See <https://www.deeplearningbook.org/>, chapter 6.5

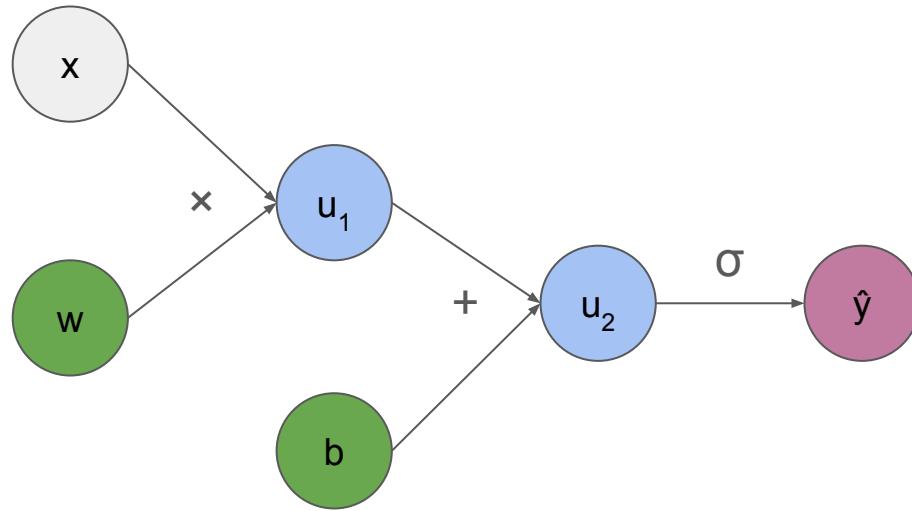
# Computation Graphs



We will represent computations as graphs

- Each nodes represents a variable: input, output or intermediate
- All non-input variables are results of atomic operations on other variables
- What could the computational graph on the left do?

# Computation Graphs



*Notice that we introduced auxiliary variables that we don't normally have.*

We will represent computations as graphs

- Each node represents a variable: input, output or intermediate
- All non-input variables are results of atomic operations on other variables
- What could the computational graph on the left do?

*Logistic regression prediction*

# Chain Rule

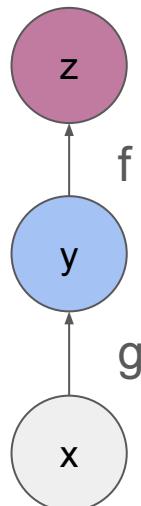
$$z = f(g(x))$$

$$y = g(x)$$

$$z = f(y)$$

$$\frac{dz}{dx} =$$

*All variables are scalars*



# Chain Rule

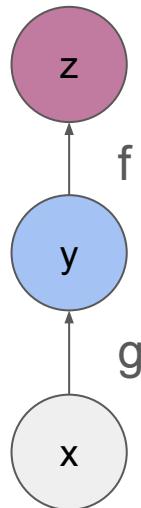
$$z = f(g(x))$$

$$y = g(x)$$

$$z = f(y)$$

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$

*All variables are scalars*



# Chain Rule

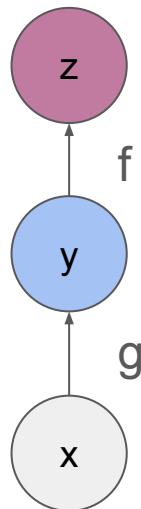
$$z = f(g(x))$$

$$y = g(x)$$

$$z = f(y)$$

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$

*All variables are scalars*



$$\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \cdot \frac{\partial y_j}{\partial x_i}$$

*If y and x are vectors*

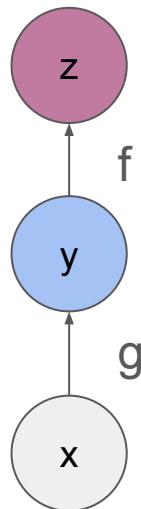
# Chain Rule

$$z = f(g(x))$$

$$y = g(x)$$

$$z = f(y)$$

$$\frac{dz}{dx} = \frac{dz}{dy} \cdot \frac{dy}{dx}$$



*All variables are scalars*

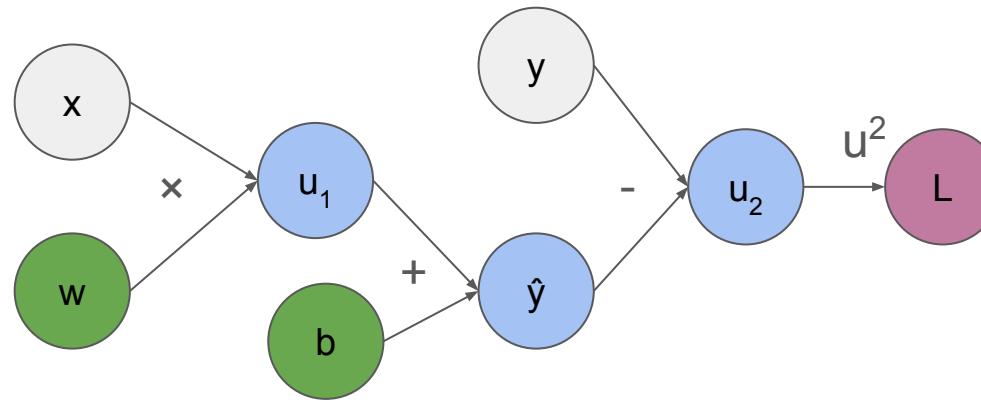
$$\frac{\partial z}{\partial x_i} = \sum_j \frac{\partial z}{\partial y_j} \cdot \frac{\partial y_j}{\partial x_i}$$

*If  $y$  and  $x$  are vectors*

*Jacobian matrix*

$$\nabla_x z = \left( \frac{\partial y}{\partial x} \right)^T \nabla_y z$$

# Backpropagation example: linear regression



Forward pass: evaluate  
the computation graph

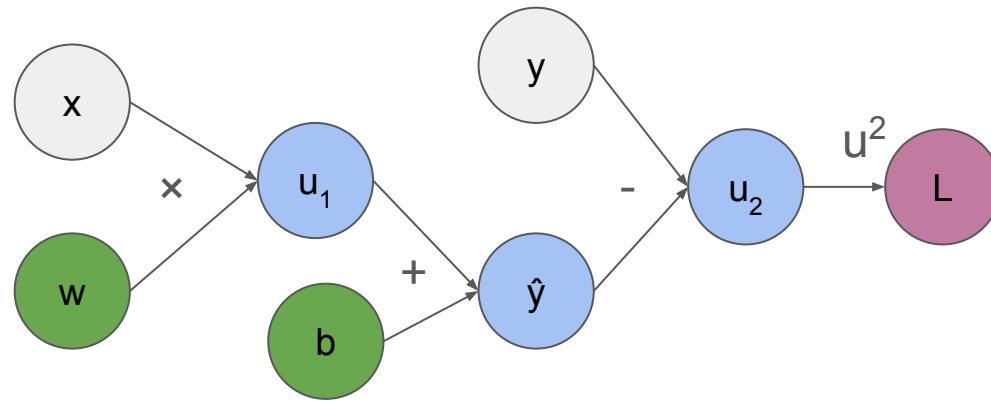
$$u_1 = xw$$

$$\hat{y} = u_1 + b$$

$$u_2 = y - \hat{y}$$

$$L = u_2^2$$

# Backpropagation example: linear regression



Backward pass: reverse  
the order of computation

Forward pass: evaluate  
the computation graph

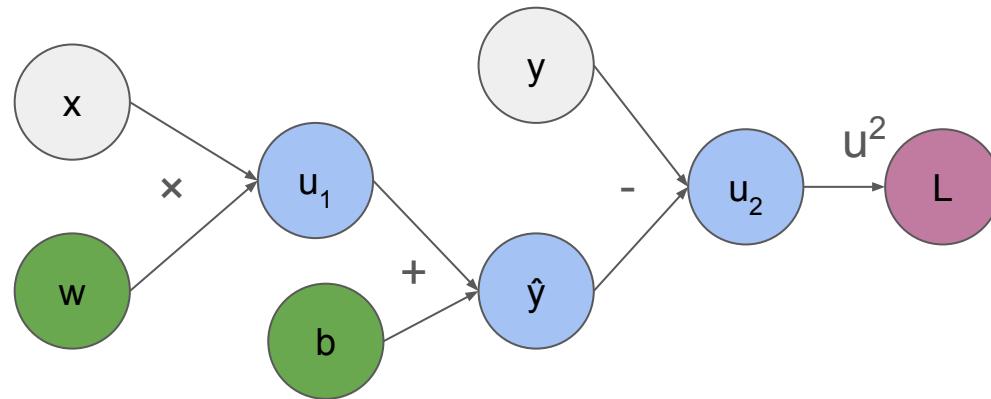
$$u_1 = xw$$

$$\hat{y} = u_1 + b$$

$$u_2 = y - \hat{y}$$

$$L = u_2^2 \quad \frac{\partial L}{\partial u_2} = 2u_2$$

# Backpropagation example: linear regression



Forward pass: evaluate  
the computation graph

$$u_1 = xw$$

$$\hat{y} = u_1 + b$$

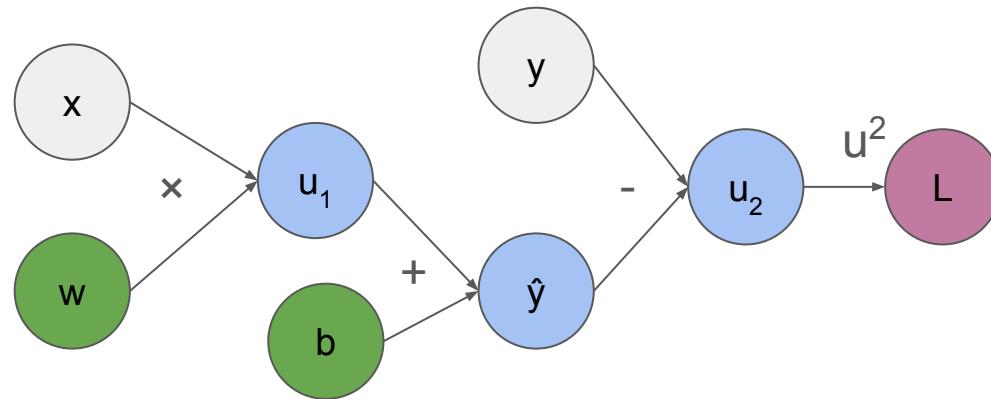
$$u_2 = y - \hat{y}$$

$$L = u_2^2$$

$$\frac{\partial L}{\partial u_2} = 2u_2 \quad \frac{\partial L}{\partial \hat{y}} = \frac{\partial L}{\partial u_2} \cdot \frac{\partial u_2}{\partial \hat{y}} = -2u_2$$

Backward pass: reverse  
the order of computation

# Backpropagation example: linear regression



Backward pass: reverse  
the order of computation

Forward pass: evaluate  
the computation graph

$$u_1 = xw$$

$$\hat{y} = u_1 + b$$

$$u_2 = y - \hat{y}$$

$$L = u_2^2$$

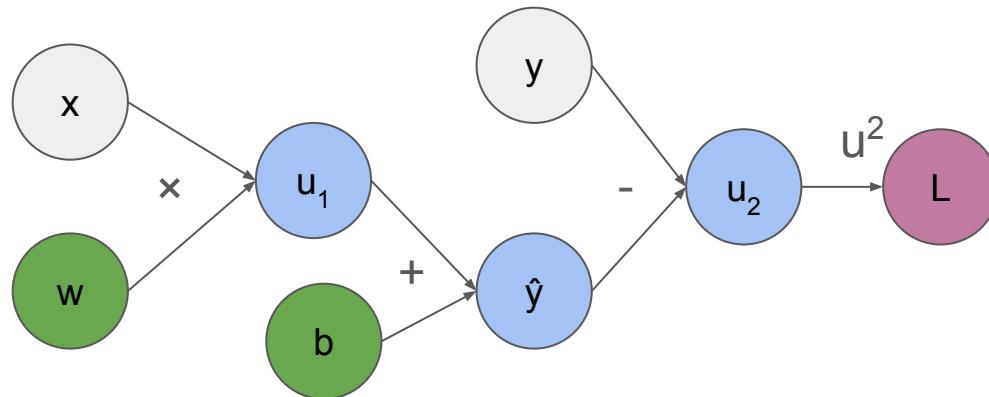
$$\frac{\partial L}{\partial u_2} = 2u_2$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial b} = -2u_2$$

$$\frac{\partial L}{\partial u_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial u_1} = -2u_2$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial L}{\partial u_2} \cdot \frac{\partial u_2}{\partial \hat{y}} = -2u_2$$

# Backpropagation example: linear regression



Forward pass: evaluate  
the computation graph

$$u_1 = xw$$

$$\hat{y} = u_1 + b$$

$$u_2 = y - \hat{y}$$

$$L = u_2^2$$

$$\frac{\partial L}{\partial u_2} = 2u_2$$

Backward pass: reverse  
the order of computation

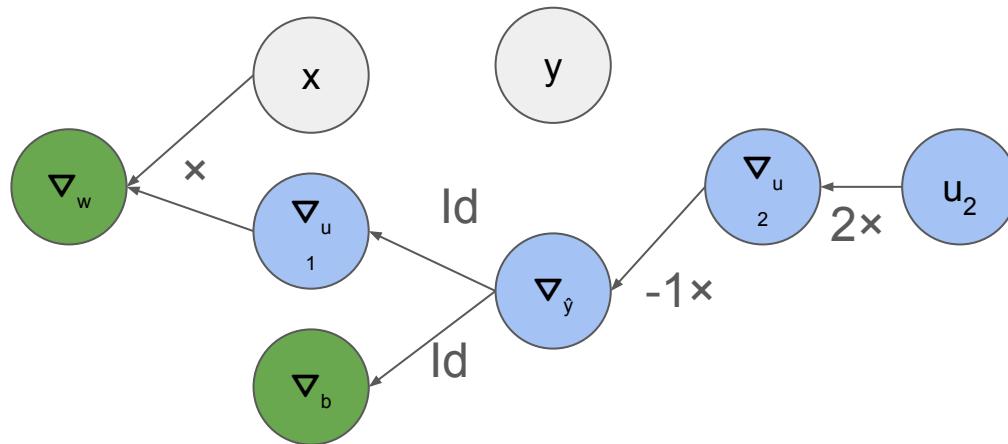
$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial u_1} \cdot \frac{\partial u_1}{\partial w} = -2u_2 \cdot x$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial b} = -2u_2$$

$$\frac{\partial L}{\partial u_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial u_1} = -2u_2$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial L}{\partial u_2} \cdot \frac{\partial u_2}{\partial \hat{y}} = -2u_2$$

# Backpropagation example: linear regression



Forward pass: evaluate  
the computation graph

$$u_1 = xw$$

$$\hat{y} = u_1 + b$$

$$u_2 = y - \hat{y}$$

$$L = u_2^2$$

$$\frac{\partial L}{\partial u_2} = 2u_2$$

Backward pass: reverse  
the order of computation

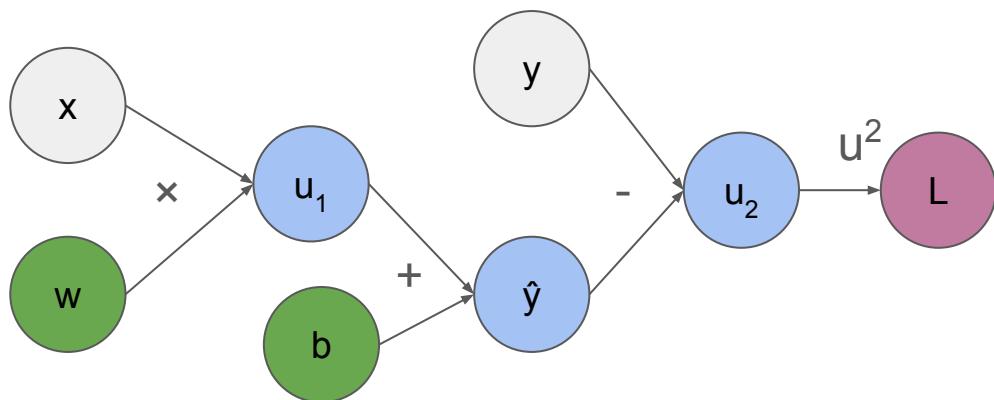
$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial u_1} \cdot \frac{\partial u_1}{\partial w} = -2u_2 \cdot x$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial b} = -2u_2$$

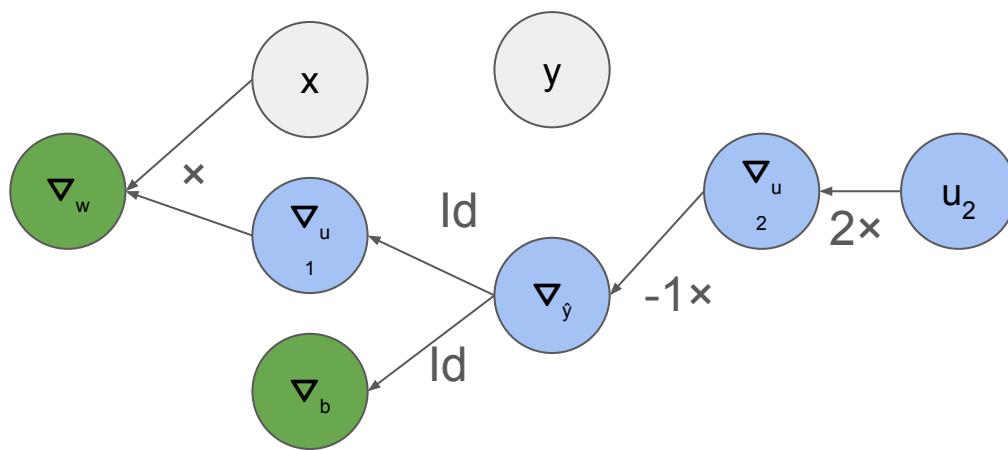
$$\frac{\partial L}{\partial u_1} = \frac{\partial L}{\partial \hat{y}} \cdot \frac{\partial \hat{y}}{\partial u_1} = -2u_2$$

$$\frac{\partial L}{\partial \hat{y}} = \frac{\partial L}{\partial u_2} \cdot \frac{\partial u_2}{\partial \hat{y}} = -2u_2$$

# Backpropagation example: linear regression



Forward pass

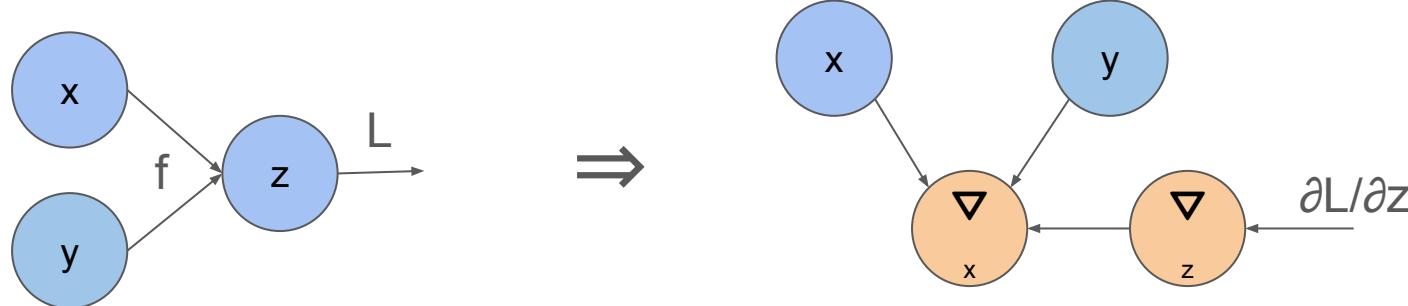


Backward pass

# Backward computation graph

We can always construct the backward graph algorithmically

- Nodes:
  - Keep all the nodes from the forward pass
  - Add a gradient node for each of the forward nodes
- Edges:
  - Backward of the original edges + a few rules



# Backward computation graph

We can always construct the backward graph algorithmically

- Nodes:
  - Keep all the nodes from the forward pass
  - Add a gradient node for each of the forward nodes
- Edges:
  - Backward of the original edges + a few rules



# Backward computation graph

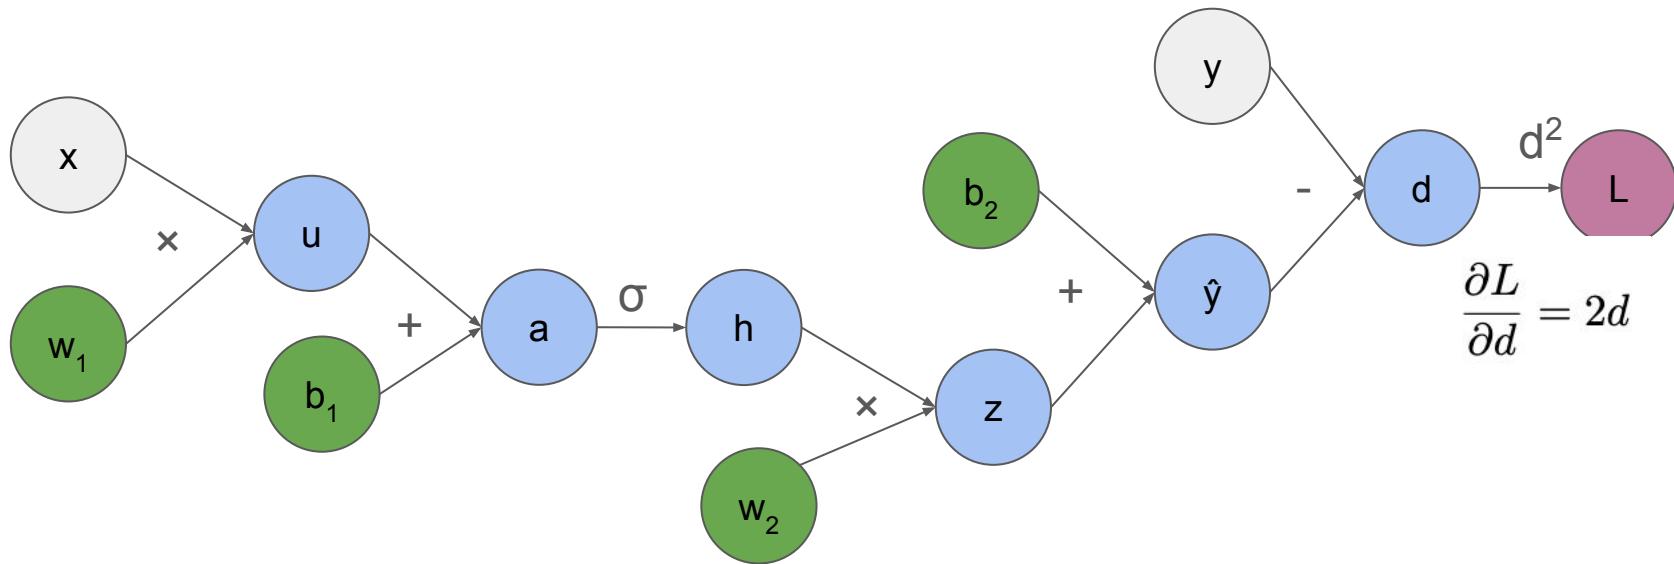
We can always construct the backward graph algorithmically

- Nodes:
  - Keep all the nodes from the forward pass
  - Add a gradient node for each of the forward nodes
- Edges:
  - Backward of the original edges + a few rules

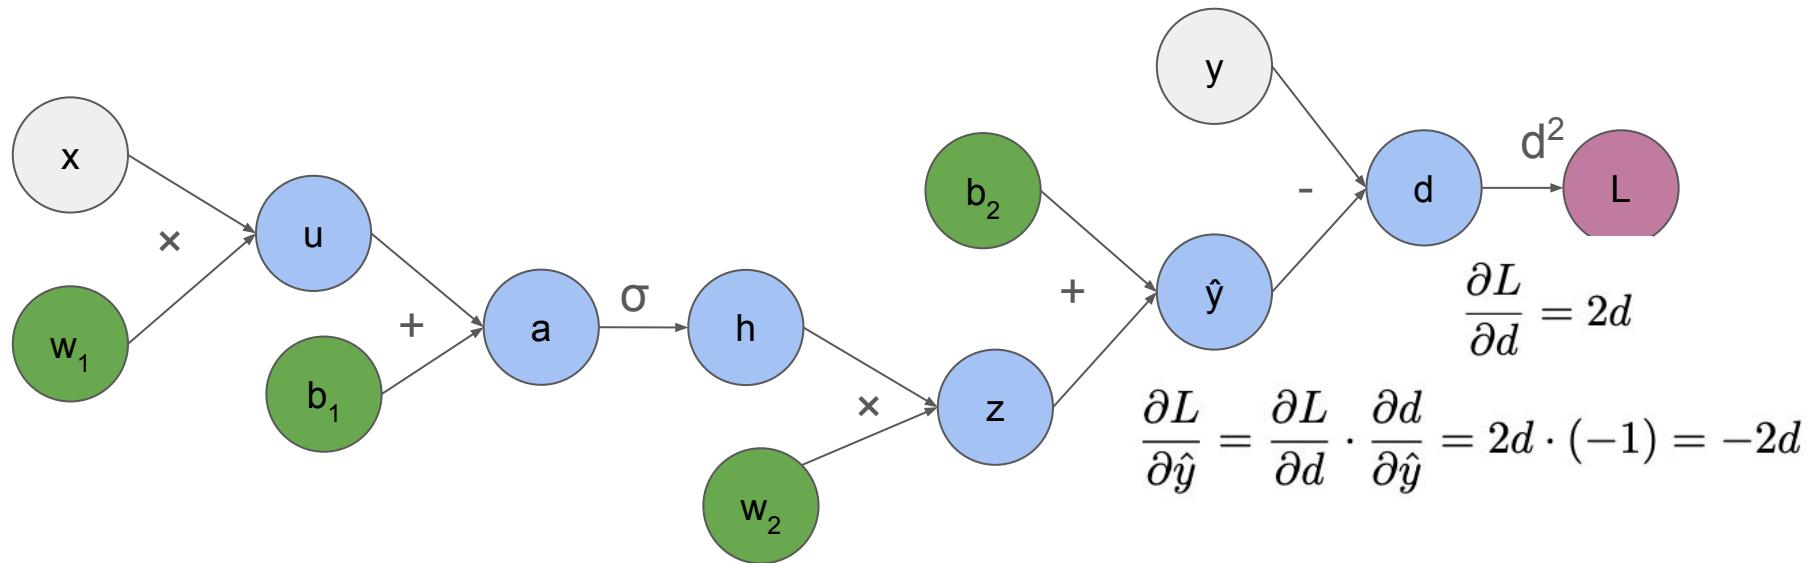
In general, backward pass:

- 2x the nodes, 2x memory
- Need to keep around all intermediate values
- Backward cost  $\leq c * \text{forward cost}$ ,  $c$  is 2-4 for NNs

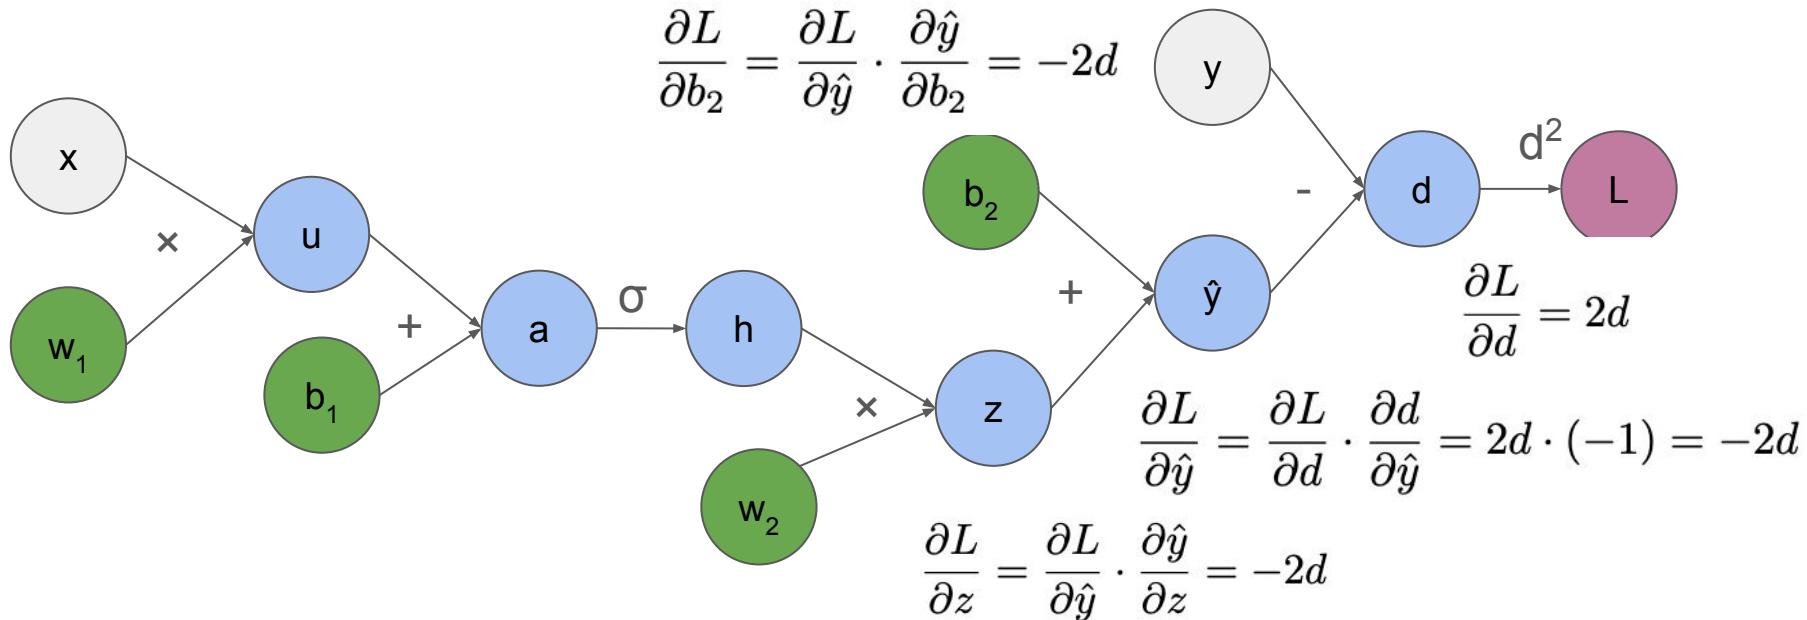
# Backprop: MLP (scalar)



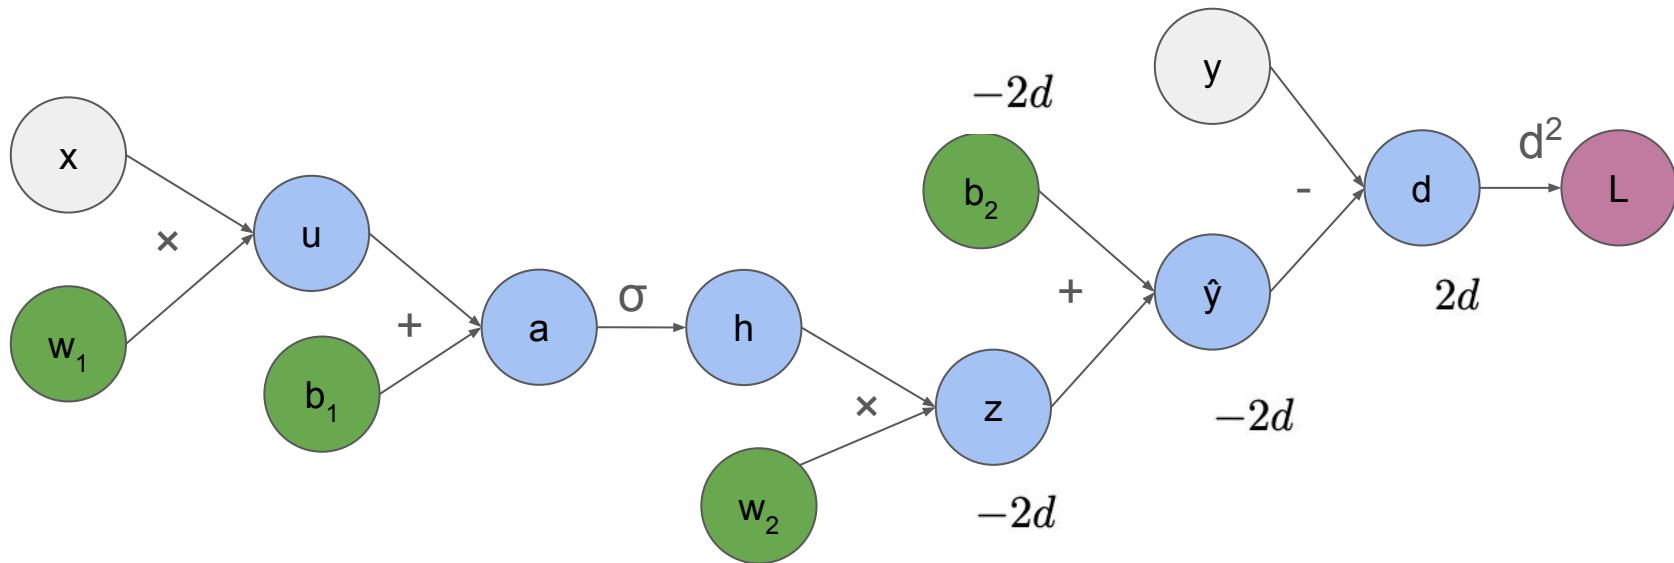
# Backprop: MLP (scalar)



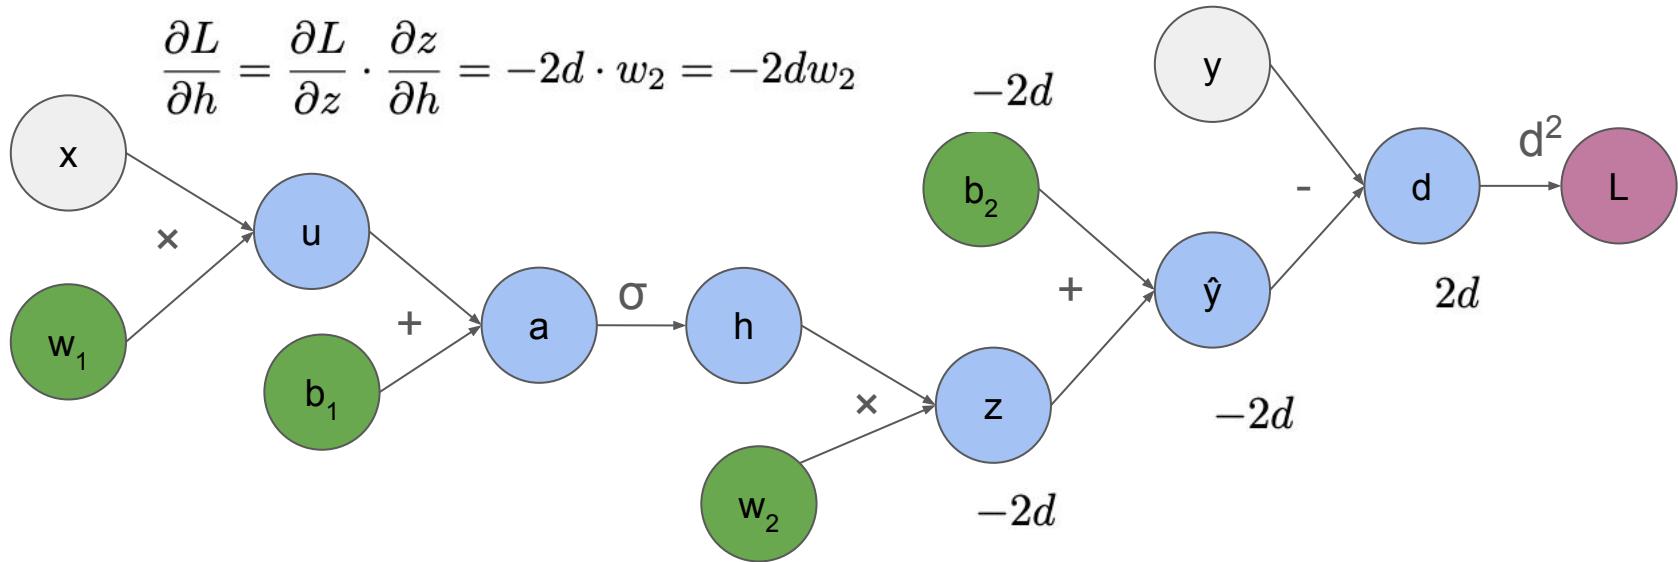
# Backprop: MLP (scalar)



# Backprop: MLP (scalar)

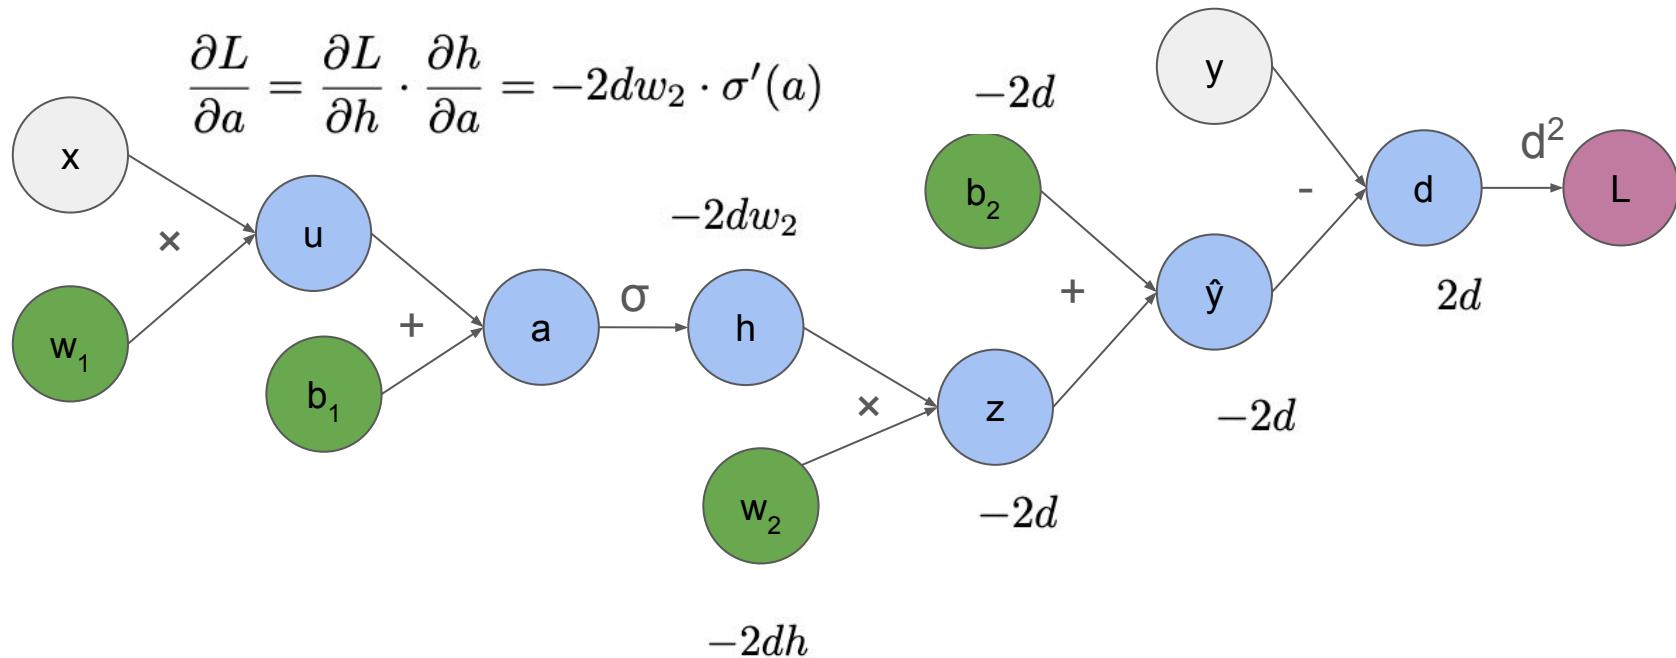


# Backprop: MLP (scalar)



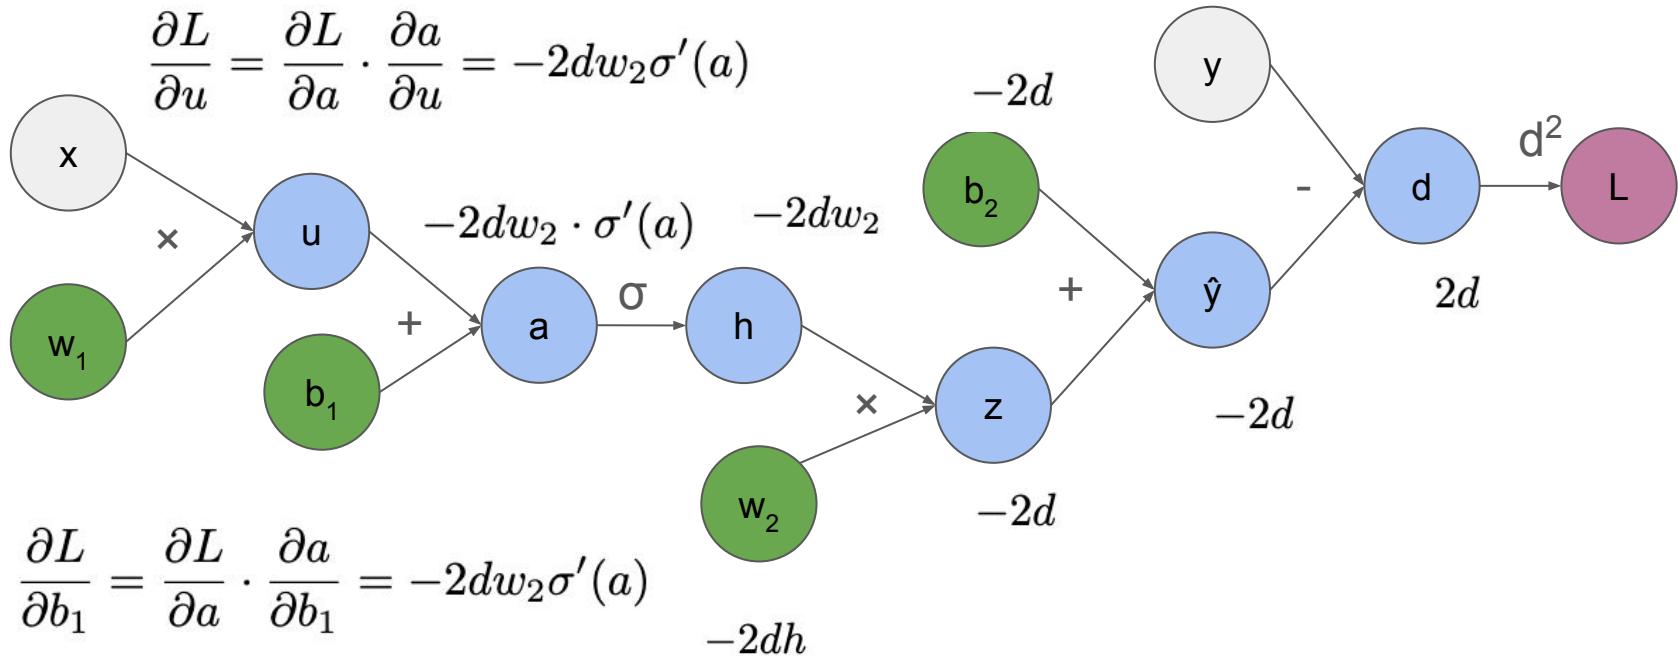
$$\frac{\partial L}{\partial w_2} = \frac{\partial L}{\partial z} \cdot \frac{\partial z}{\partial w_2} = -2d \cdot h = -2dh$$

# Backprop: MLP (scalar)



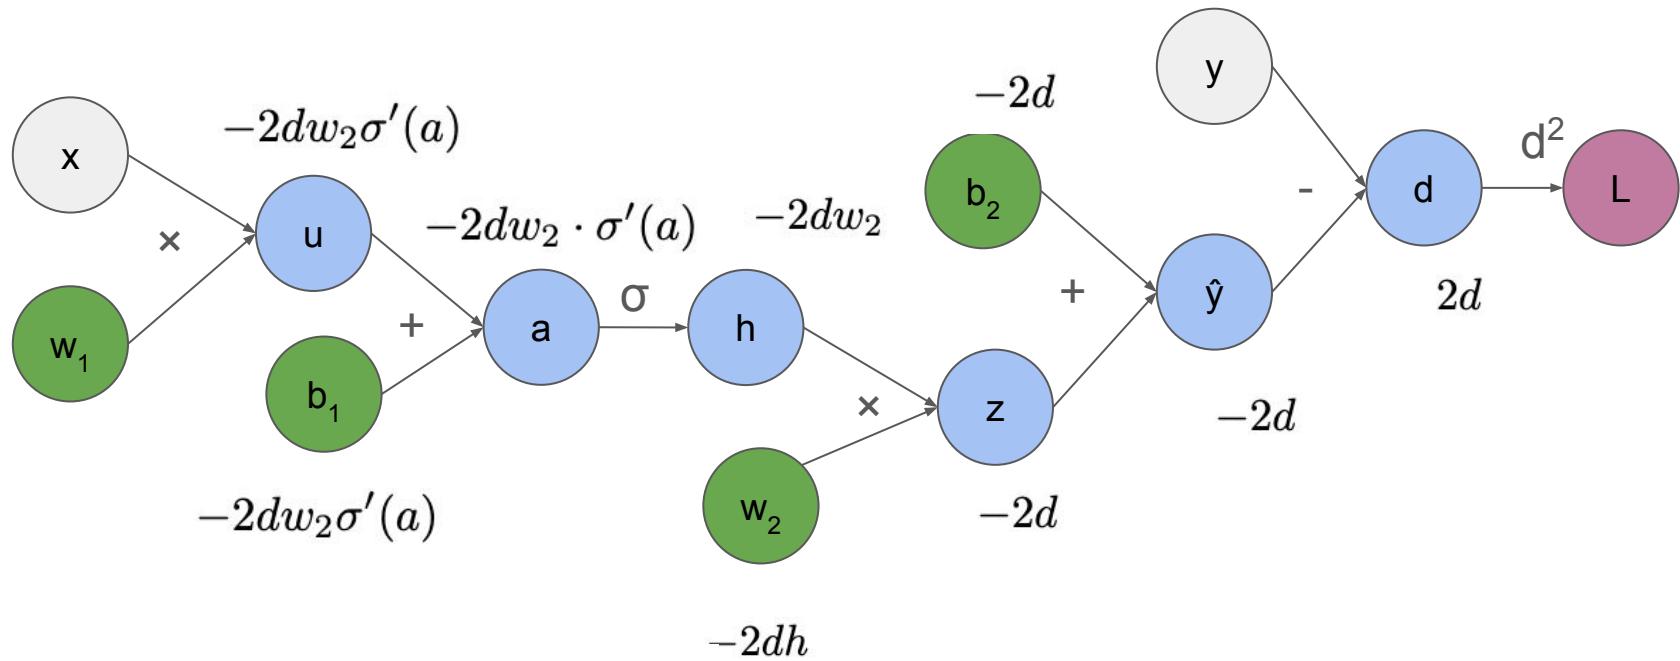
# Backprop: MLP (scalar)

$$\frac{\partial L}{\partial u} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial u} = -2dw_2\sigma'(a)$$



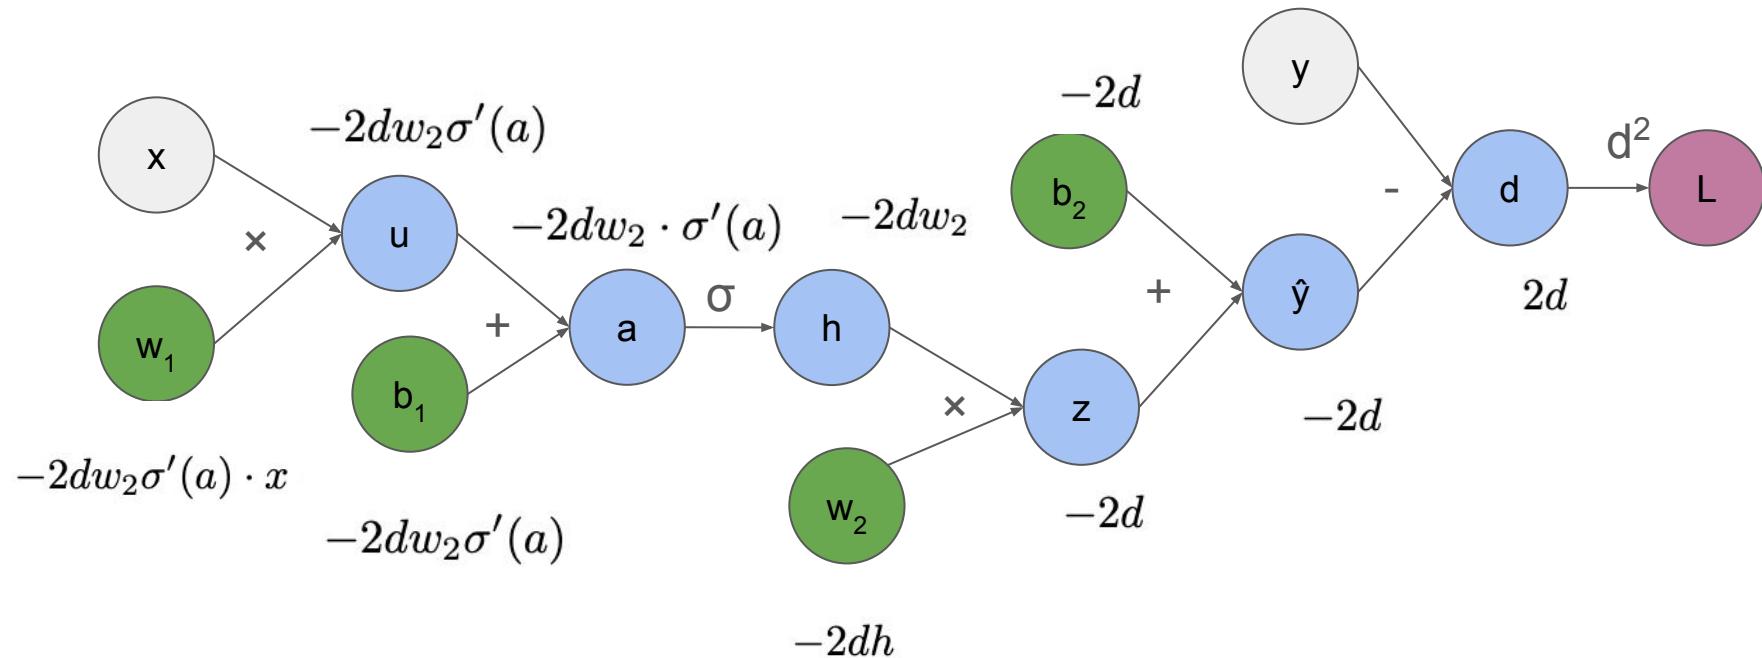
$$\frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial b_1} = -2dw_2\sigma'(a)$$

# Backprop: MLP (scalar)



$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial u} \cdot \frac{\partial u}{\partial w_1} = -2dw_2\sigma'(a) \cdot x$$

# Backprop: MLP (scalar)



We computed the gradients with respect to all the parameters

- We can now update all the parameters with gradient descent

# Backprop: MLP general case

$$x \in \mathbb{R}^{n_{in}}, W_1 \in \mathbb{R}^{n_h \times n_{in}}, b_1 \in \mathbb{R}^{n_h}$$

$$h \in \mathbb{R}^{n_h}, W_2 \in \mathbb{R}^{n_{out} \times n_h}, b_2 \in \mathbb{R}^{n_{out}}$$

$$y, \hat{y}, d \in \mathbb{R}^{n_{out}}, L \in \mathbb{R}$$

$$u = W_1 x$$

$$a = u + b_1$$

$$h = \sigma(a)$$

$$z = W_2 h$$

$$\hat{y} = z + b_2$$

$$d = \hat{y} - y$$

$$L = d^T d = \|d\|^2$$

$$\frac{\partial L}{\partial L} = 1$$

$$\frac{\partial L}{\partial d} = 2d$$

$$\frac{\partial L}{\partial \hat{y}} = 2d$$

$$\frac{\partial L}{\partial b_2} = 2d$$

$$\frac{\partial L}{\partial z} = 2d$$

$$\frac{\partial L}{\partial W_2} = 2d h^T$$

$$\frac{\partial L}{\partial h} = W_2^T (2d)$$

$$\frac{\partial L}{\partial a} = \frac{\partial L}{\partial h} \odot \sigma'(a)$$

$$\frac{\partial L}{\partial b_1} = \frac{\partial L}{\partial a}$$

$$\frac{\partial L}{\partial u} = \frac{\partial L}{\partial a}$$

$$\frac{\partial L}{\partial W_1} = \frac{\partial L}{\partial u} x^T$$

Forward pass

Backward pass

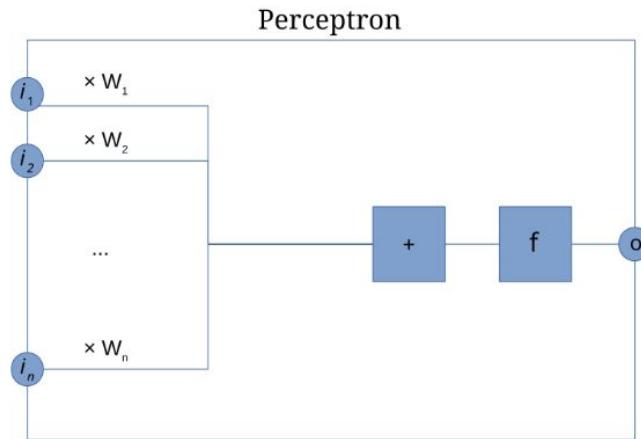


# History

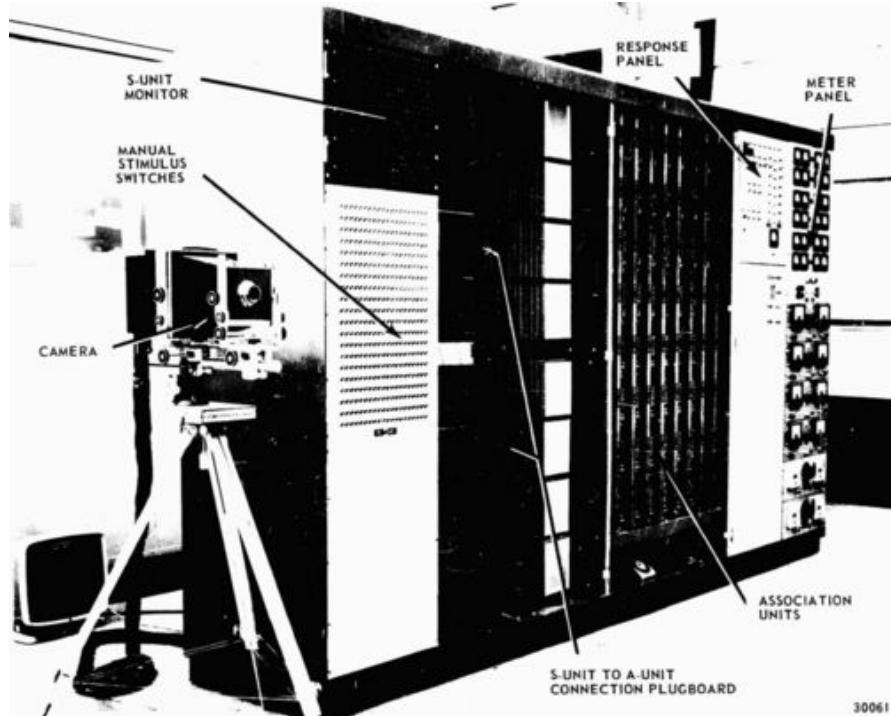
0.36

0.34

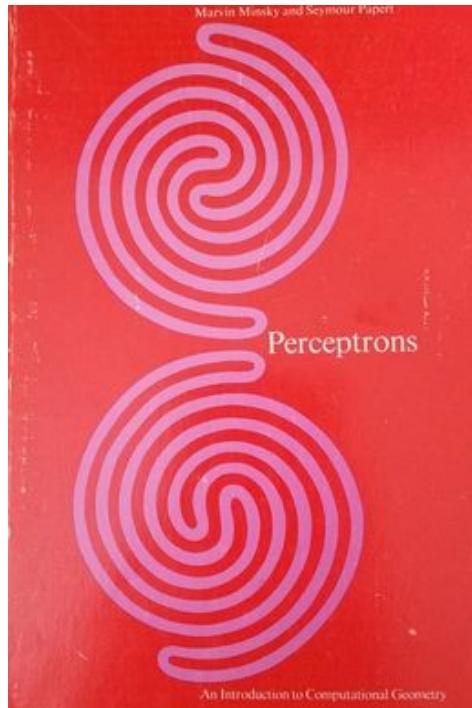
# Rosenblatt Perceptron (1957)



$$o = f\left(\sum_{k=1}^n i_k \cdot W_k\right)$$



# Marvin Minsky, Seymour Papert (1969)



Showed that perceptrons cannot represent XOR

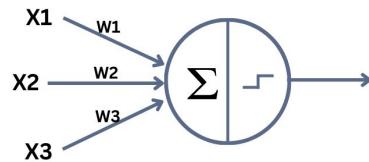


2 input XOR gate		
A	B	$A \oplus B$
0	0	0
0	1	1
1	0	1
1	1	0

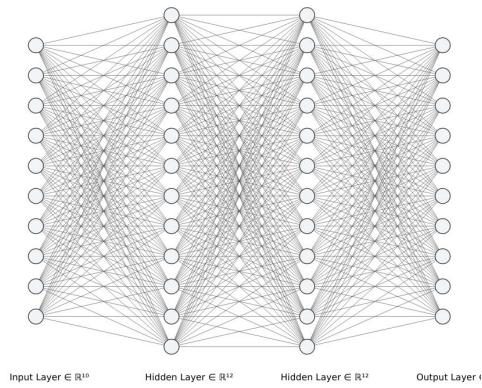
Caused AI winter

# Hinton et al: Backpropagation (1986)

Showed how to train Multi-Layer Perceptrons



Single-layer perceptron

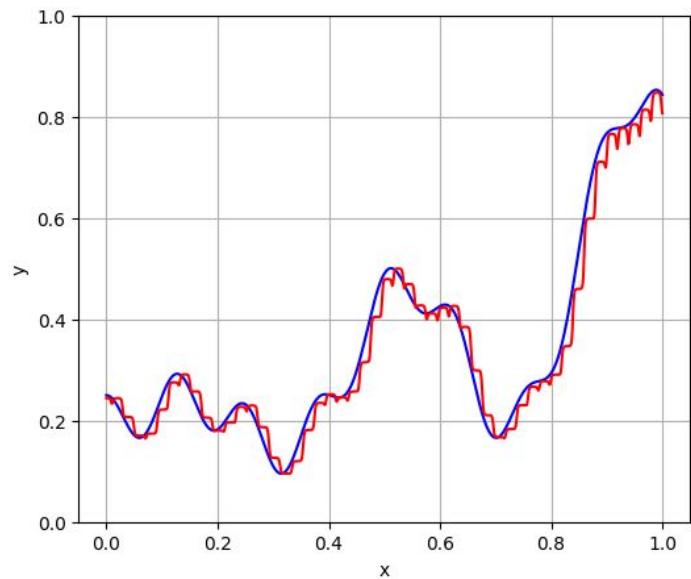
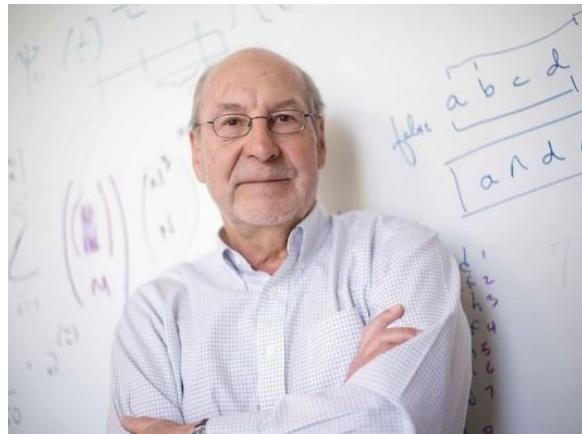


Multi-layer perceptron



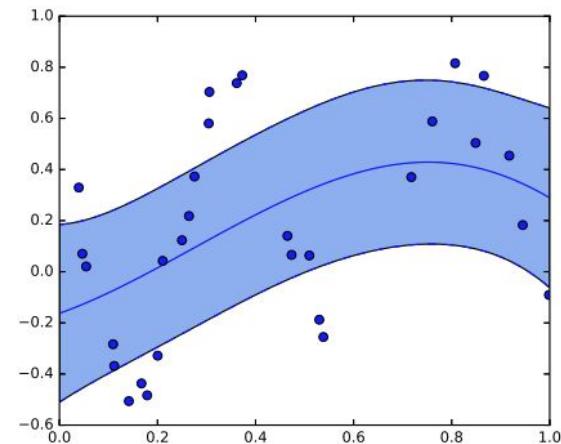
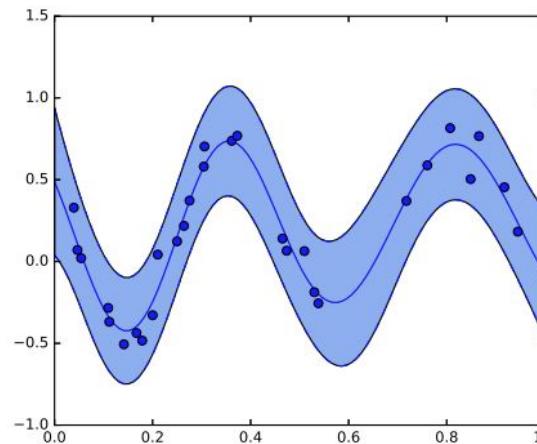
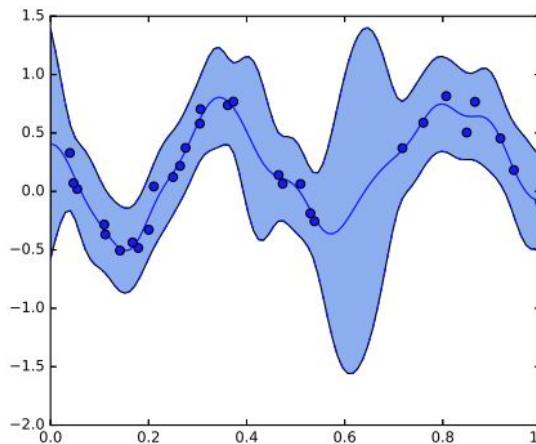
# Universal approximation theorem: Cybenko (1989)

Showed MLPs can approximate arbitrary functions



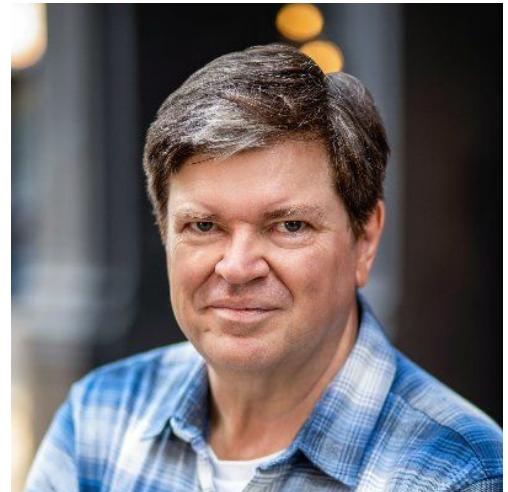
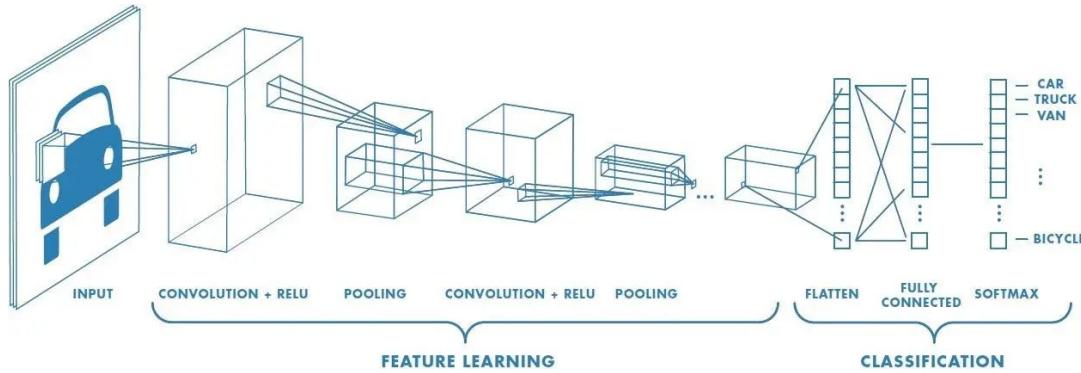
# Kernel Methods (1990-2000)

Another winter: kernel methods work better and are more interpretable



# ConvNets: Yann Lecun (1989)

Yann invented CNNs which eventually became state-of-the art for image processing



# ImageNet Moment (2012)



Ilya Sutskever

Start of the current era, neural nets outcompete all other methods on an image processing problem.



# ImageNet Moment (2012)

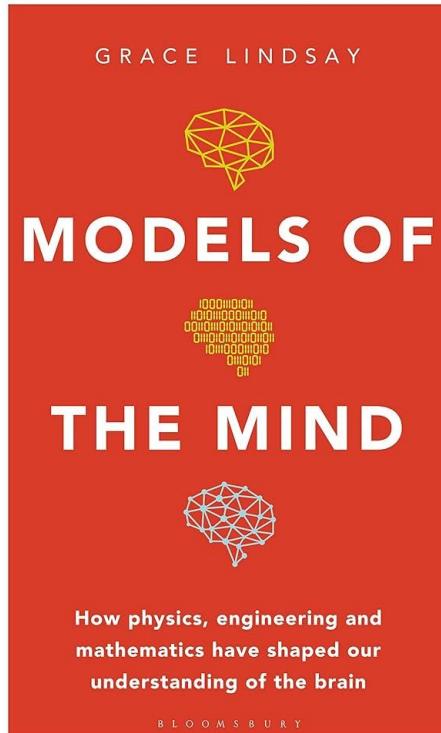


Ilya Sutskever

Start of the current era, neural nets outcompete all other methods on an image processing problem.



# More history



Great book if you want to learn more about the history of neural networks.